

Pagoda: Towards Binary Code Privacy Protection with SGX-based Execute-Only Memory

JIYONG YU¹, XINYANG GE², TRENT JAEGER³, CHRIS FLETCHER¹, WEIDONG CUI⁴



¹UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN¹


²DATABRICKS

³PENNSYLVANIA STATE UNIVERSITY



⁴MICROSOFT RESEARCH




Code Privacy in Untrusted Environments



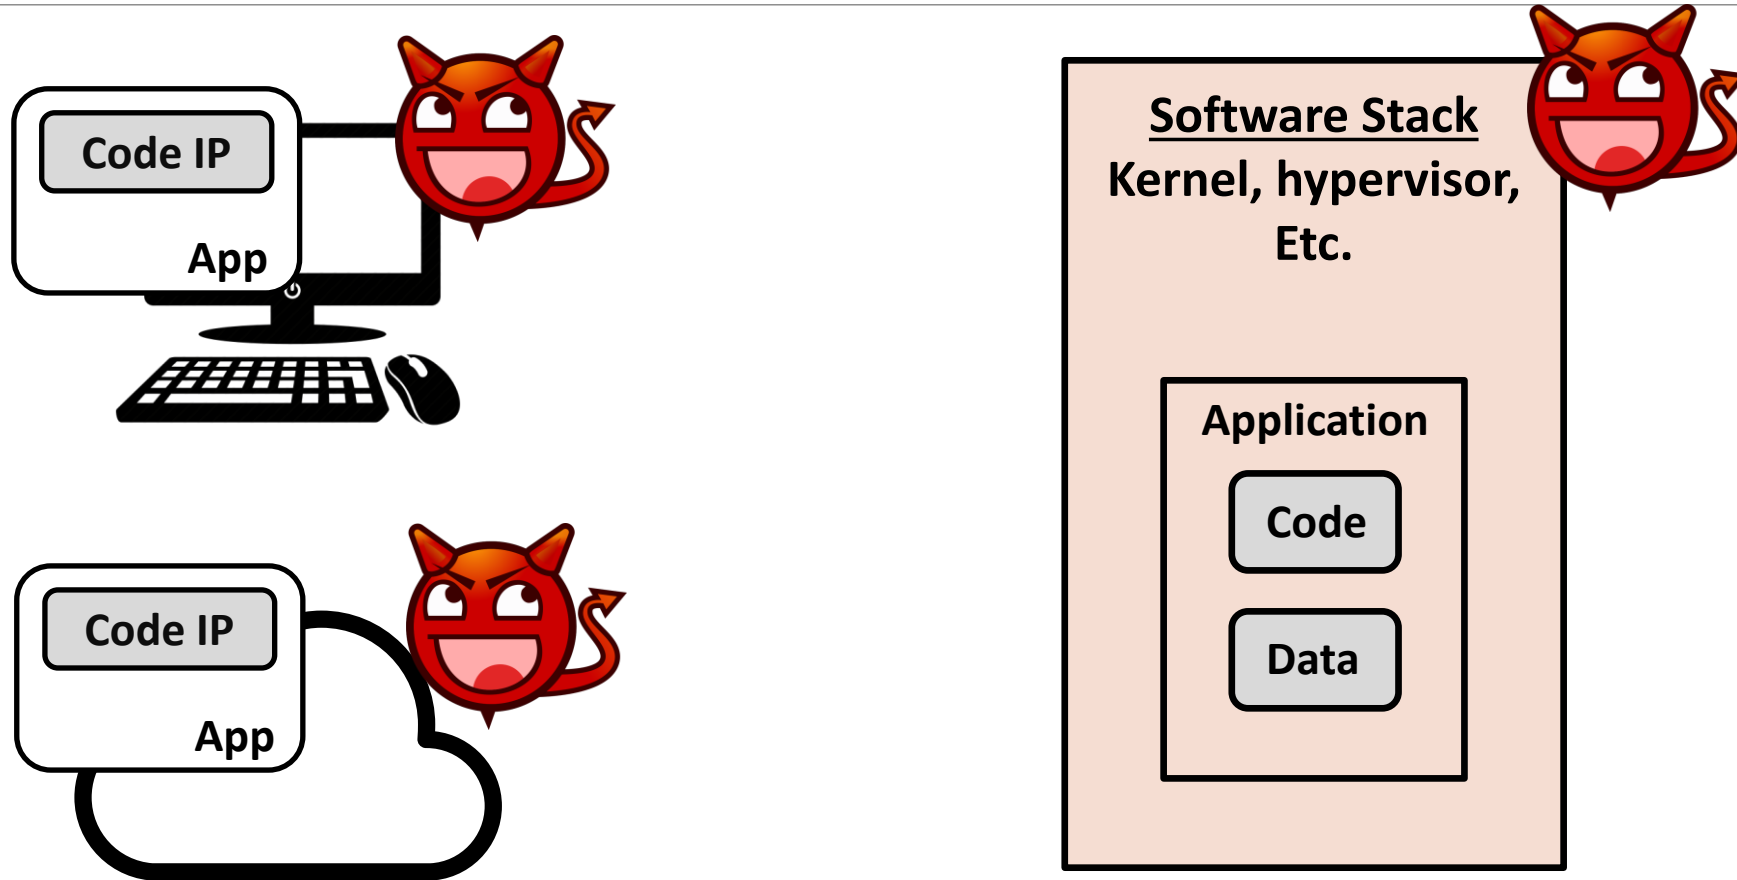
Code Privacy in Untrusted Environments



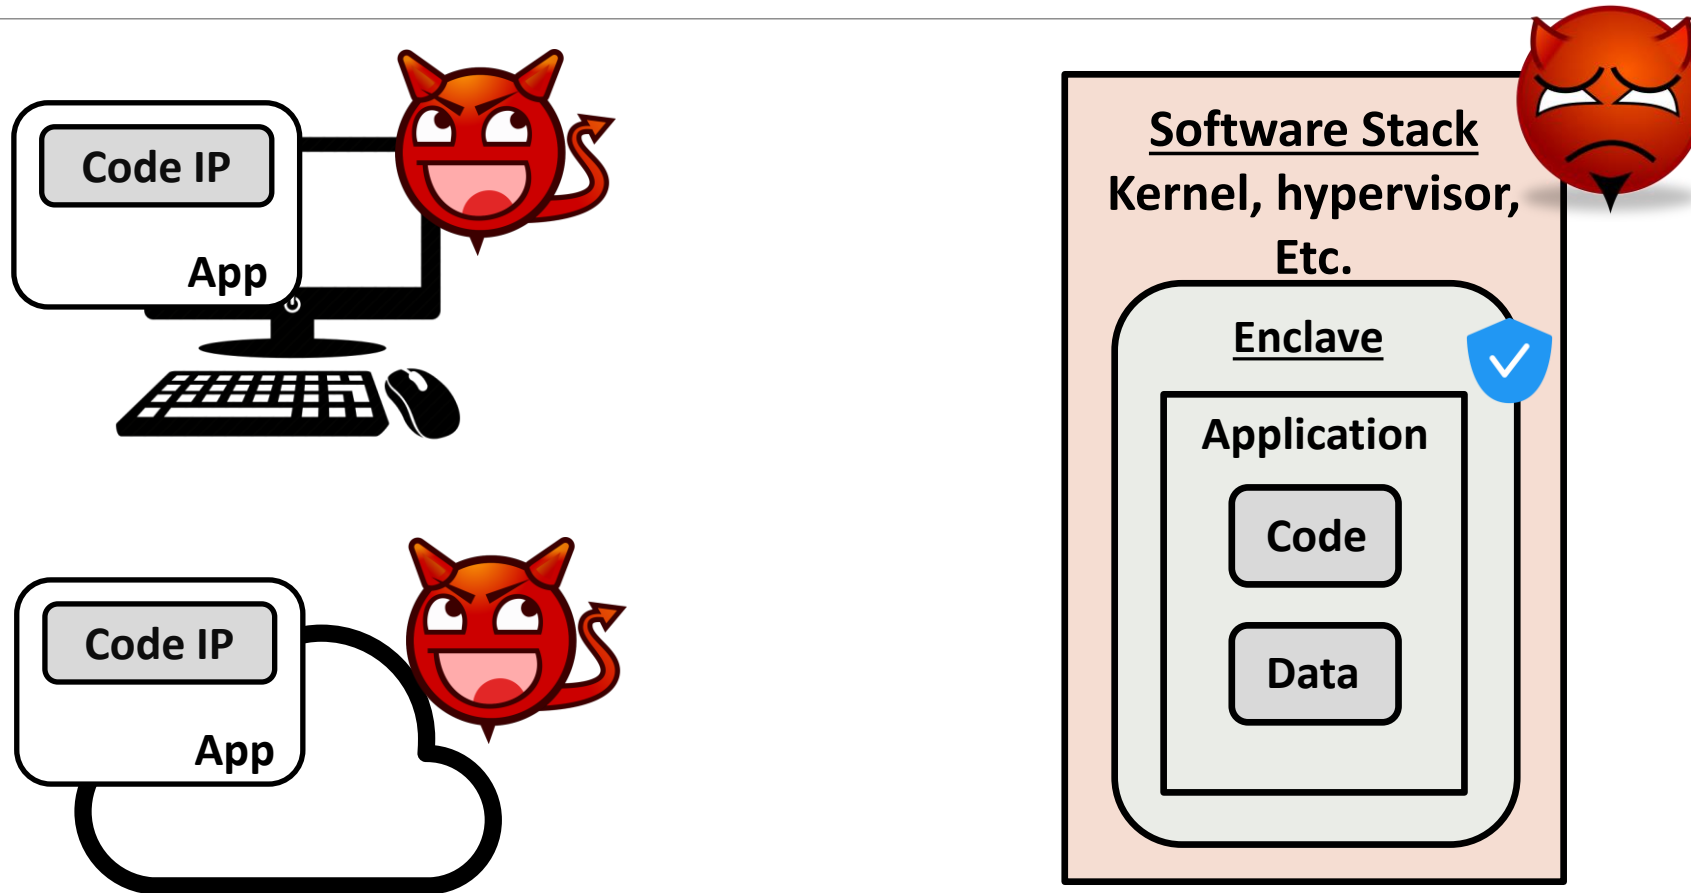
Code Privacy in Untrusted Environments




Intel SGX Comes to Rescue



Intel SGX Comes to Rescue




Protecting Large Applications using SGX is Challenging




- Enclave programming model requires significant software refactoring

Protecting Large Applications using SGX is Challenging




- Enclave programming model requires significant software refactoring
- Existing “plug-and-play” SGX frameworks (e.g., GrapheneSGX)
 - Assuming code is bug-free

Protecting Large Applications using SGX is Challenging



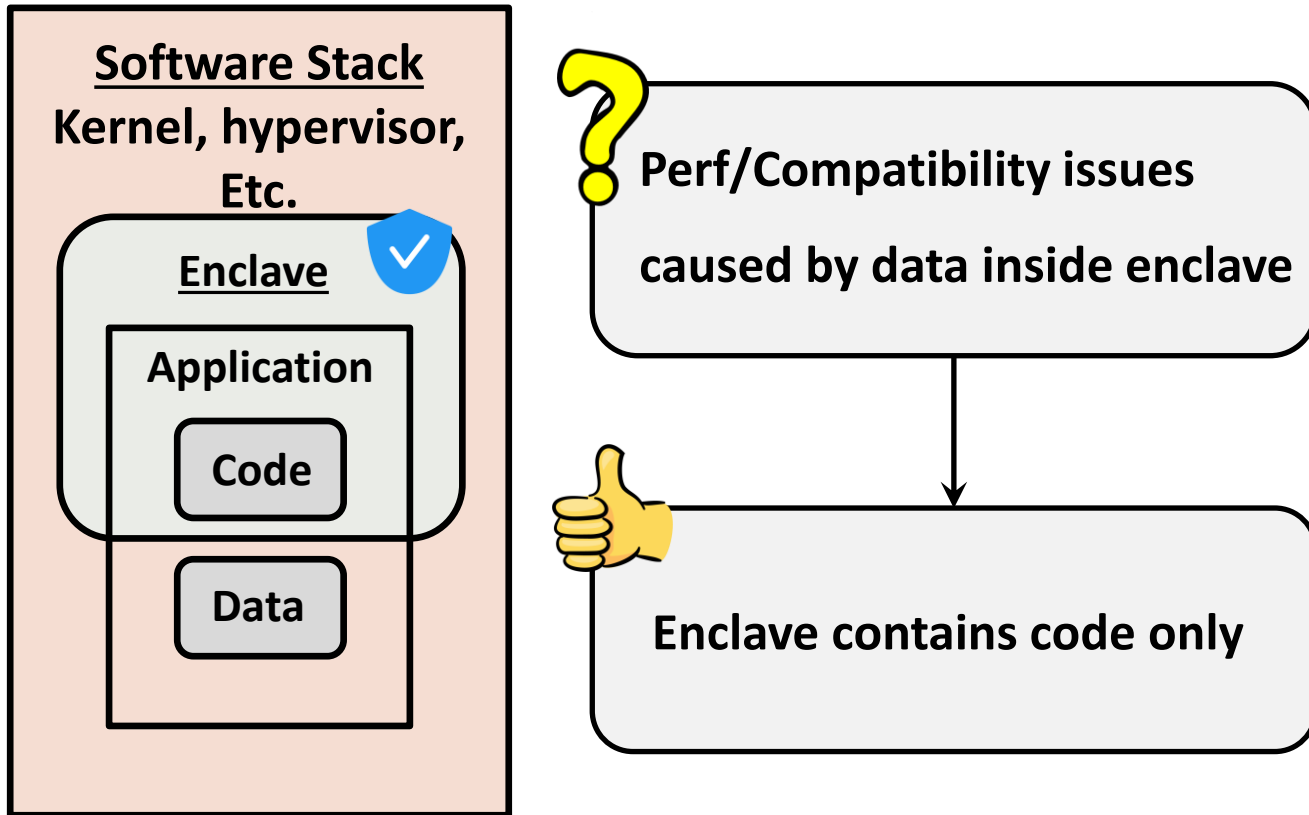
- Enclave programming model requires significant software refactoring
- Existing “plug-and-play” SGX frameworks (e.g., GrapheneSGX)
 - Assuming code is bug-free
 - Explicit data movement during system calls

Key Idea of Pagoda




? Perf/Compatibility issues
caused by data inside enclave

Key Idea of Pagoda



Key Idea of Pagoda




? Perf/Compatibility issues
caused by data inside enclave

👍 Enclave contains code only

? Vulnerabilities in code
Manipulating data → Code disclosure
(read-out)

Key Idea of Pagoda




? Perf/Compatibility issues
caused by data inside enclave

👍 Enclave contains code only

? Vulnerabilities in code
Manipulating data → Code disclosure
(read-out)

👍 Code in enclave must always be
execute-only

Key Idea of Pagoda




? Perf/Compatibility issues
caused by data inside enclave

👍 Enclave contains code only

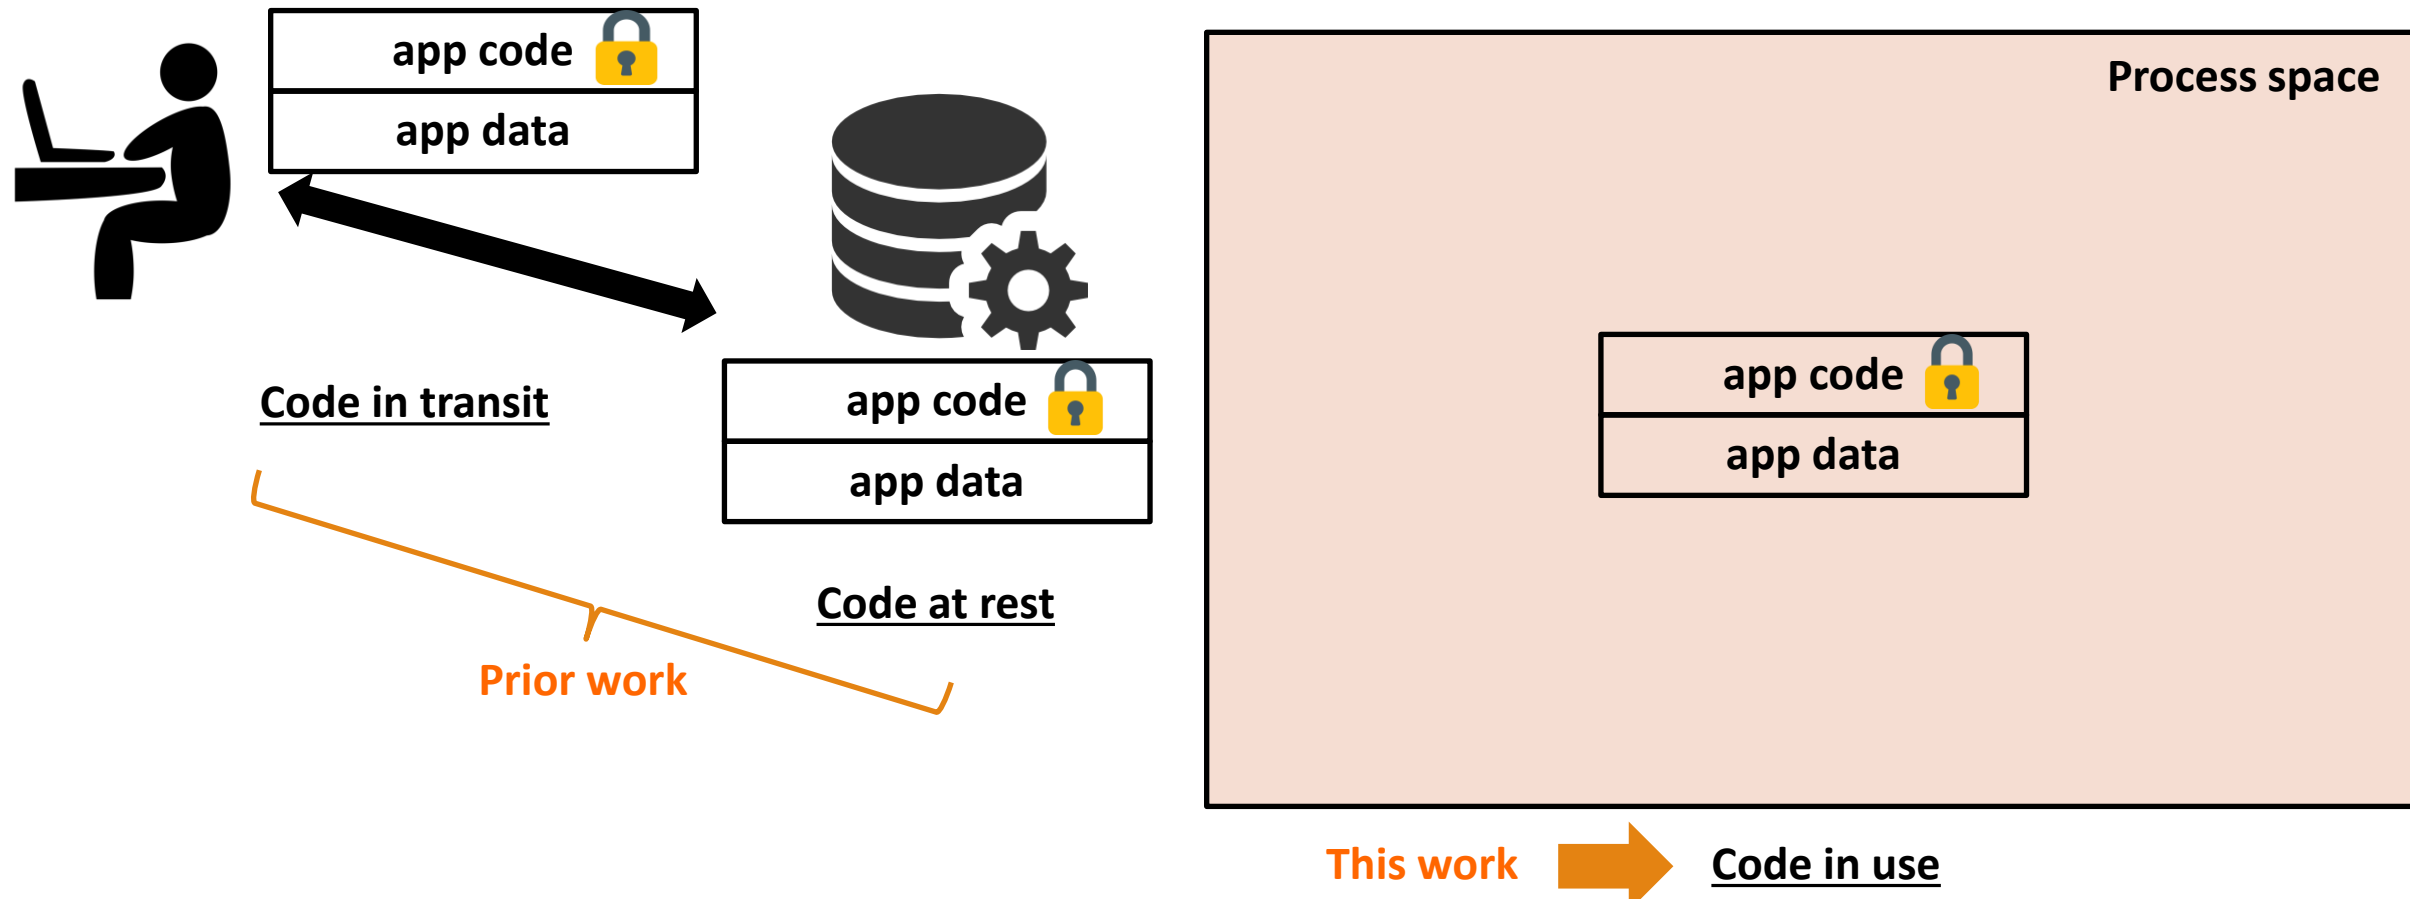
? Vulnerabilities in code
Manipulating data → Code disclosure
(read-out)

👍 Code in enclave must **always** be
execute-only

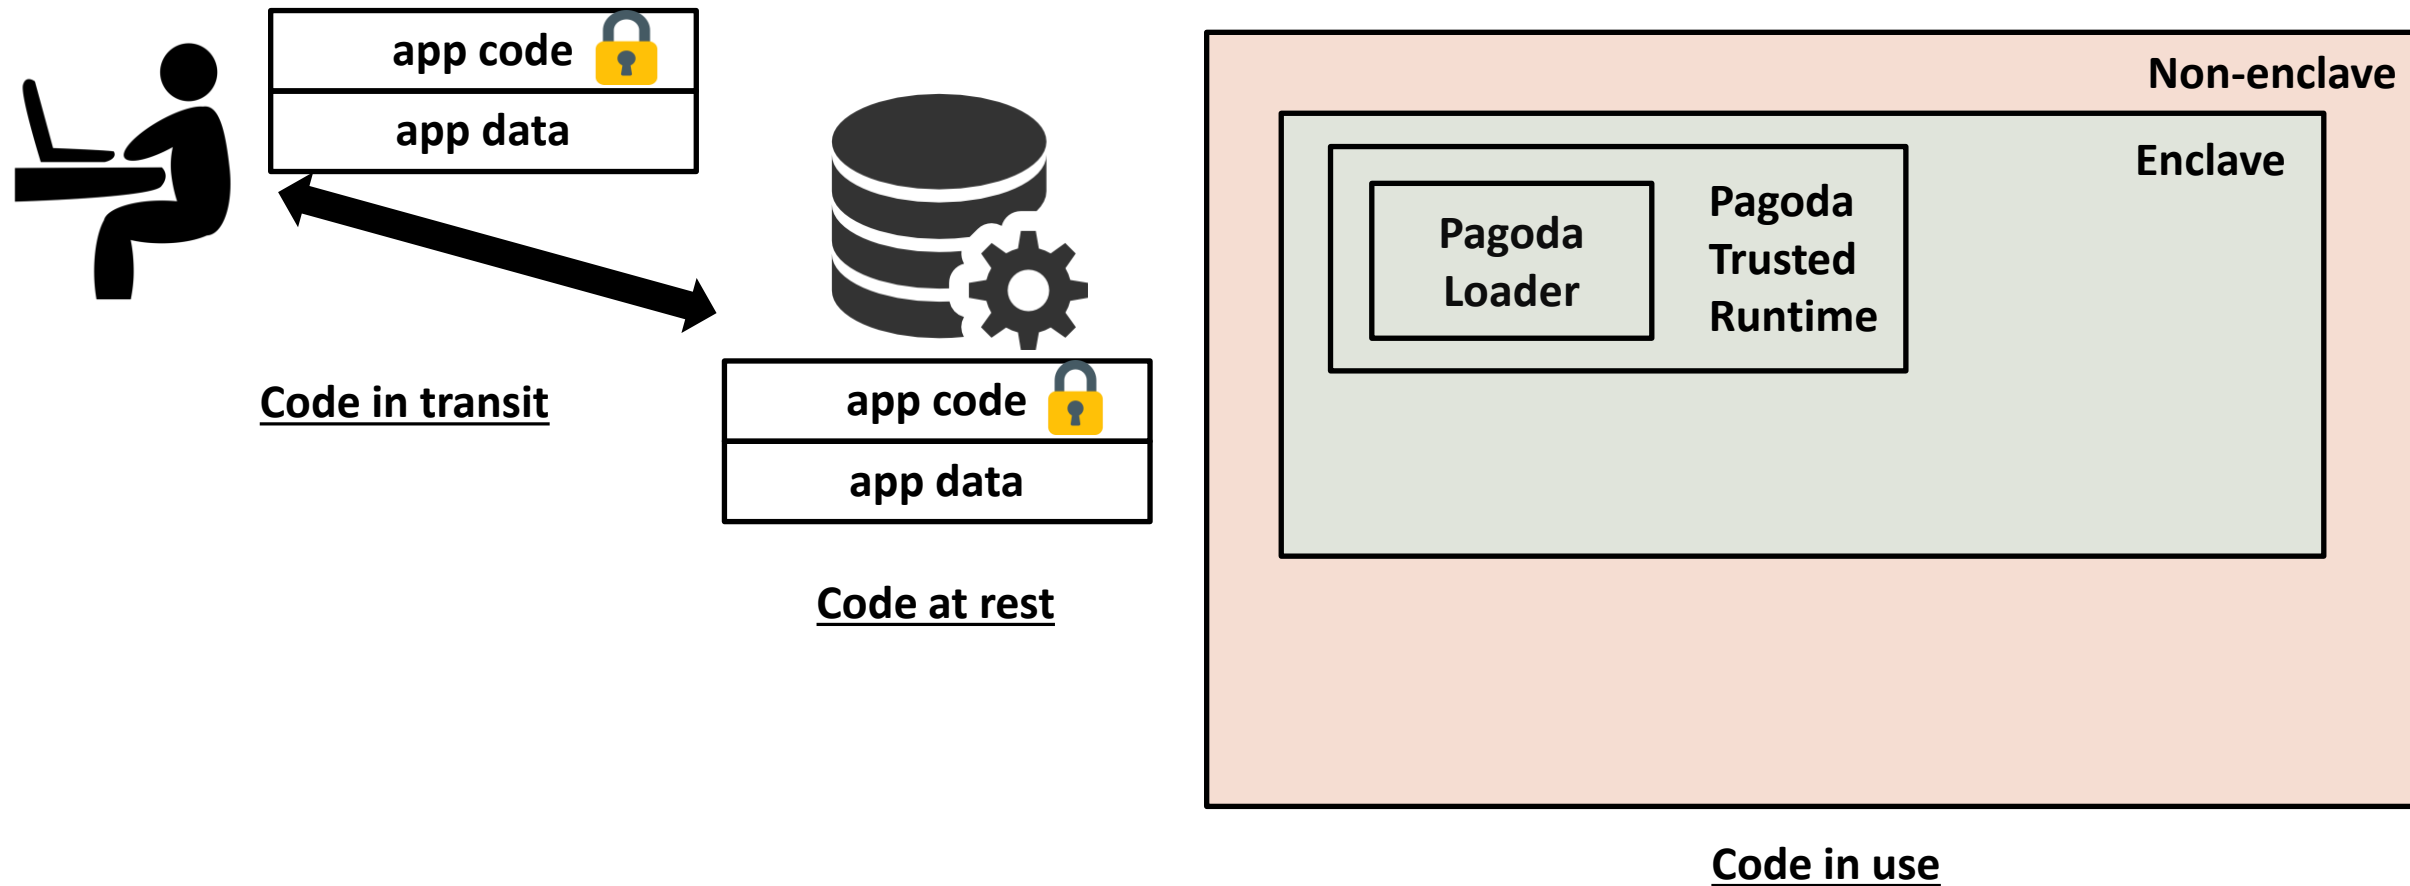
Deploying a Protected Software



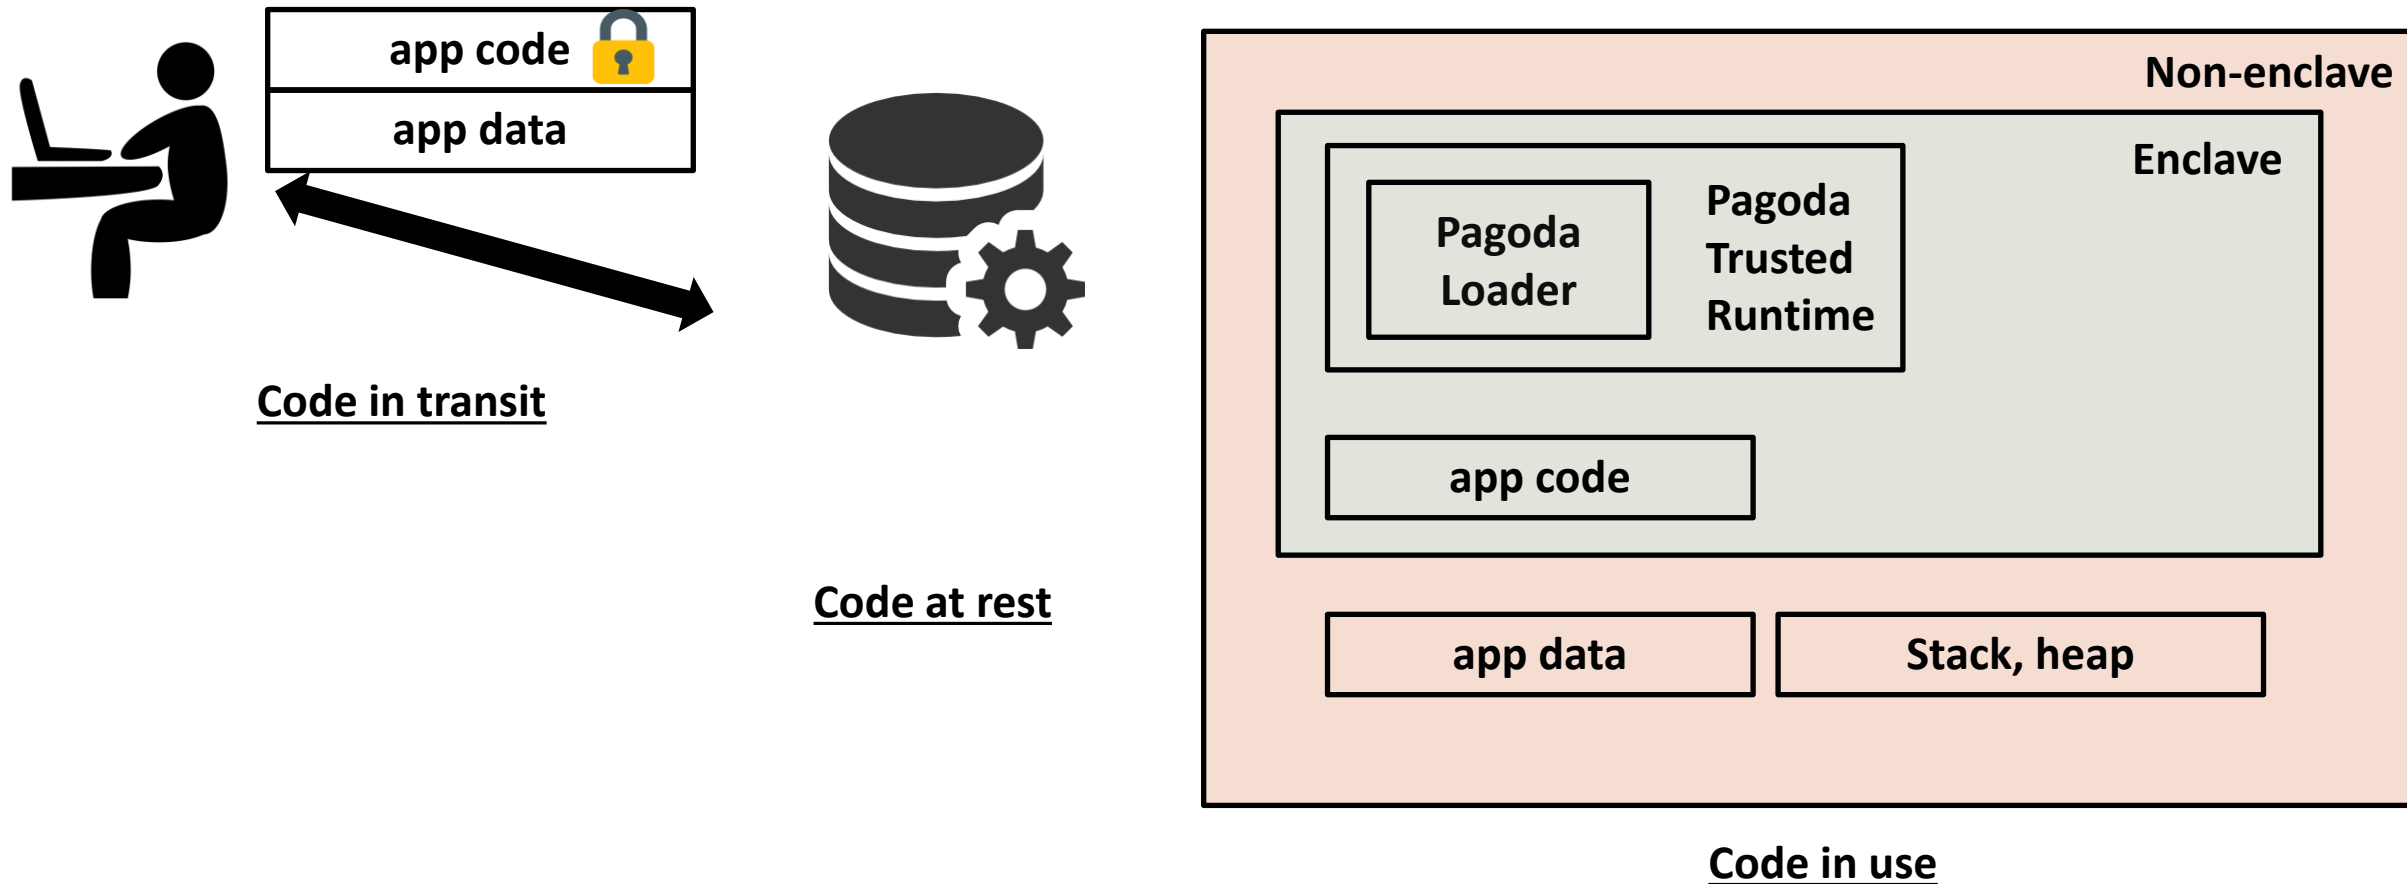
Deploying a Protected Software



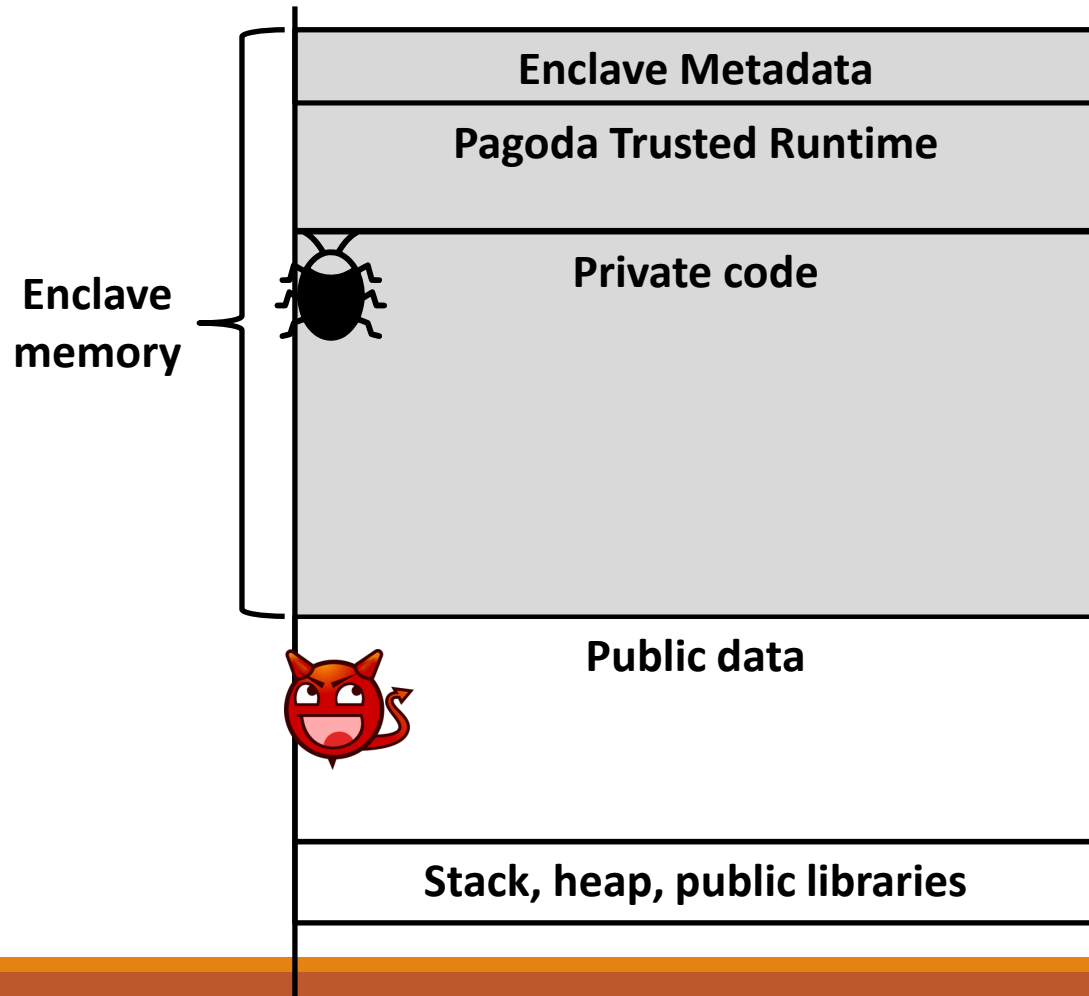
Deploying a Protected Software



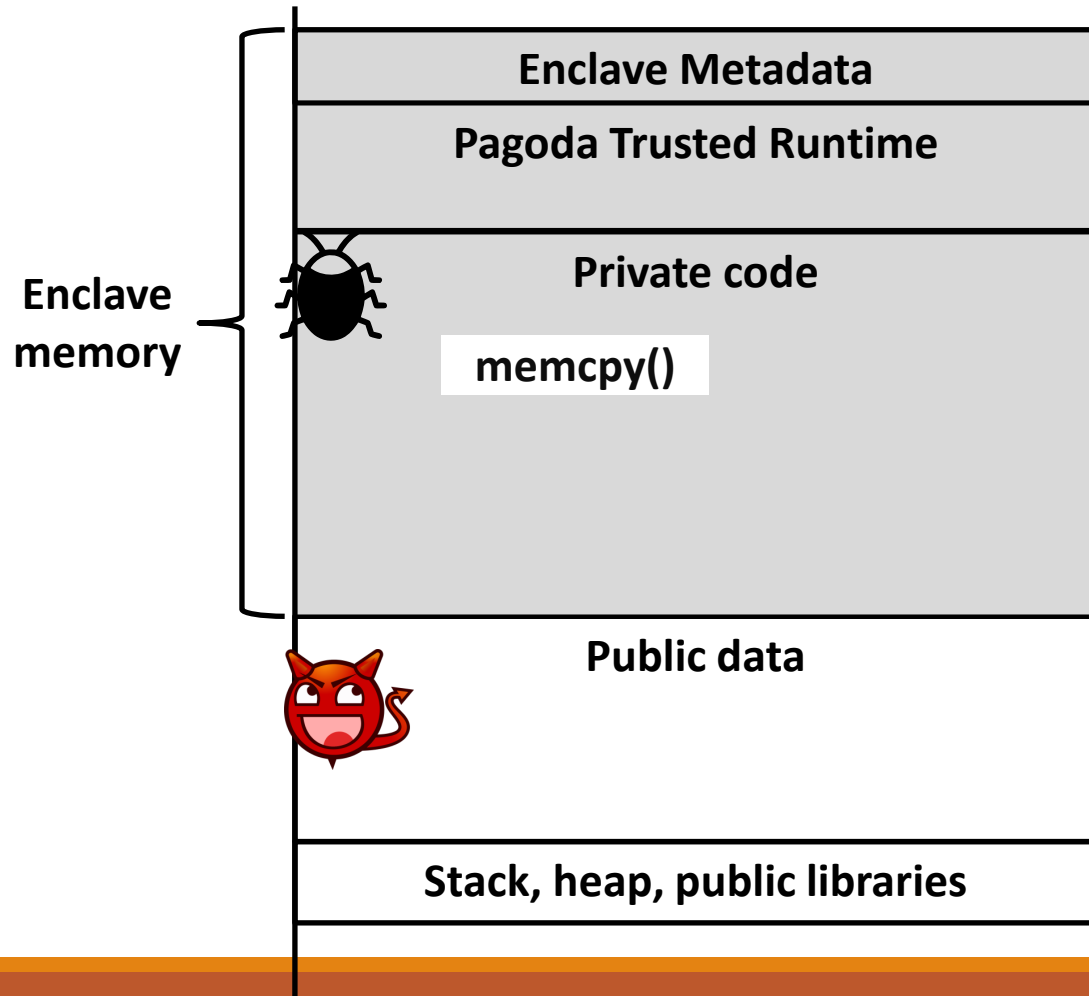
Deploying a Protected Software



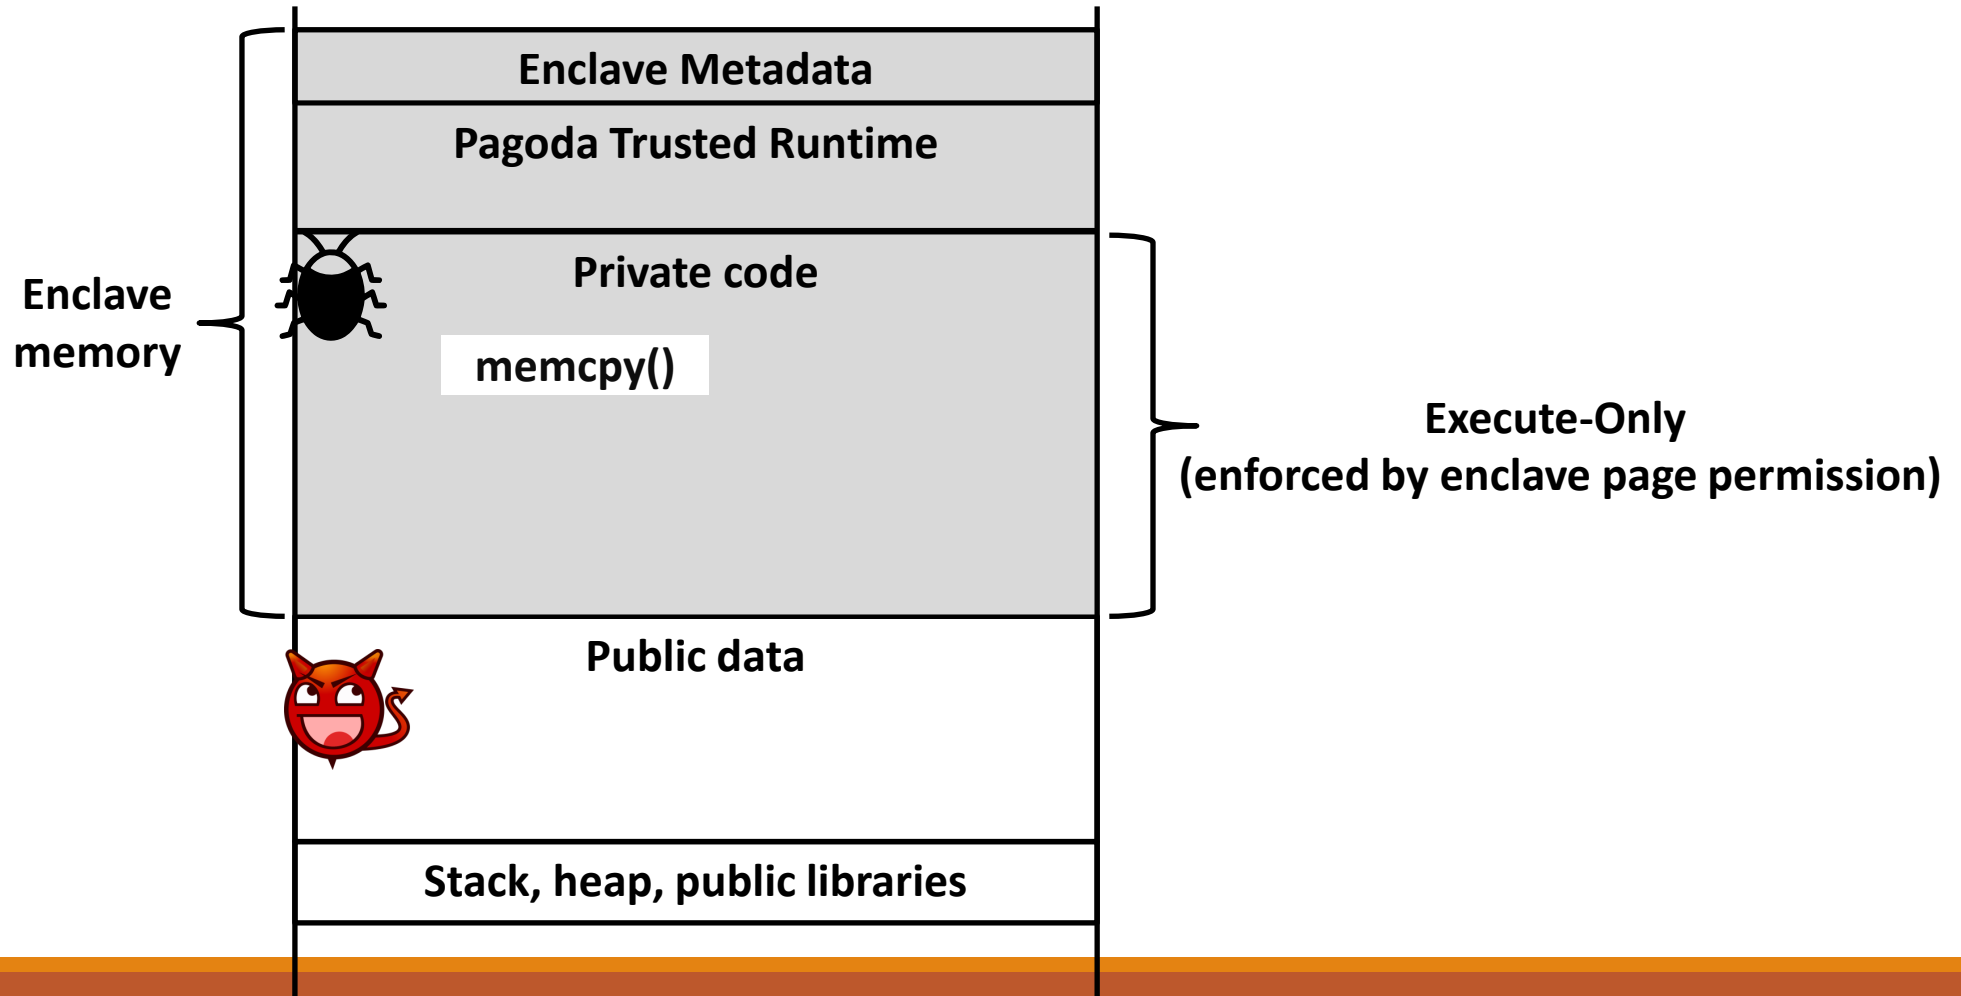
Enforcing Execute-Only Property



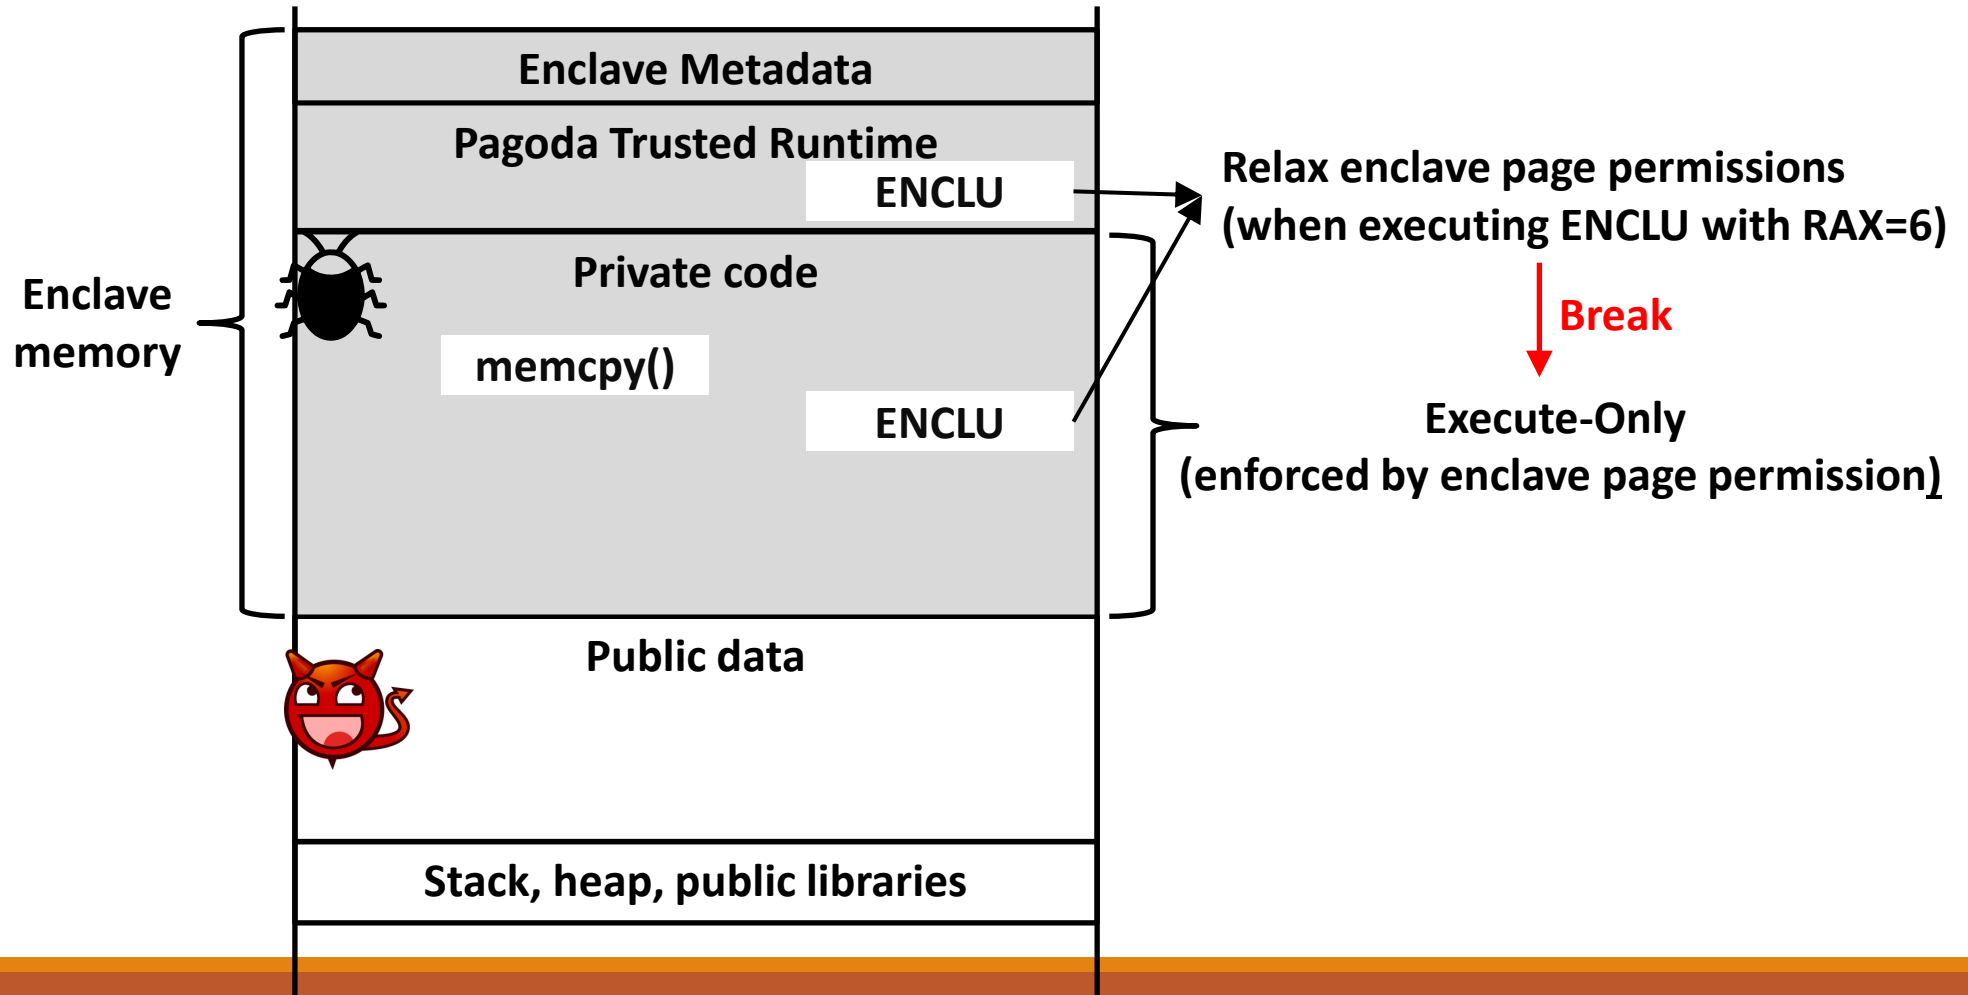
Enforcing Execute-Only Property



Enforcing Execute-Only Property

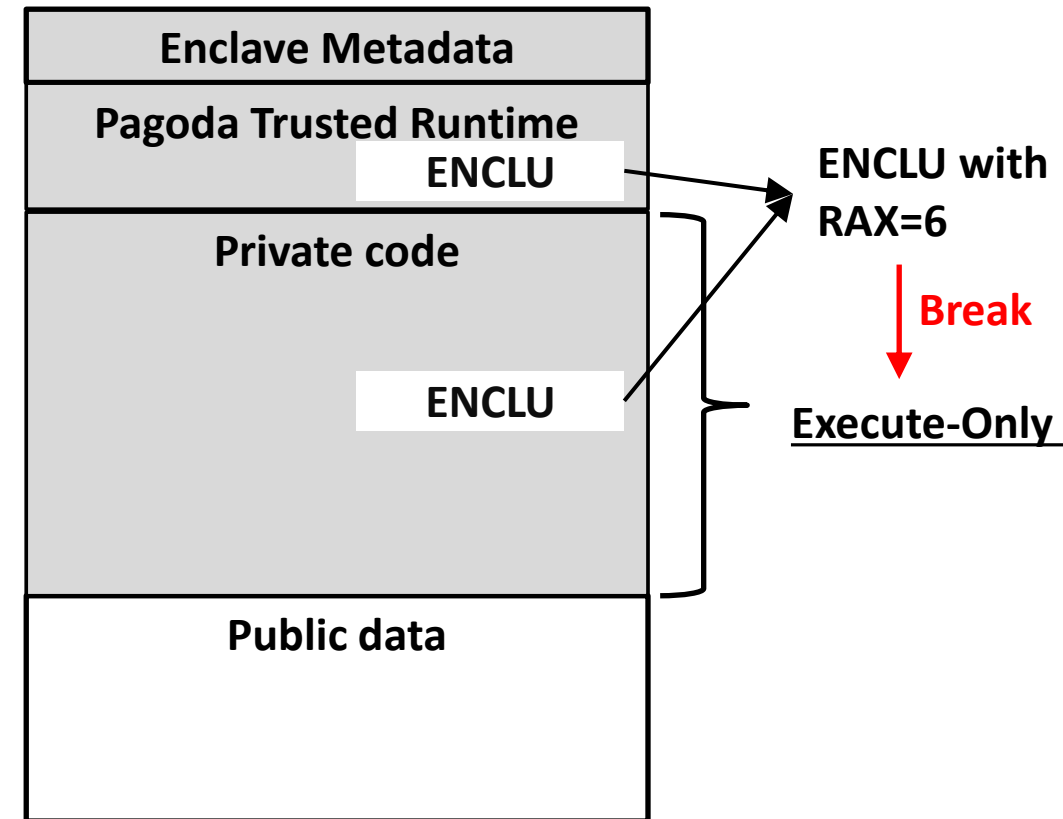


Preventing Relaxation of Execute-Only



Preventing Relaxation of Execute-Only

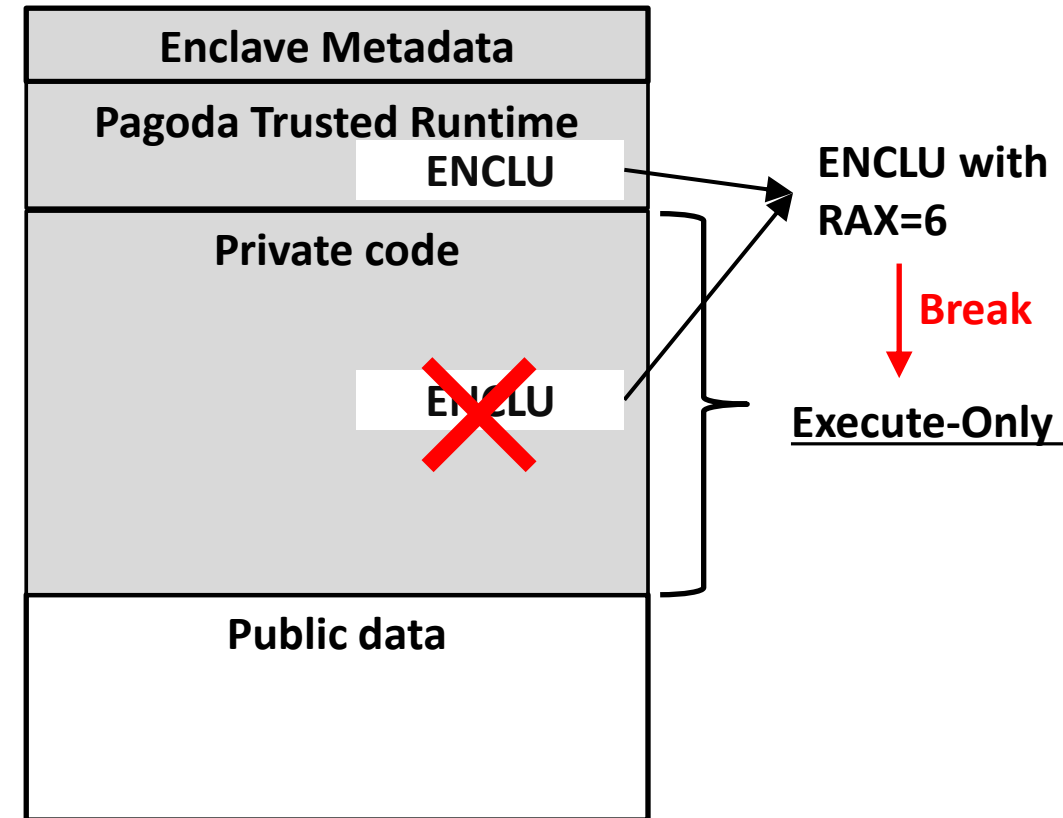
- Strategies:



Preventing Relaxation of Execute-Only

- Strategies:

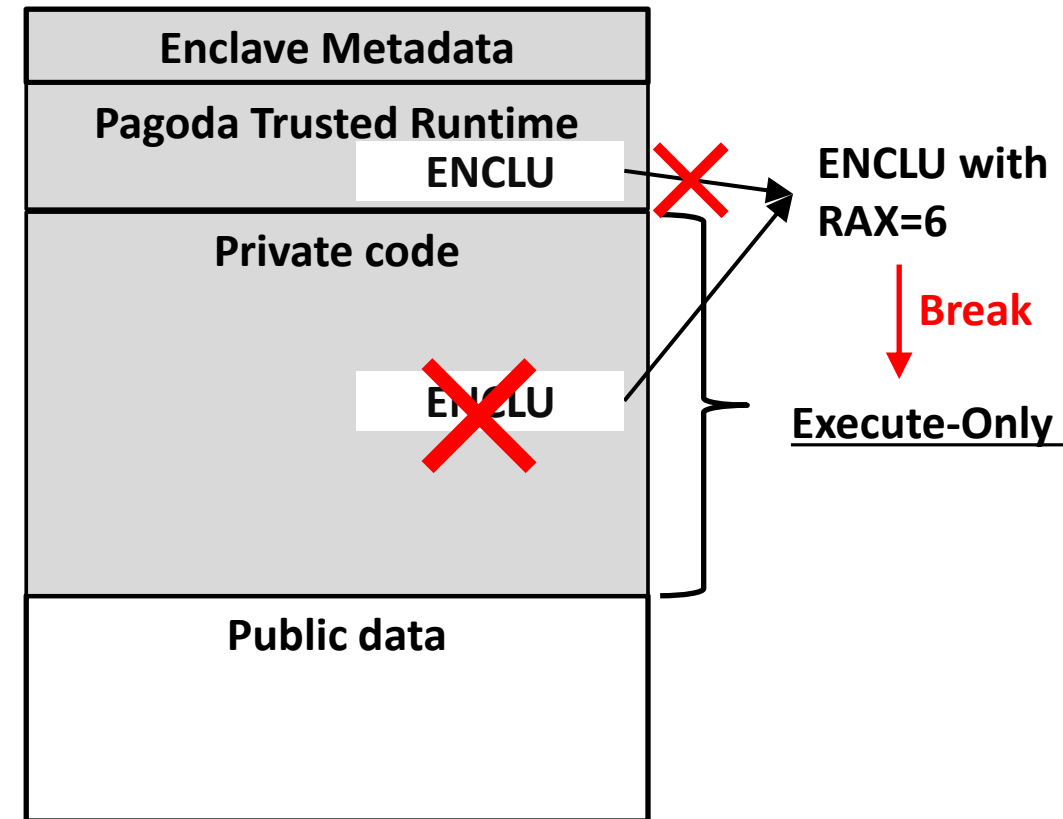
1. No ENCLU (3-byte sequence {0x0F, 0x01, 0xD7}) in the application code



Preventing Relaxation of Execute-Only

■ Strategies:

1. No ENCLU (3-byte sequence {0x0F, 0x01, 0xD7}) in the application code
2. All ENCLU in the Pagoda Trusted Runtime cannot be used with RAX=6



Preventing Relaxation of Execute-Only

```
/* Pagoda trusted runtime */
```

```
--enclu-----
```

```
  cmp rax, 6 // check if ENCLU is  
              // used w/ rax=6
```

```
  jne cont
```

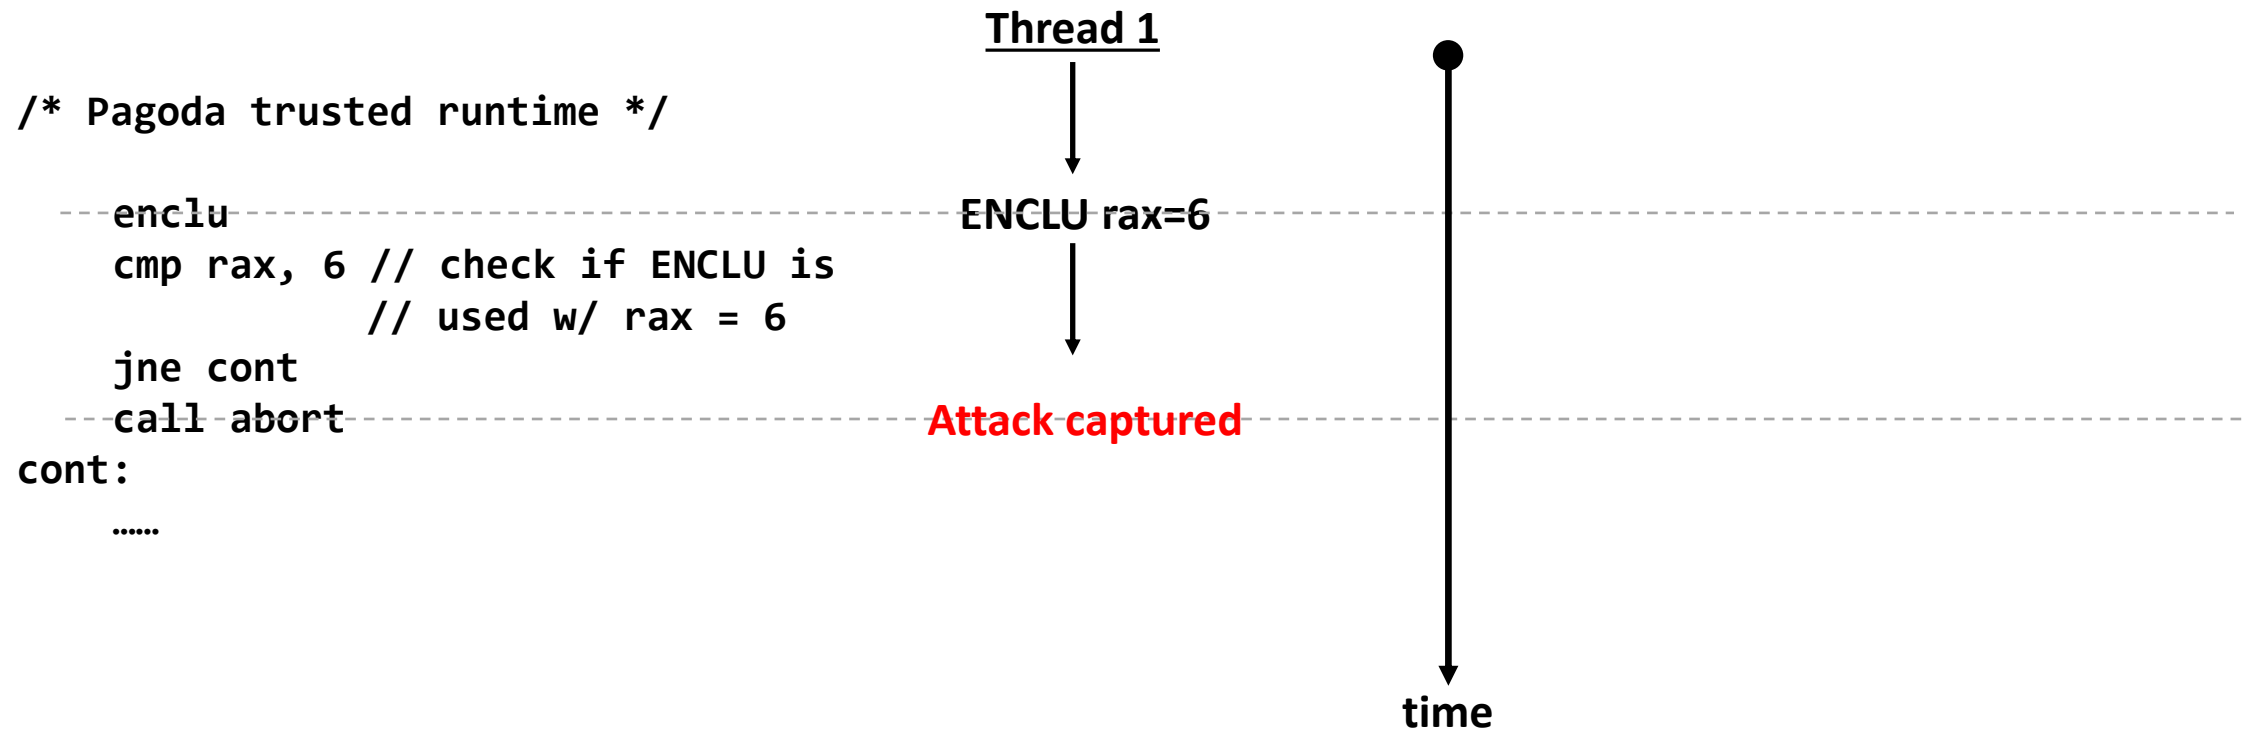
```
--call abort-----
```

```
cont:
```

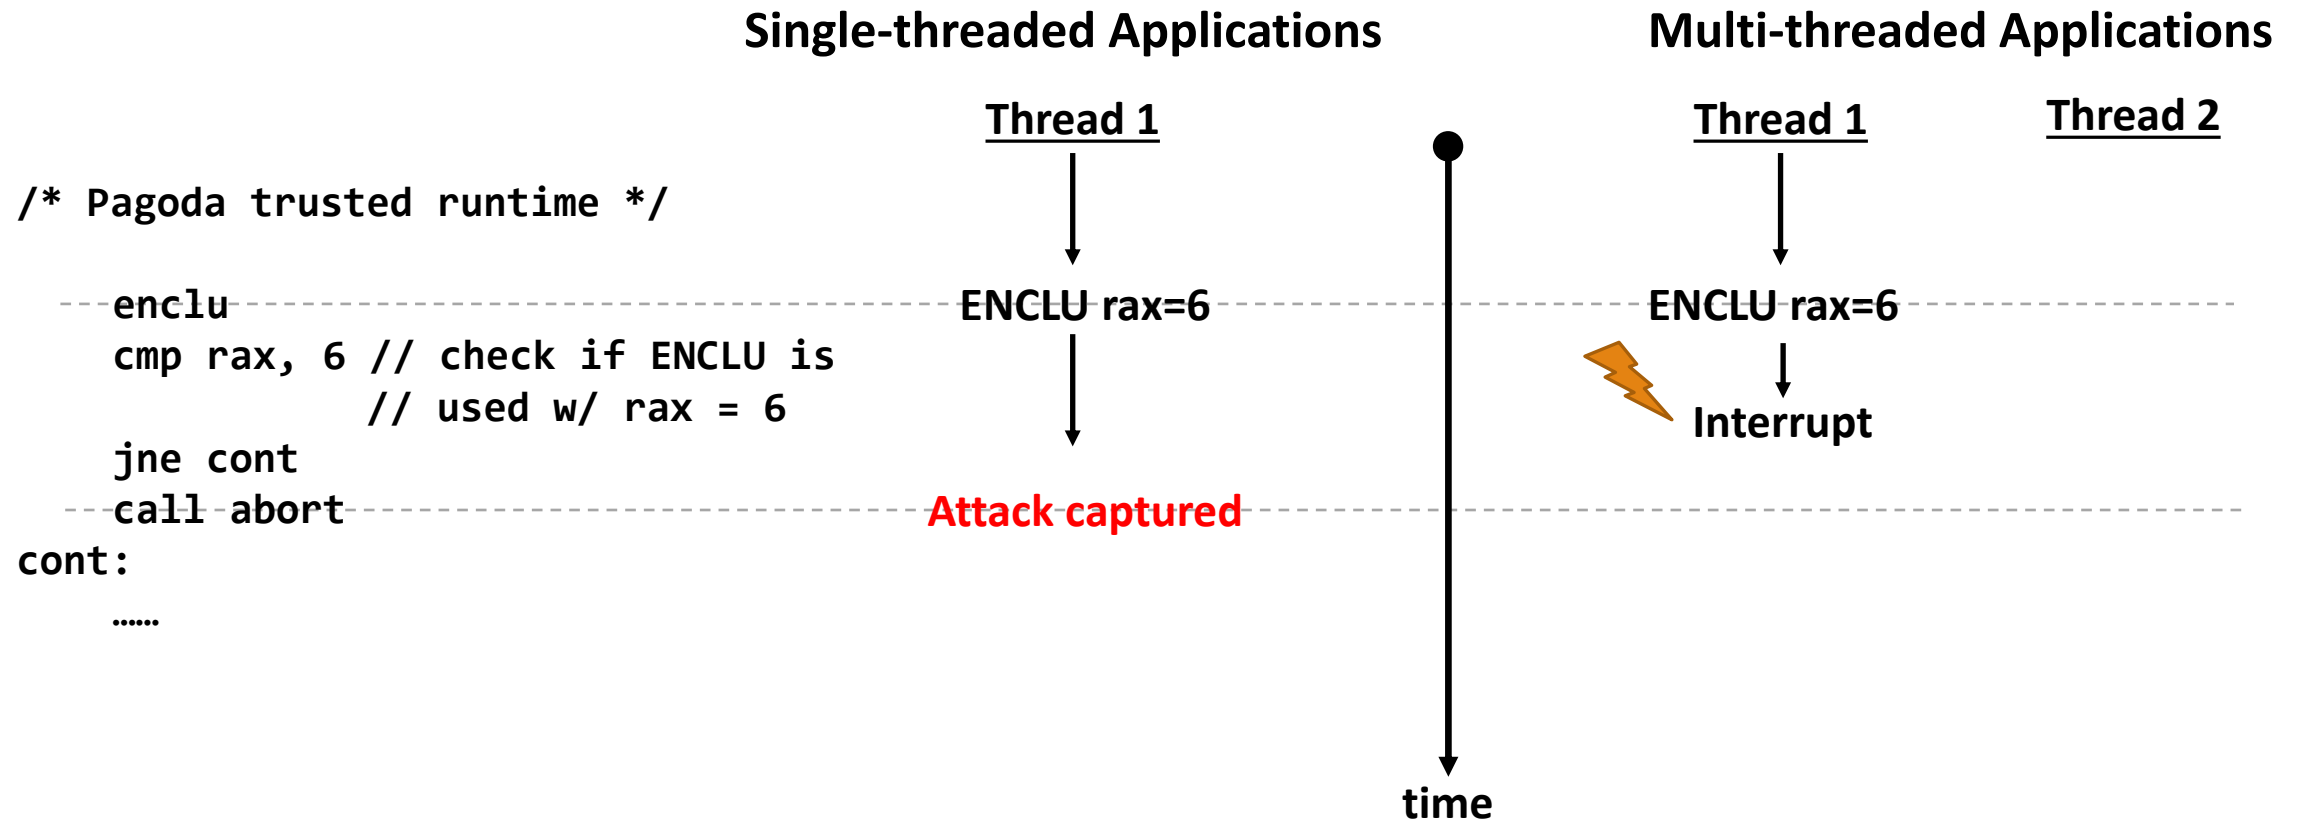
```
  .....
```

Preventing Relaxation of Execute-Only

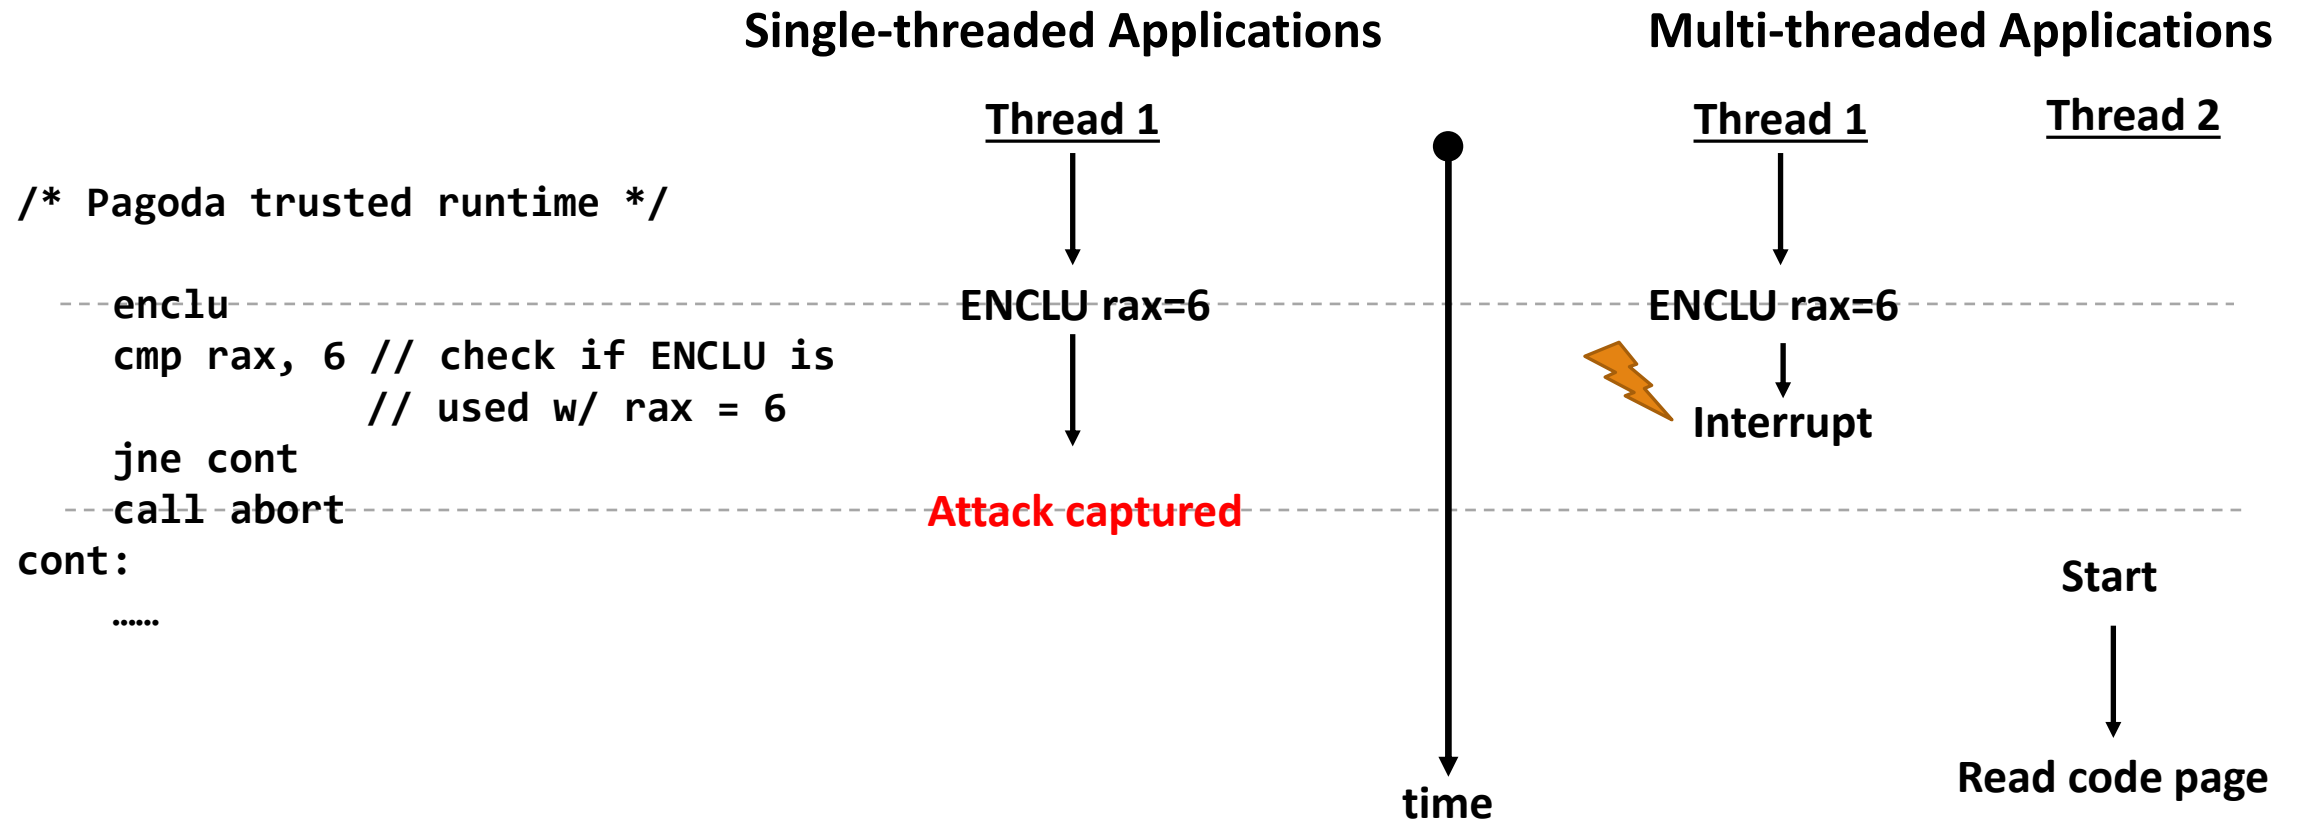
Single-threaded Applications



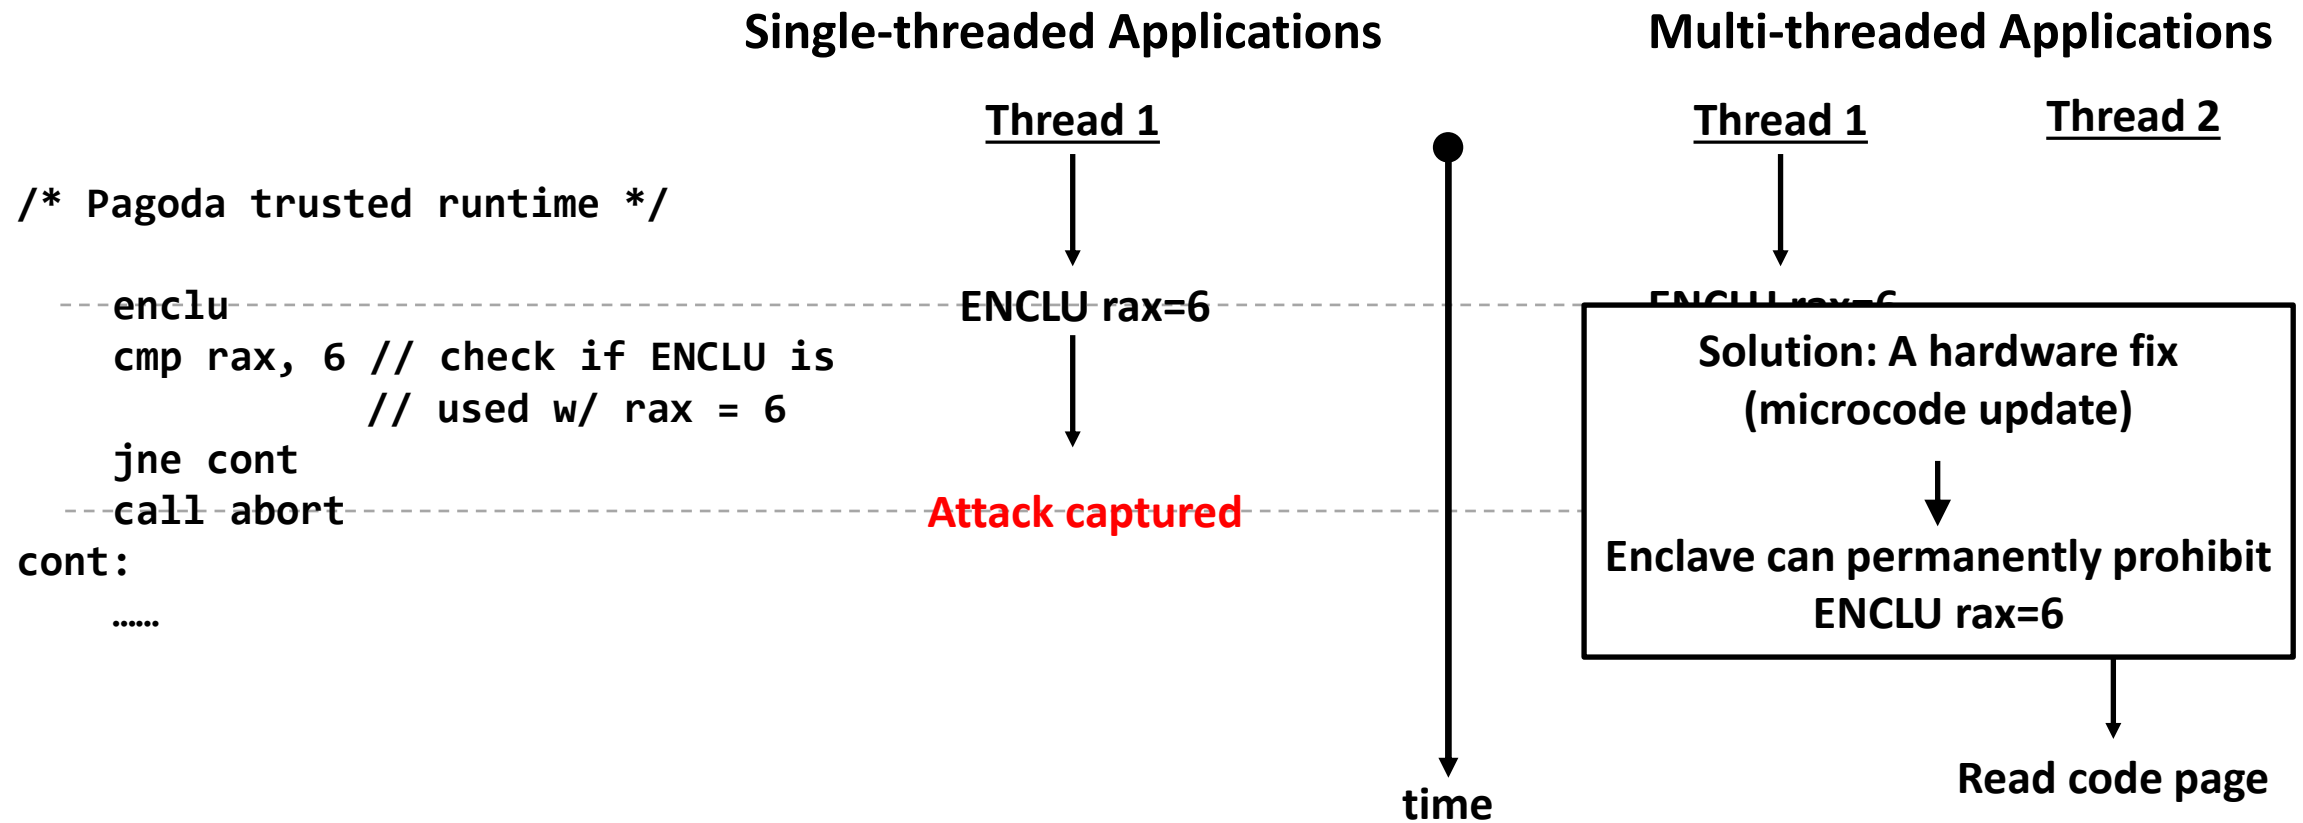
Preventing Relaxation of Execute-Only



Preventing Relaxation of Execute-Only



Preventing Relaxation of Execute-Only



Additional Features

- Fast cross-enclave-boundary calls
- Supporting different types of applications using dynamic linking/loading
- Multi-threading, signal-handling, etc.

Evaluation

Benchmark	Type	Performance Metric	Pagoda / Native Linux
SPEC CPU 2017	Standardized benchmark	Execution time	102.1%
Lighttpd	Server applications	Peak Throughput	28.9%
Memcached			51.1%
Quake	Desktop applications	Frame-per-second (FPS)	86.2%
Witchblast			97.0%

Conclusion

This work proposes Pagoda, the first practical and high-performance code privacy protection system based on SGX.

Q & A

Backup slides

Evaluation – SPEC CPU 2017

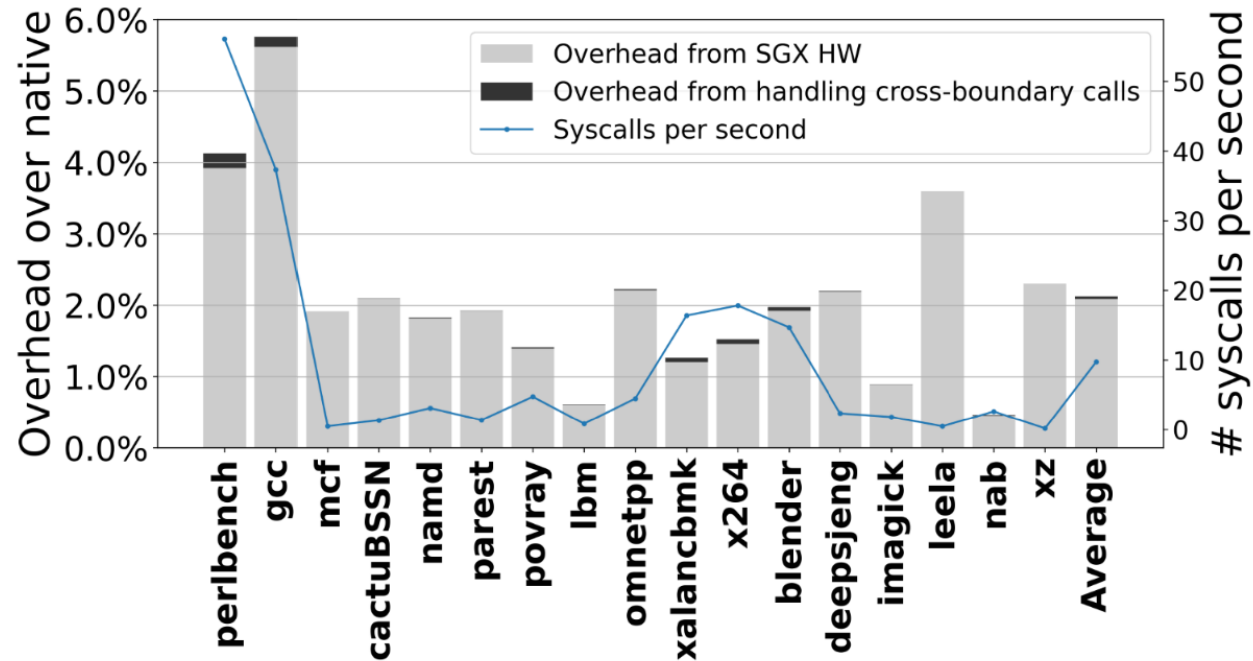


Fig. 5: The breakdown of Pagoda's performance overhead over native Linux execution for SPEC2017 benchmarks.

Evaluation – Server Applications

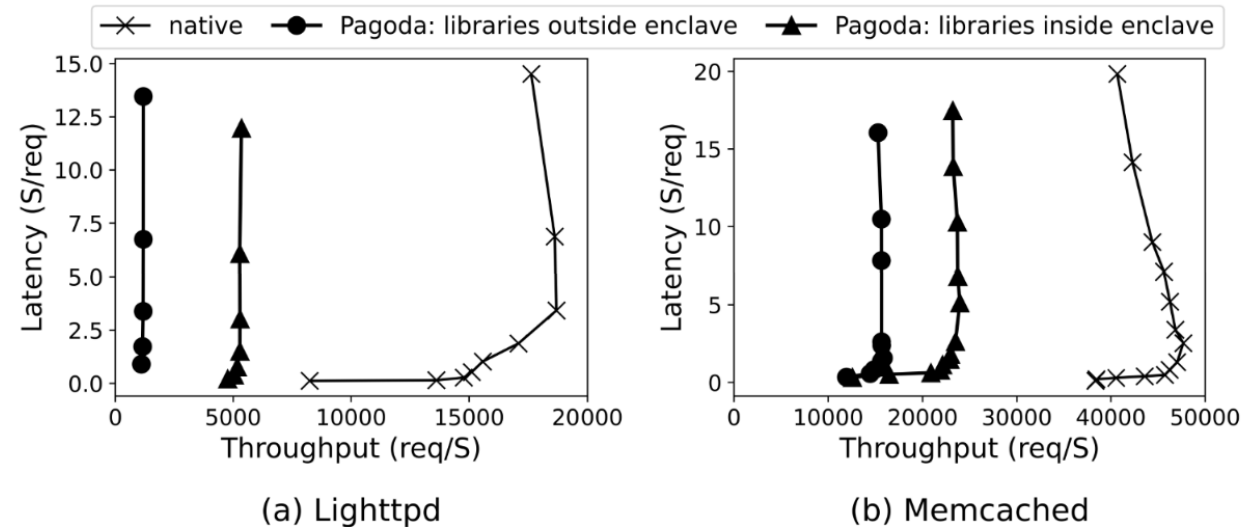


Fig. 6: Throughput vs. Latency of Lighttpd and Memcached. We run both applications with three configurations: bare-metal Linux, Pagoda with all shared libraries outside the enclave (treating all libraries as public), Pagoda with all libraries inside the enclave (treating all libraries as private).

Evaluation – Desktop (Gaming) Applications

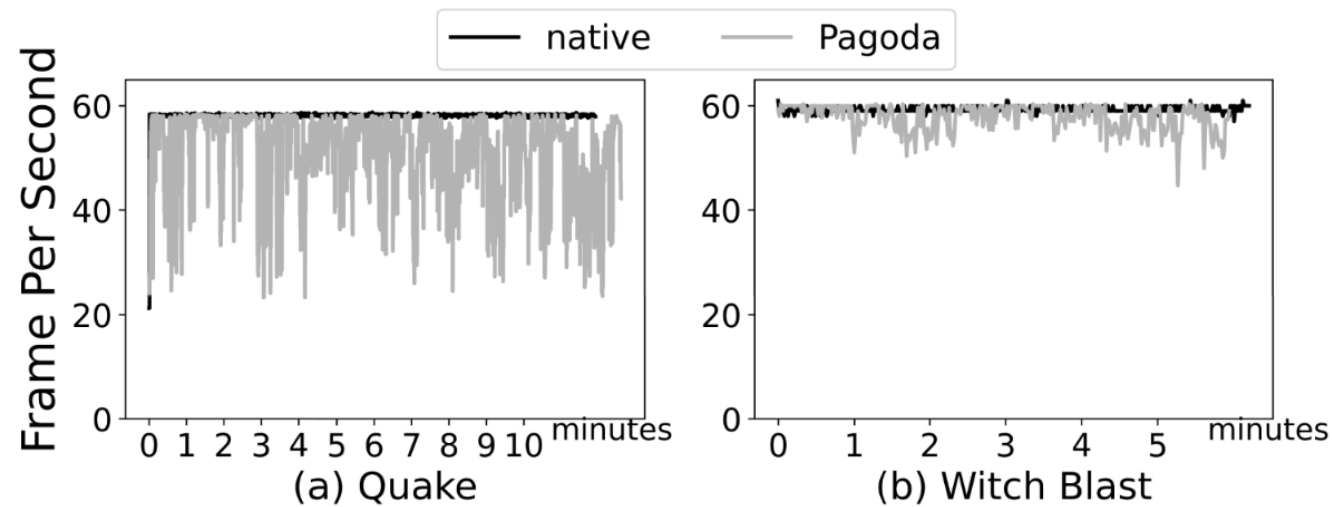
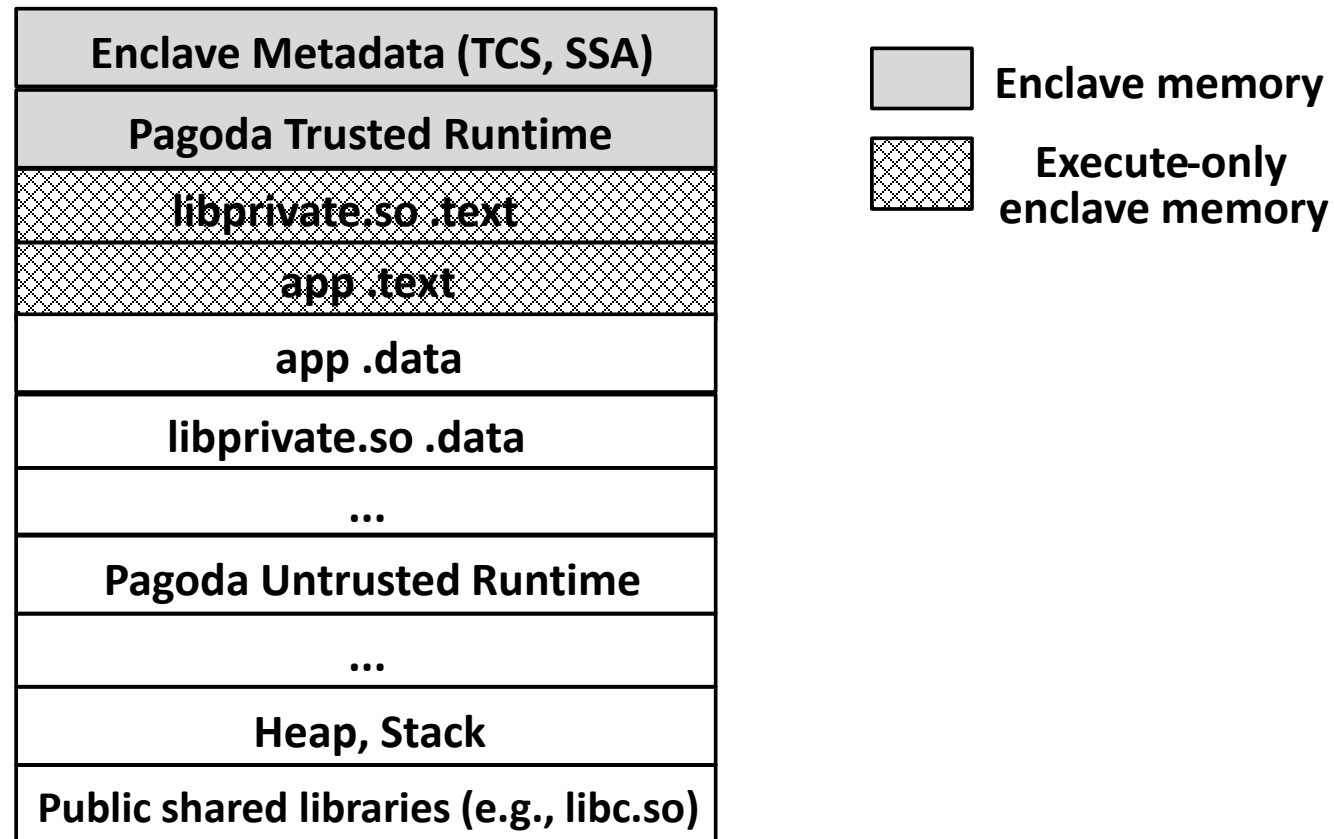


Fig. 7: Comparing the change of Frame-Per-Second over time between native Linux execution and Pagoda.

Pagoda memory layout

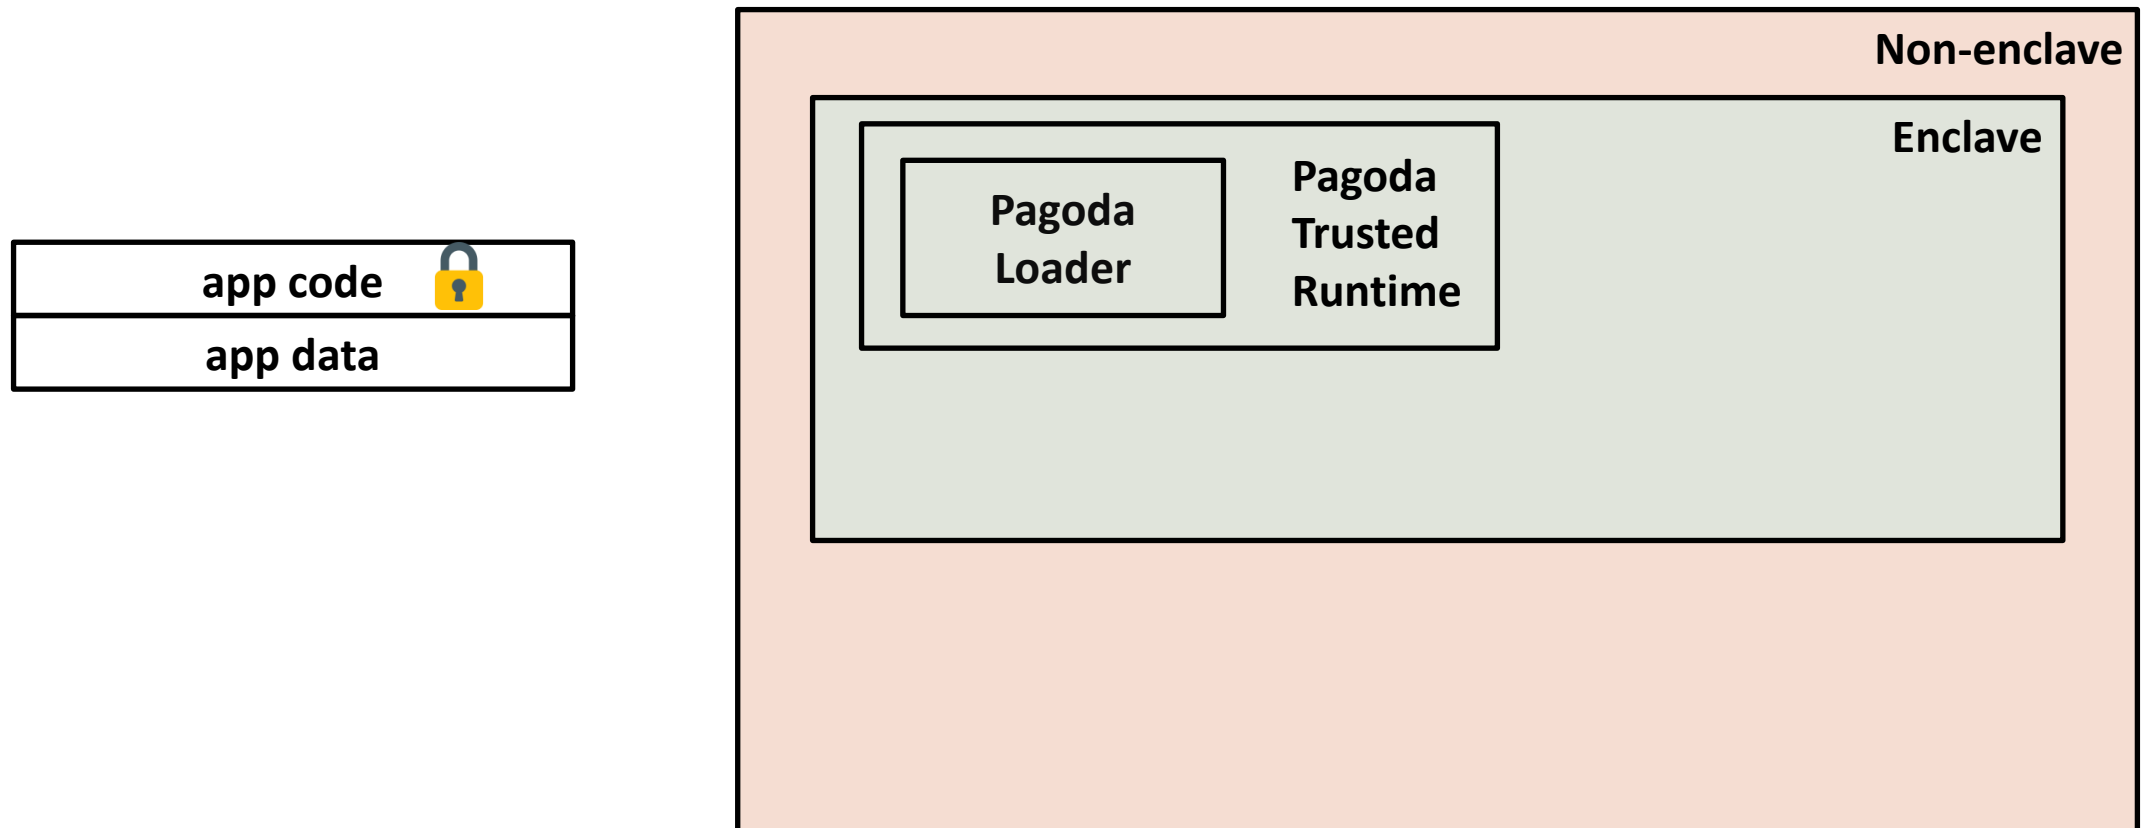


Direct vs. Indirect attack

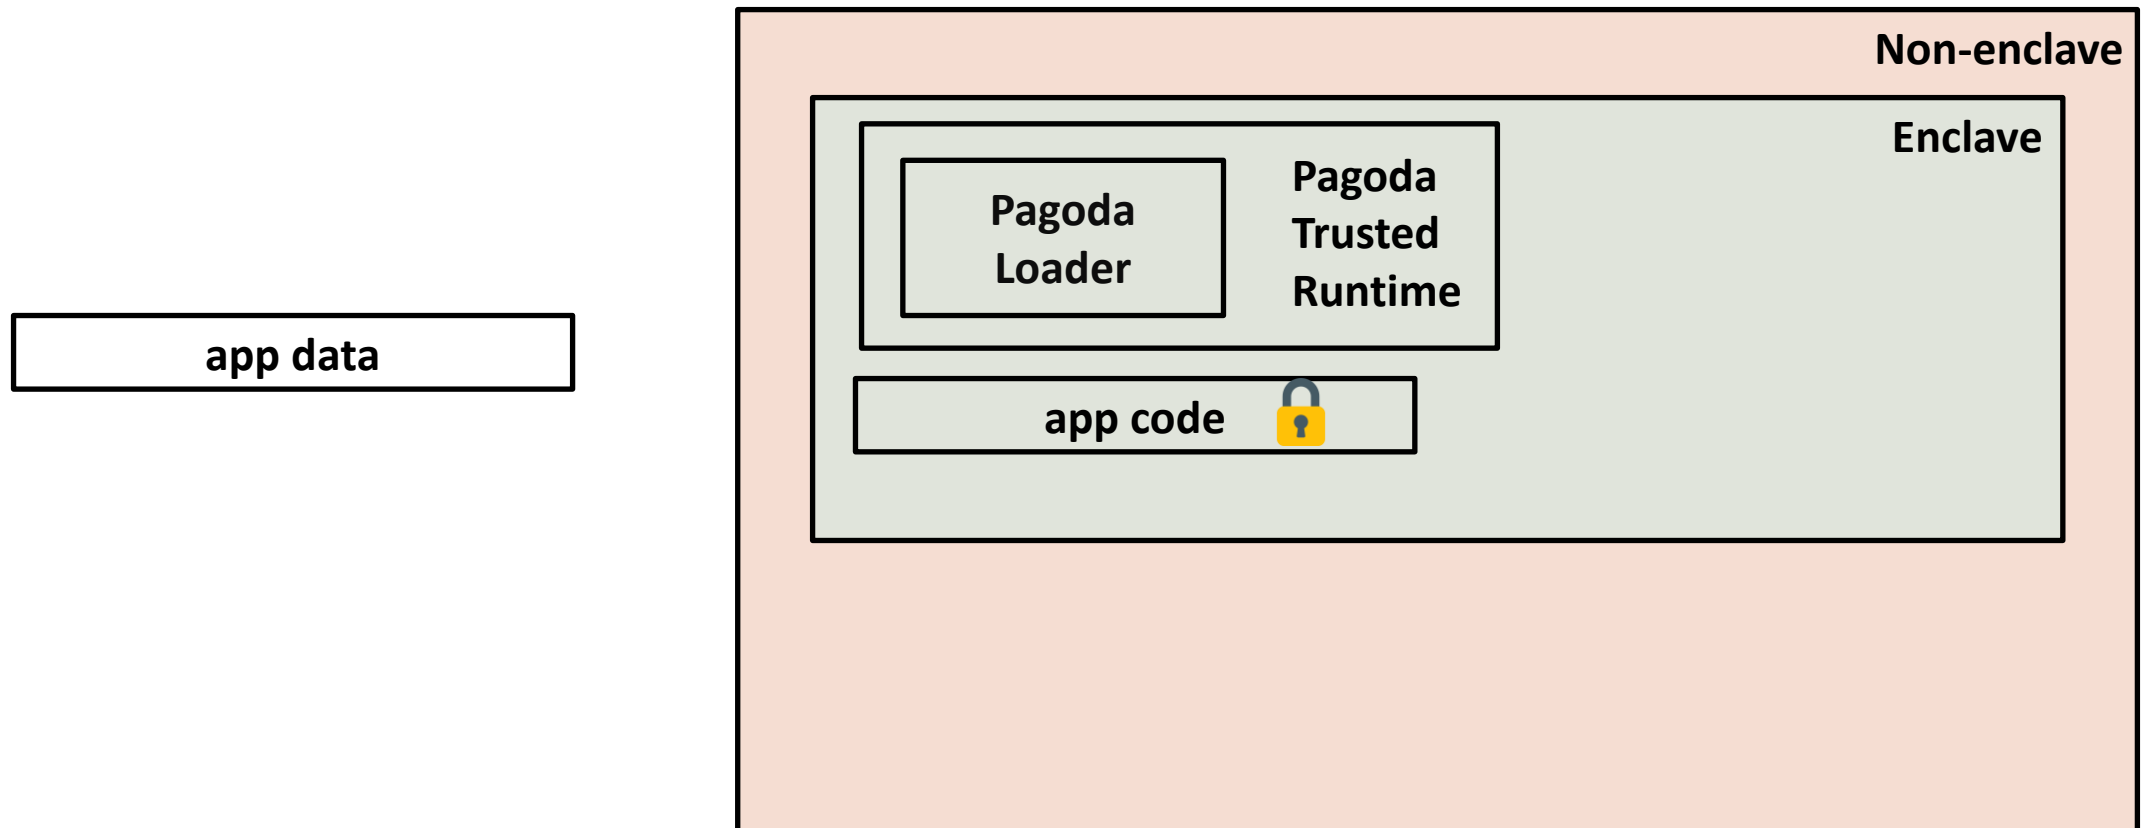
Direct attack starts w/ direct read of code. All attack discussion so far is about this type

Indirect attack starts w/ execute code & some form of observation of the system (data memory, uarch usage) This require huge effort, non-trivial, but possible (to at least reconstruct some information about the code), future work. See discussion section in the paper.

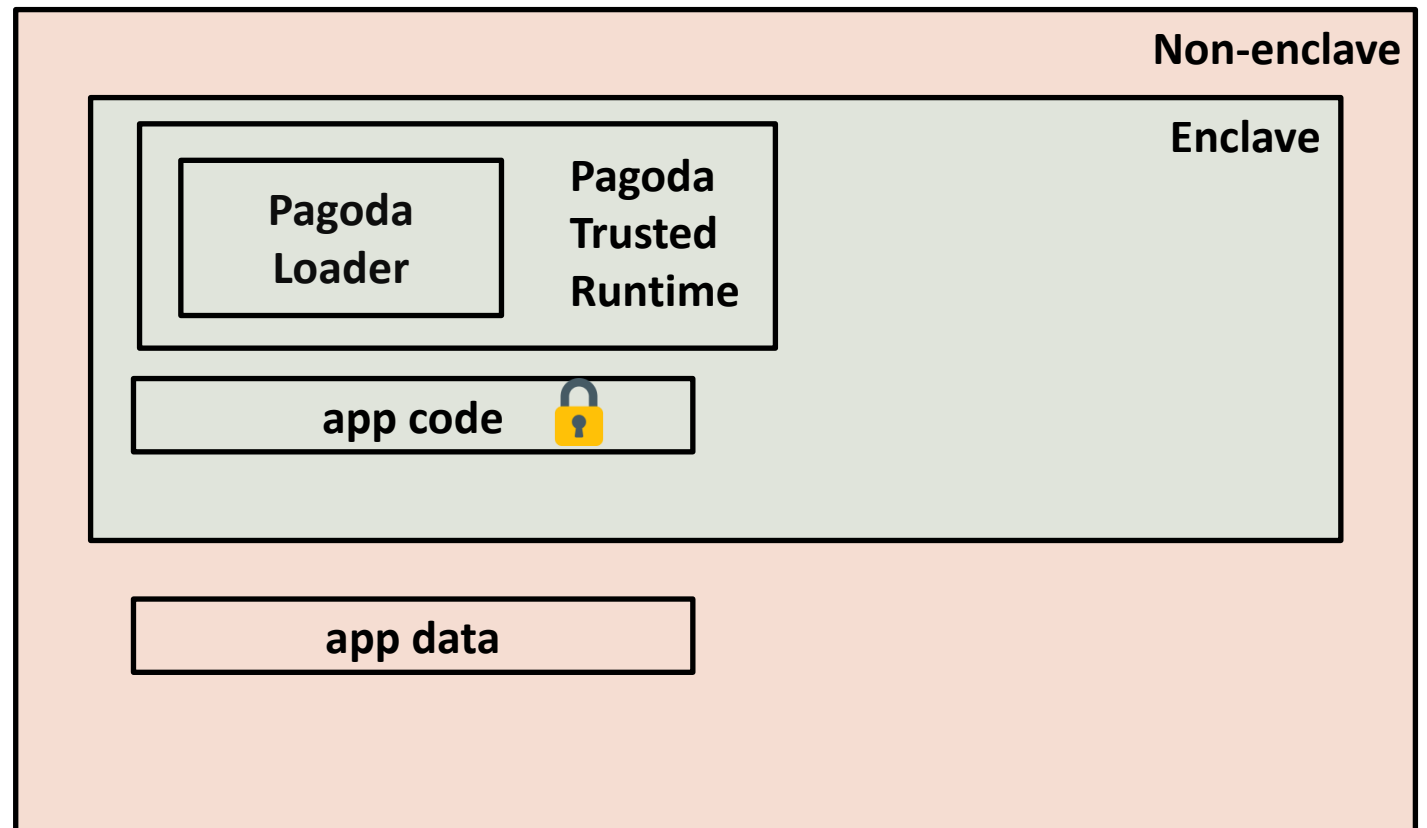
Deploying a Protected Software



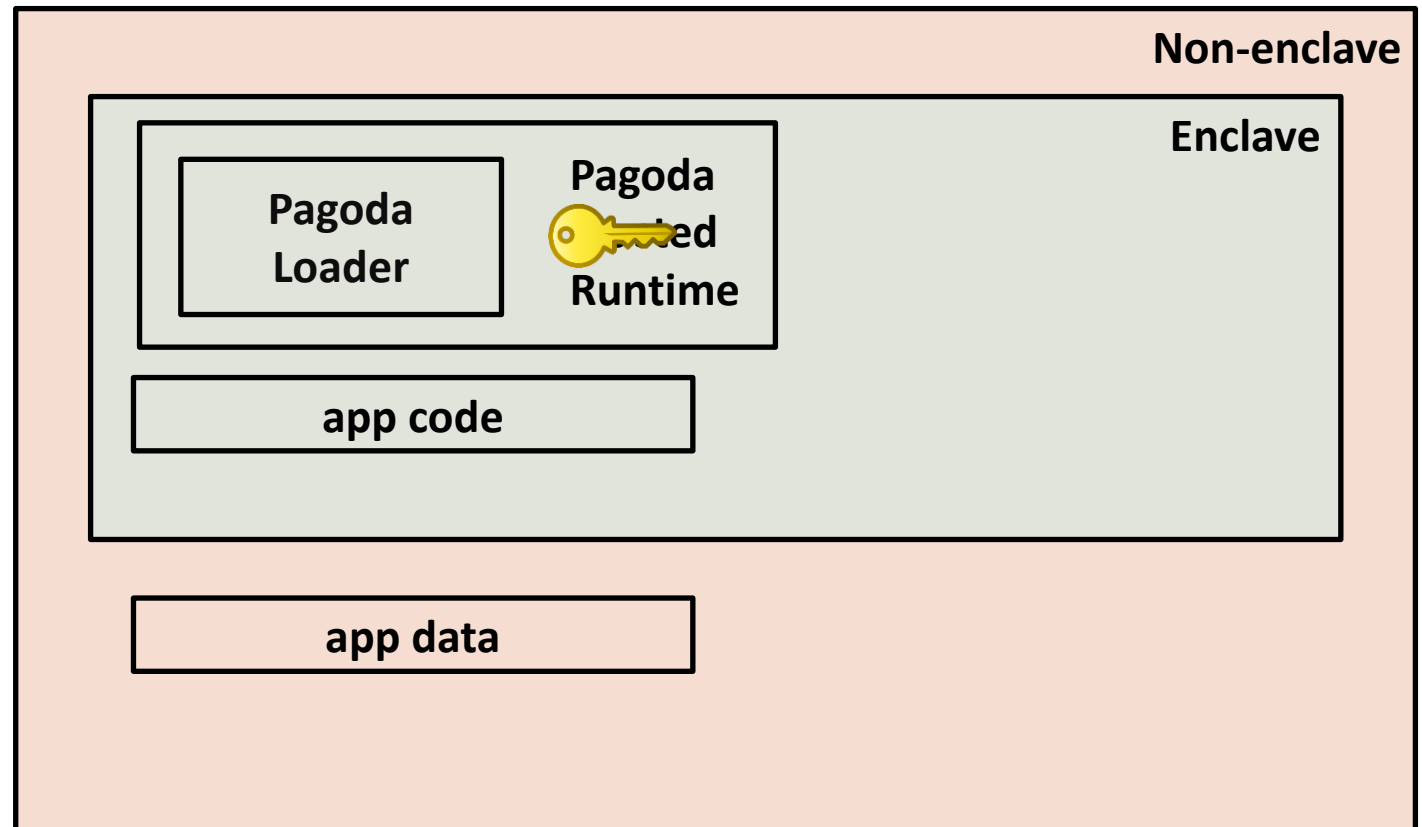

Deploying a Protected Software



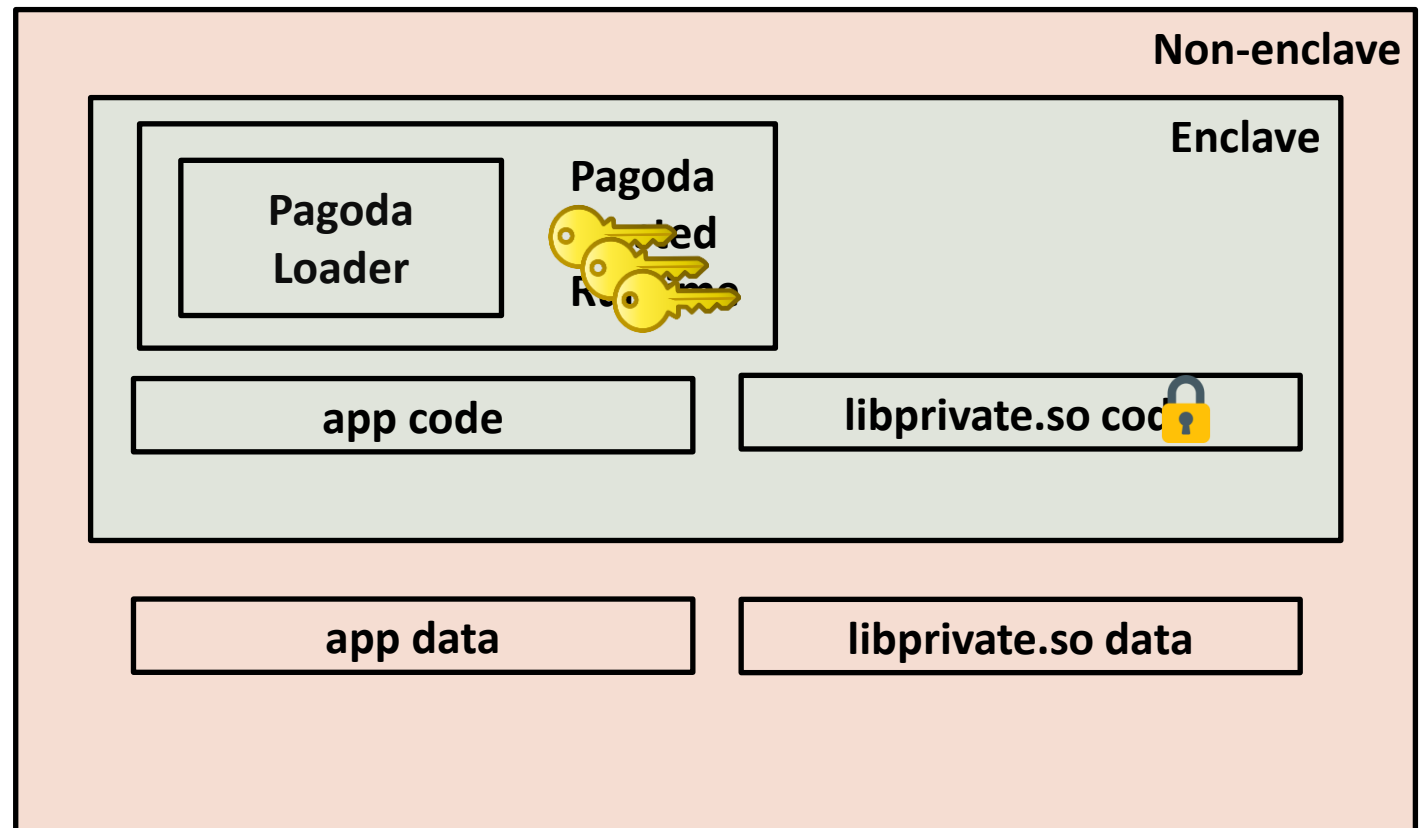

Deploying a Protected Software



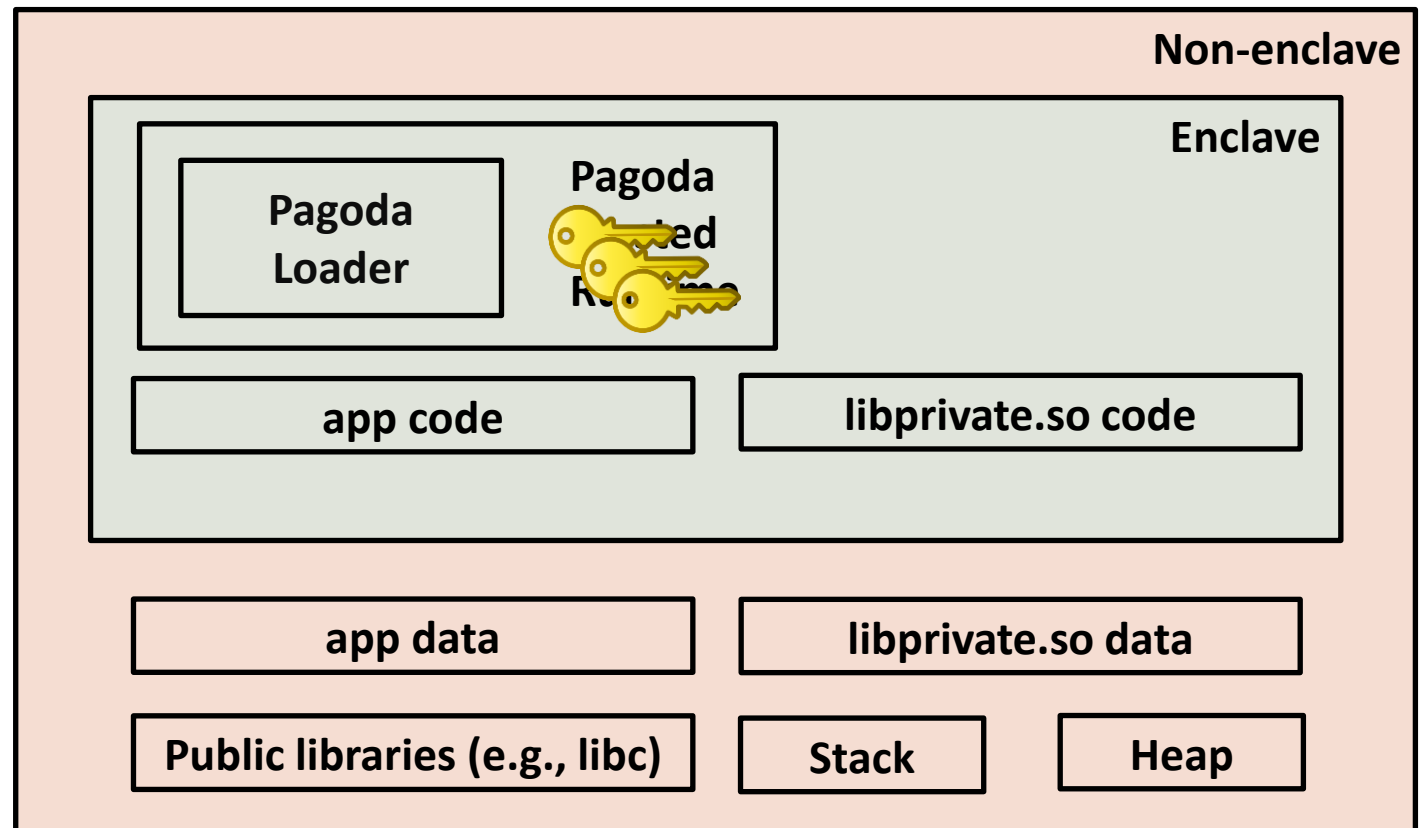

Deploying a Protected Software



Deploying a Protected Software



Deploying a Protected Software



Meeting notes

1. Add an overview slide: objective: lifetime code privacy. Make a graph depicts in transit, at rest, in use. Say it's the standard way to fix in transit/at rest (prior works) with remote attestation/encryption. We focus on in-use in Pagoda. In 10-min version, replace all the slides on this topic with this overview slide.
2. keep trusted runtime term, add explanation what it does and loader is part of it
3. 23->24: copy the figure, when saying $\frac{1}{2}$, point to the encl in the figure, make it clear that we eliminate encl in app binary, and regulate the encl in Pagoda runtime
4. cut eval slide details (NUC...)
5. server app graph: more requests along the curve
6. backup slide #1: detailed memory layout, explain what we do for making this layout
7. backup slide #2: discussion section, direct/indirect attack

Direct attack starts w/ direct read of code. All attack discussion so far is about this type

Indirect attack starts w/ execute code & some form of observation of the system (data memory, uarch usage) This require huge effort, non-trivial, but possible (to at least reconstruct some information about the code), future work. See discussion section in the paper.

Question: what's the contribution of the paper:

- 1) insight: for code privacy, just put code in enclave -> great compatibility
- 2) for insight, we need to address some problems: a) functionality. See backup slide, and the communication stuff b) security: a practical assumption (vulnerable code), we tackle all vulnerability problems to use of XOM. Also boil down all ways to break XOM into preventing the use of EMDPE. And study how to solve this problem for single/multi-thread app

Chris's notes

- minor: since xinyang was at MSR when he did pagoda, not sure we want the databricks logo
- slide 4: say for both settings, the software stack is untrusted. you said "platform", which could also mean HW.
- slide 9: I don't see the most important thing on the slide "enclave code is in TCB"
- slide 12: it's not just moving data/in out. that doesn't really cover the syscall benefits.
- slide 12: it's not as straightforward as XOM. you still need to ensure XOM is maintained throughout app lifetime. we should allude to this being a problem we will have to solve, but not say (yet) how we will solve it.
- slide 18: suggested organization: say we need to protect code throughout its lifetime = in use + in transit + at rest. here is how we deal with each: ... (and make it clear that code in transit/at rest isn't our contribution)
- slide 25: the red X is a bit vague. maybe spell out "attack detected"?
- slide 31: you might want to say something about syscalls, because you don't mention earlier how they are performed.
- slide 32: why are the graphs bending backwards on themselves!?
- slide 33: it's a bit counter-intuitive that the overcall to syscalls would create such a large dip.

Weidong, Xinyang
