



Speculative Taint Tracking (STT): A Comprehensive Protection for Speculatively Accessed Data

JIYONG YU, MENGJIA YAN, ARTEM KHYZHA*,

ADAM MORRISON*, JOSEP TORRELLAS, CHRISTOPHER W. FLETCHER

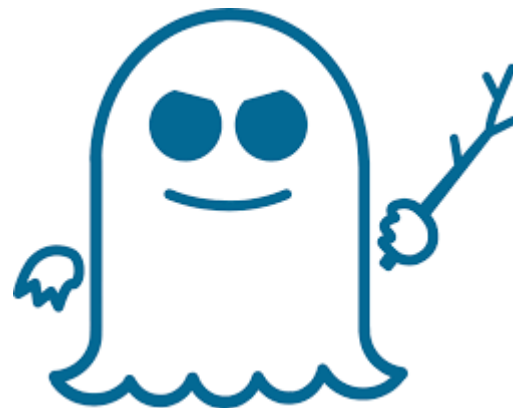
UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

*TEL AVIV UNIVERSITY

Processors are Insecure



Processors are Insecure



Speculative Execution Attacks

```
// Spectre Variant 1
```

```
if (addr < N) { // speculation
```

```
    // access instruction
```

```
    spec_val = load [addr];
```

```
    // covert channel
```

```
    load [spec_val];
```

```
}
```

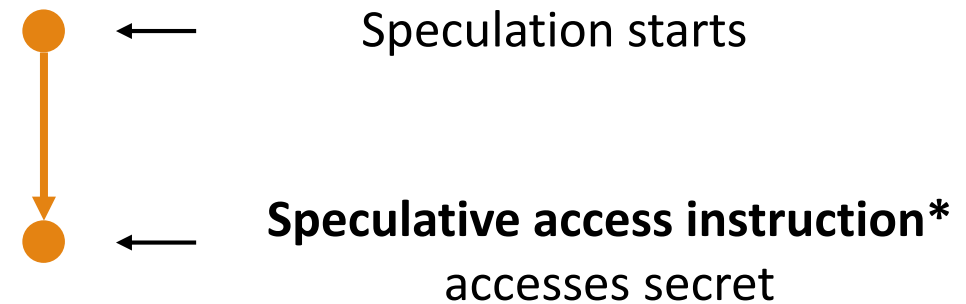


Speculation starts

Speculative Execution Attacks

```
// Spectre Variant 1

if (addr < N) {    // speculation
    // access instruction
    spec_val = load [addr];
    // covert channel
    load [spec_val];
}
```



*: Kiriansky, Vladimir, et al. "DAWG: A defense against cache timing attacks in speculative execution processors." *MICRO-51*, 2018.

Speculative Execution Attacks

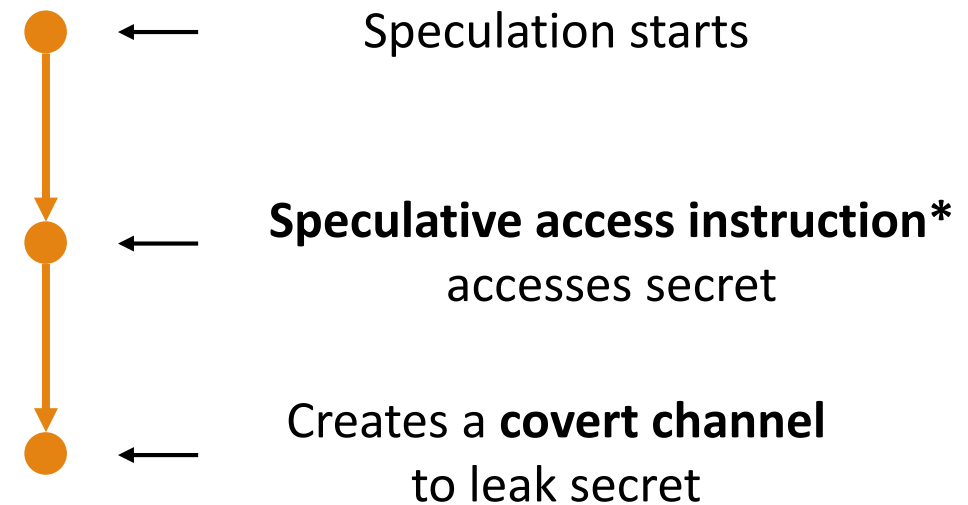
```
// Spectre Variant 1

if (addr < N) {    // speculation

    // access instruction
    spec_val = load [addr];

    // covert channel
    load [spec_val];

}
```



*: Kiriansky, Vladimir, et al. "DAWG: A defense against cache timing attacks in speculative execution processors." *MICRO-51*, 2018.

Speculative Execution Attacks

```
// Spectre Variant 1
```

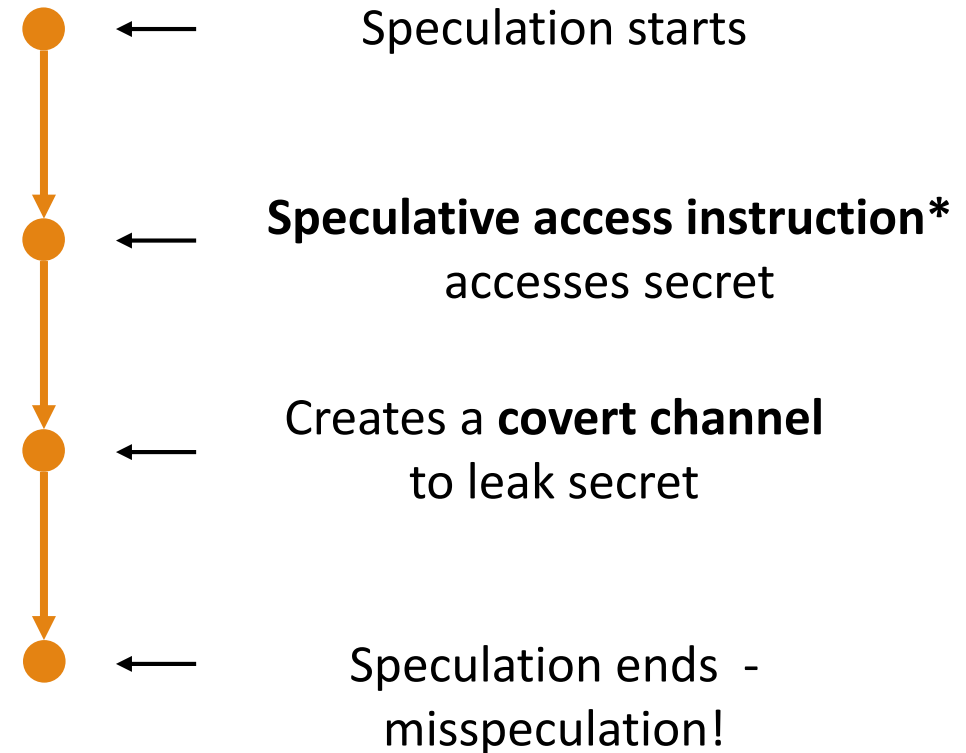
```
if (addr < N) { // speculation
```

addr =
N+1

```
// access instruction  
spec_val = load [addr];
```

```
// covert channel  
load [spec_val];
```

```
}
```



*: Kiriansky, Vladimir, et al. "DAWG: A defense against cache timing attacks in speculative execution processors." *MICRO-51*, 2018.

Speculative Execution Attacks

```
// Spectre Variant 1
```

```
if (addr < N) { // speculation
```

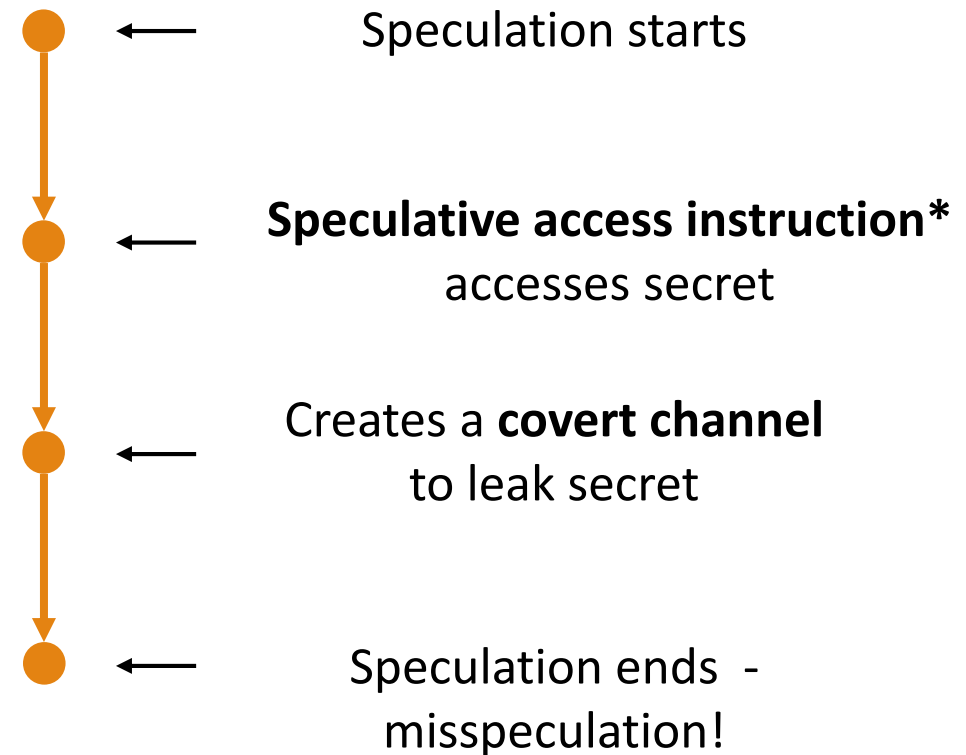
```
    // access instruction
```

```
    spec_val = load [addr];
```

```
    // covert channel
```

```
    load [spec_val];
```

```
}
```



*: Kiriansky, Vladimir, et al. "DAWG: A defense against cache timing attacks in speculative execution processors." *MICRO-51*, 2018.

Main Insight of STT

Main Insight of STT

“Sufficient for security: prevent secrets from reaching covert channels”

Main Insight of STT

“Sufficient for security: prevent secrets from reaching covert channels”

```

if (addr < N) {

    // access instruction
    spec_val = load [addr];

    // simple arithmetic
    spec_val = spec_val + 4;

    // cache/mem covert channel
    load [spec_val];

}

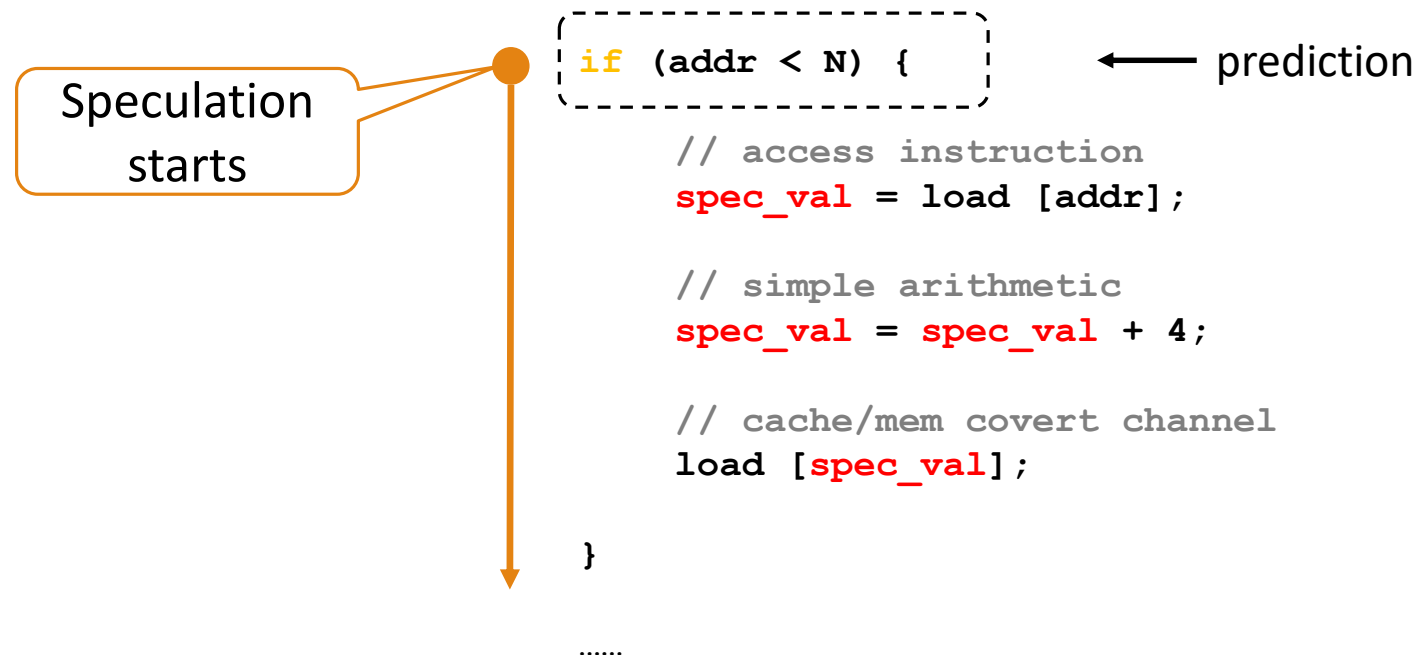
.....

```

Creates a covert channel?	Input operand is a secret?	Requires protection?

Main Insight of STT

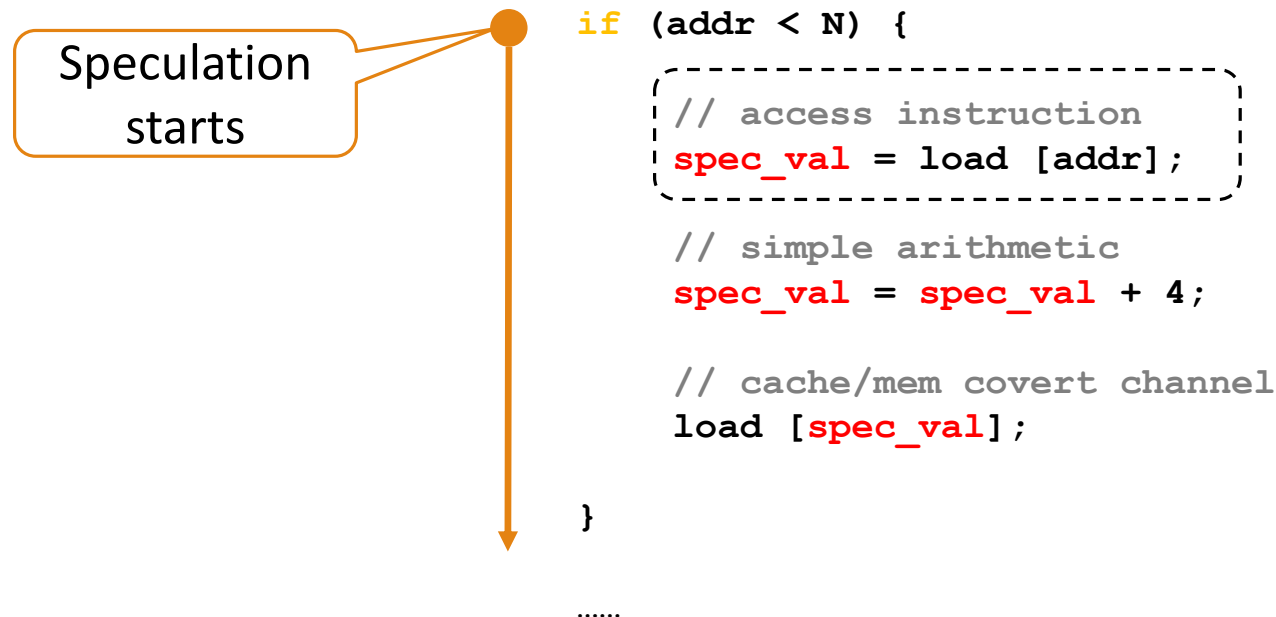
“Sufficient for security: prevent secrets from reaching covert channels”



Creates a covert channel?	Input operand is a secret?	Requires protection?

Main Insight of STT

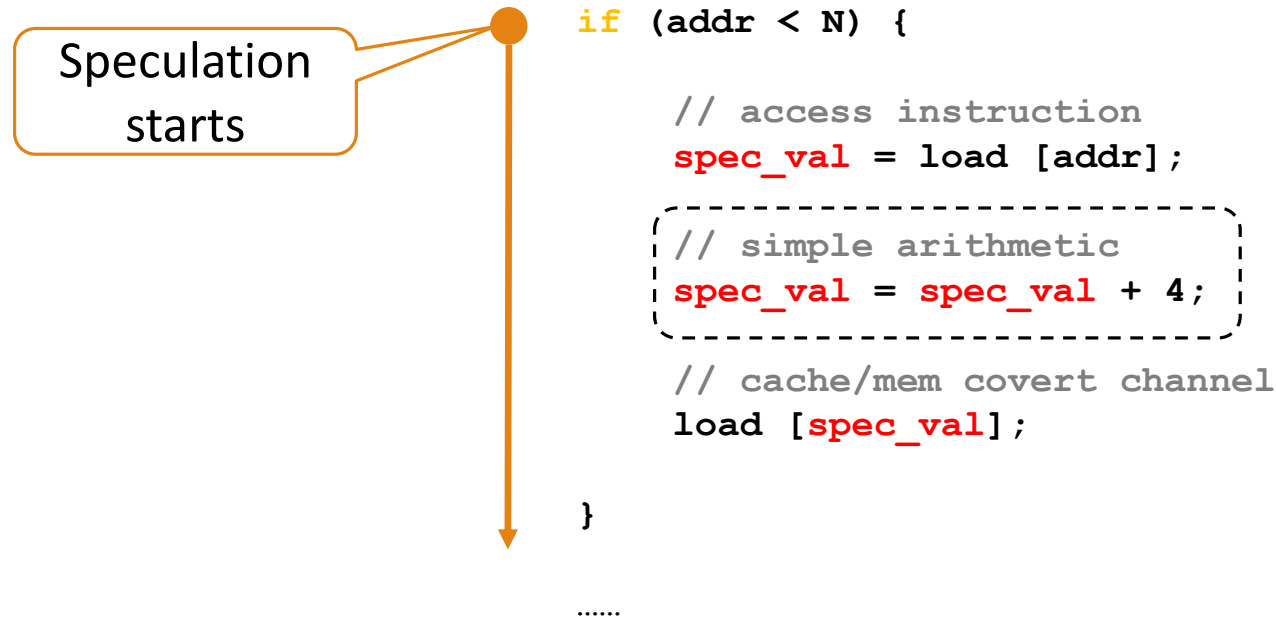
“Sufficient for security: prevent secrets from reaching covert channels”



Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No

Main Insight of STT

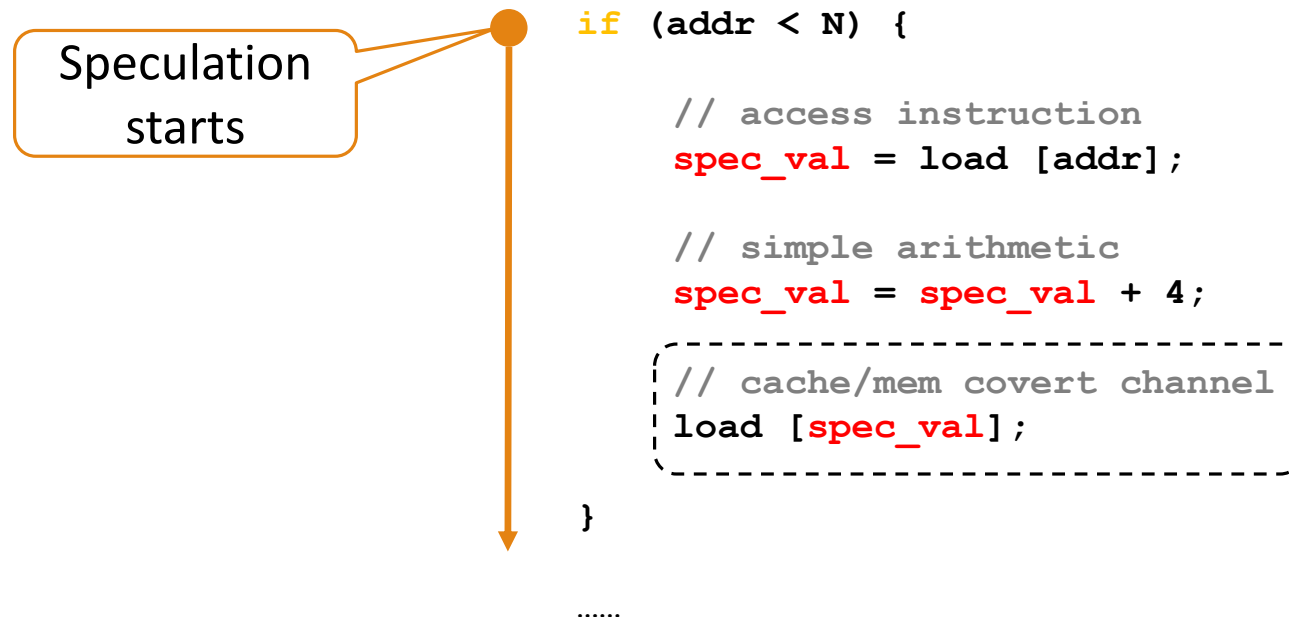
“Sufficient for security: prevent secrets from reaching covert channels”



Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No
No	Yes	No

Main Insight of STT

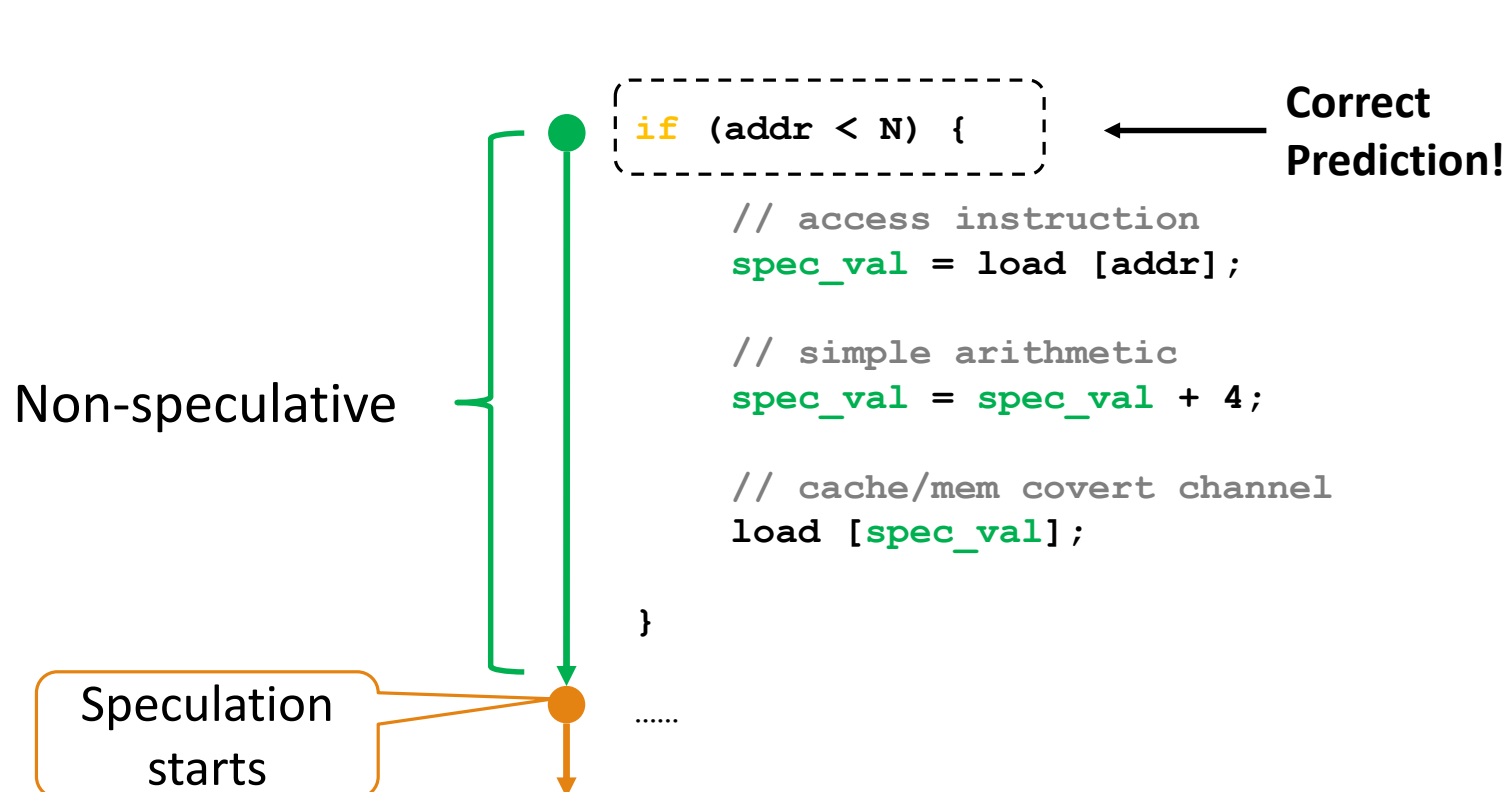
“Sufficient for security: prevent secrets from reaching covert channels”



Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No
No	Yes	No
Yes	Yes	Yes

Main Insight of STT

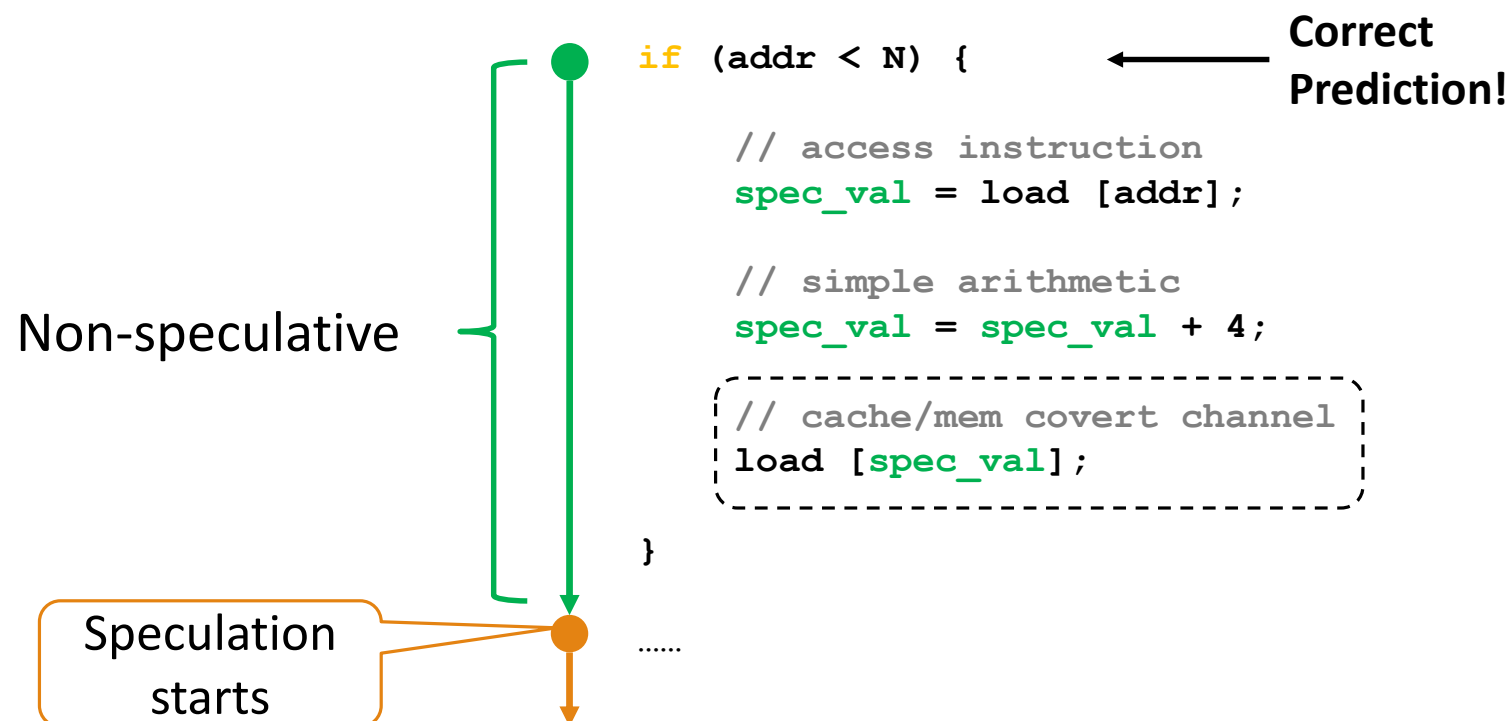
“Sufficient for security: prevent secrets from reaching covert channels”



Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No
No	No	No
Yes	No	No

Main Insight of STT

“Sufficient for security: prevent secrets from reaching covert channels”

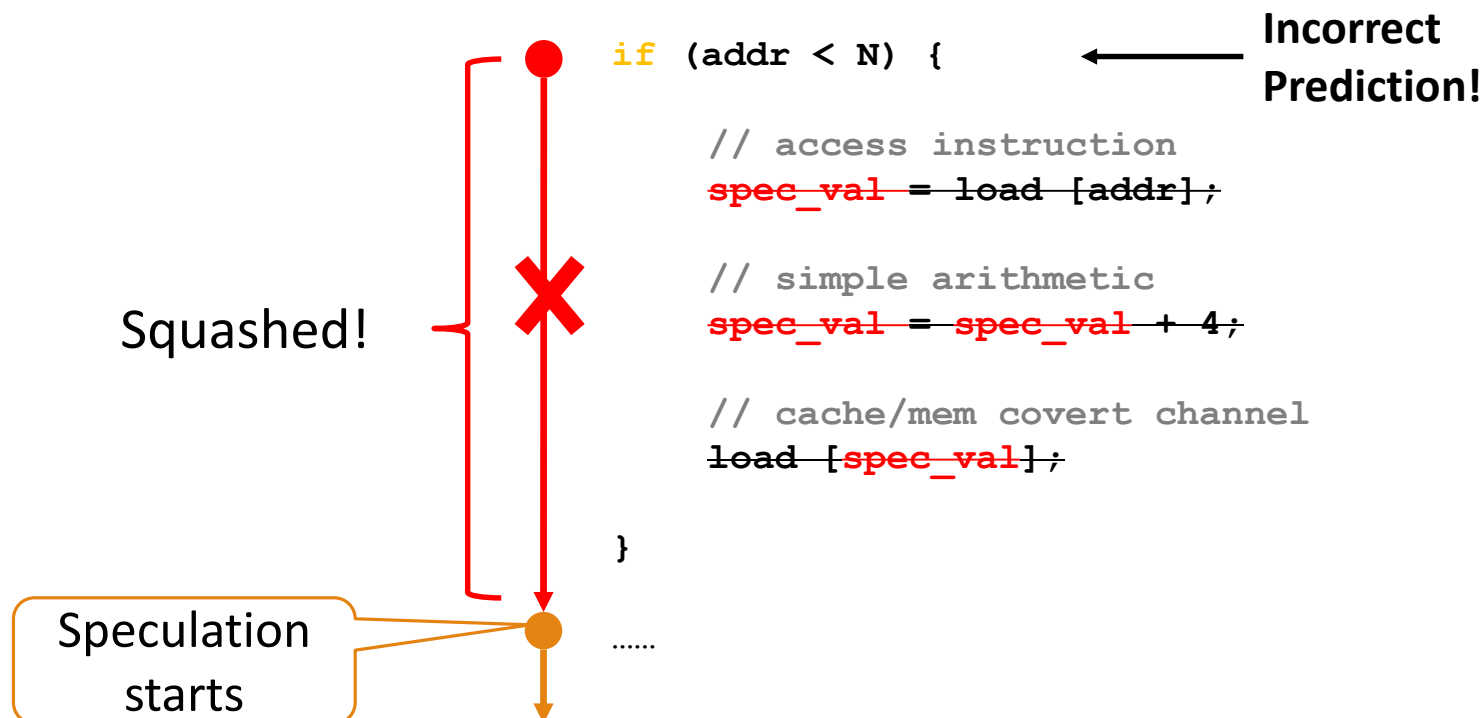


Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No
No	No	No
Yes	No	No


Main Insight of STT

“Sufficient for security: prevent secrets from reaching covert channels”


Creates a covert channel?	Input operand is a secret?	Requires protection?
Yes	No	No
No	Yes	No
Yes	Yes	Yes



Speculative Taint Tracking

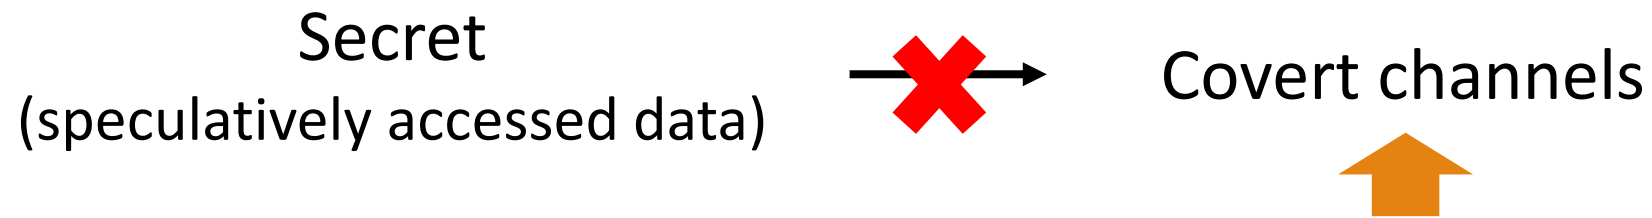
Secret
(speculatively accessed data)  Covert channels

Speculative Taint Tracking

Secret
(speculatively accessed data)  Covert channels



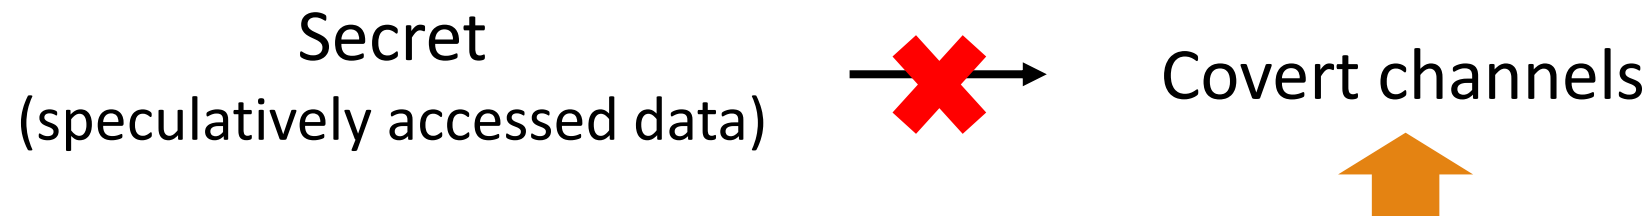
Speculative Taint Tracking



What are the covert channels?



Speculative Taint Tracking



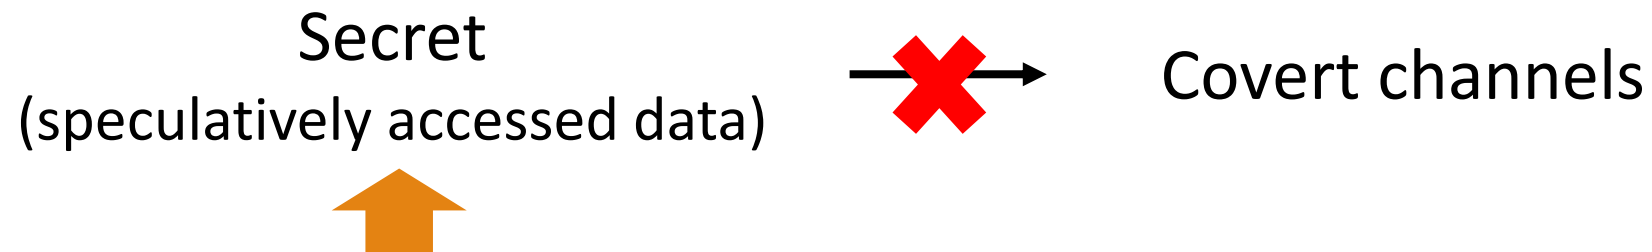
What are the covert channels?



**A new classification to understand
covert channels in speculative
machines**



Speculative Taint Tracking



What are the covert channels?

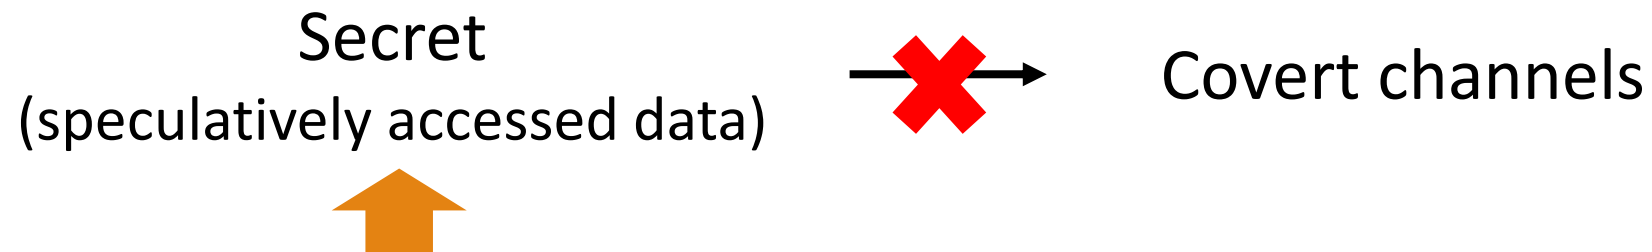


**A new classification to understand
covert channels in speculative
machines**



How to identify all the secrets?

Speculative Taint Tracking



What are the covert channels?



A new classification to understand covert channels in speculative machines

How to identify all the secrets?

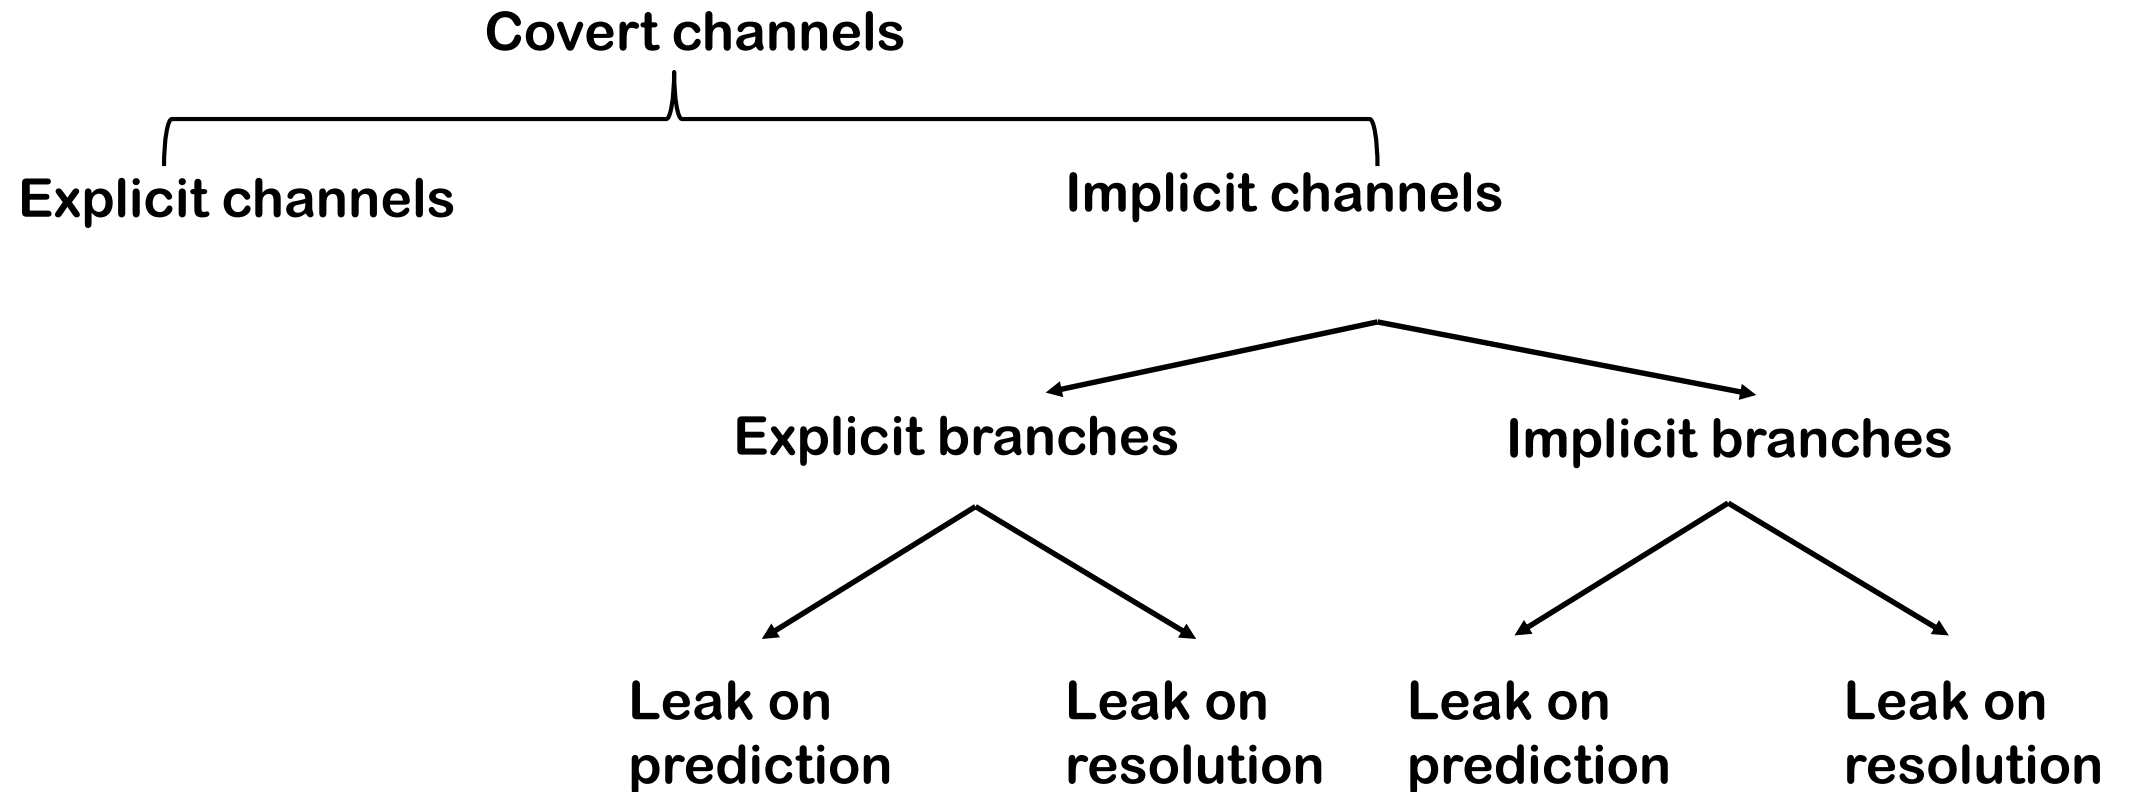


A new taint/untaint mechanism to track secrets in hardware

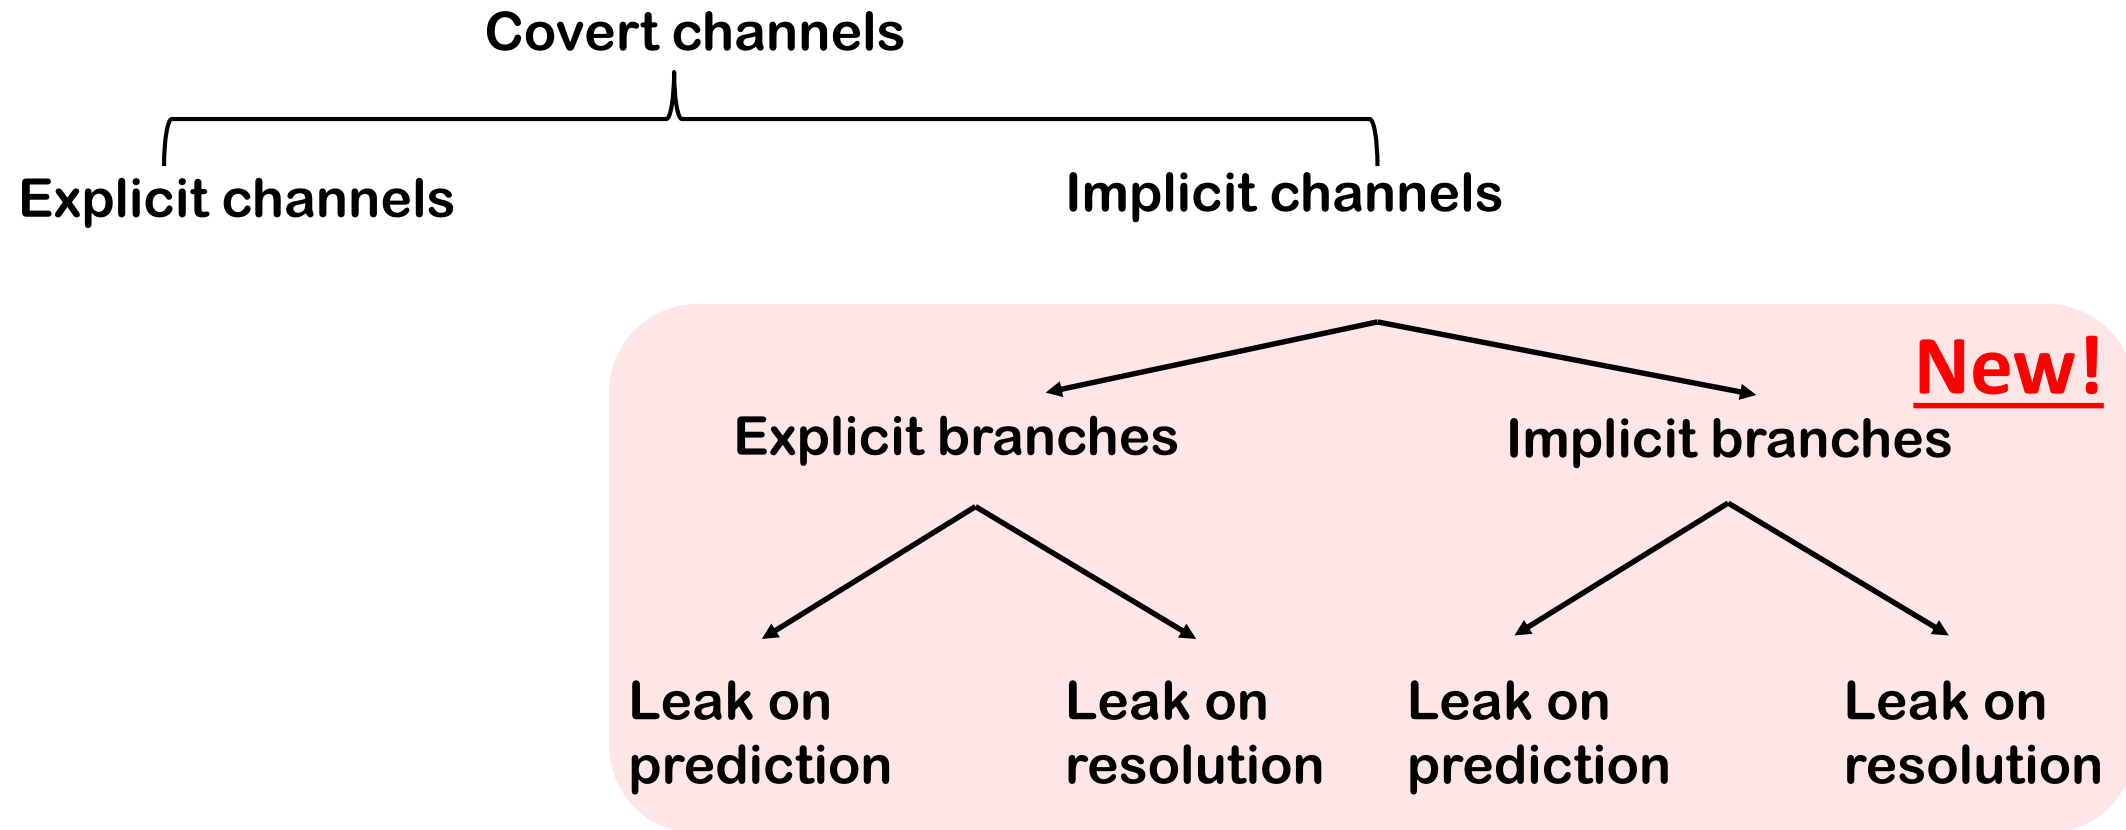


A Classification of Covert Channels in HW

Classification of Covert Channels



Classification of Covert Channels



Classification of Covert Channels

Covert channels



Explicit channels:

Secret inputs are directly leaked
by *operand-dependent*
hardware resource usage

```
load [secret];
```

Classification of Covert Channels

Covert channels



Explicit channels:

Secret inputs are directly leaked
by *operand-dependent*
hardware resource usage

Examples:

- memory loads
- data-dependent arithmetic

Classification of Covert Channels

Covert channels

Explicit channels:

Secret inputs are directly leaked
by *operand-dependent*
hardware resource usage

Examples:

- memory loads
- data-dependent arithmetic

Implicit channels:

Secret inputs are indirectly
leaked by *how (or that) one or*
several instructions execute

```
secret = load [addr];  
if (secret == 1)  
    load [0x00];
```

Classification of Covert Channels

Covert channels

```
graph TD; A[Covert channels] --- B[Explicit channels]; A --- C[Implicit channels];
```

Explicit channels:

Secret inputs are directly leaked
by *operand-dependent*
hardware resource usage

Examples:

- memory loads
- data-dependent arithmetic

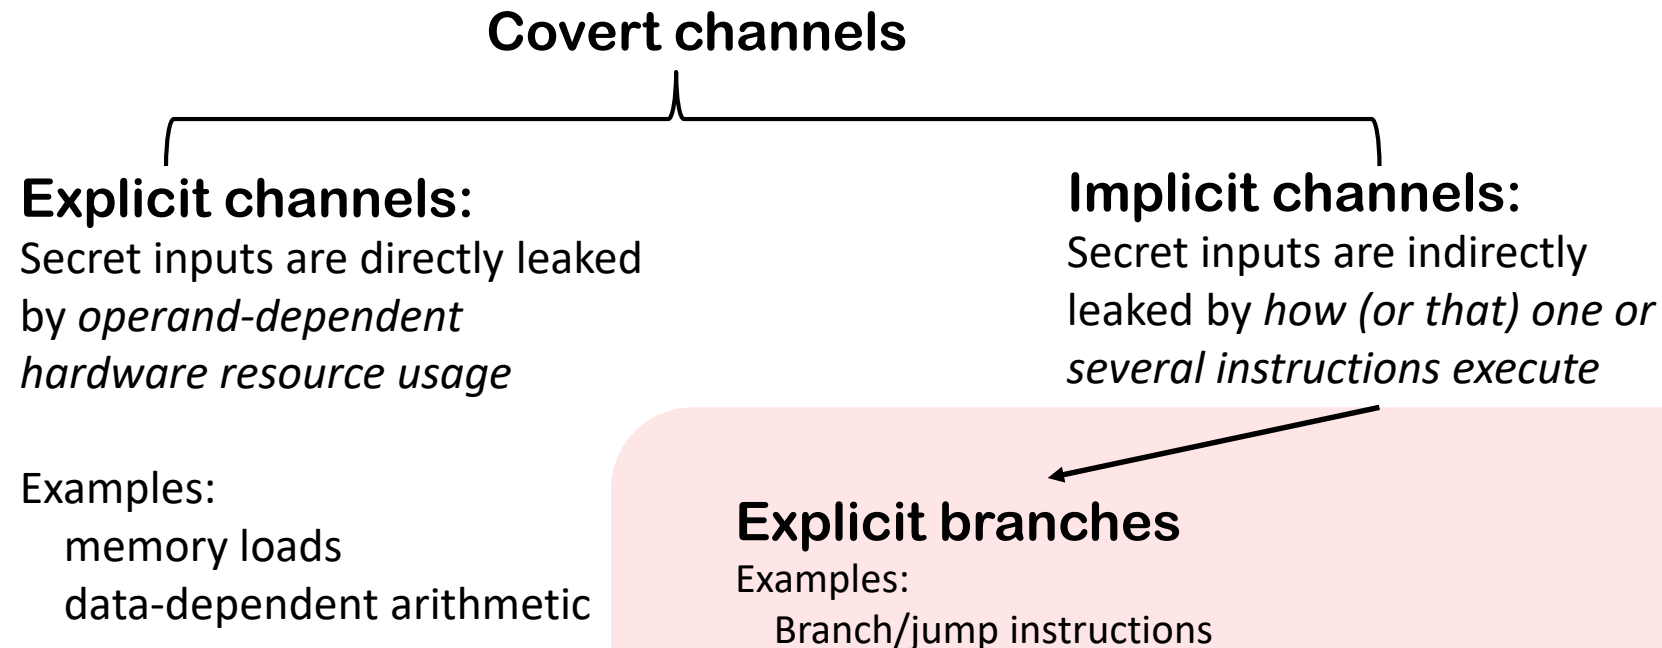
Implicit channels:

Secret inputs are indirectly
leaked by *how (or that) one or*
several instructions execute

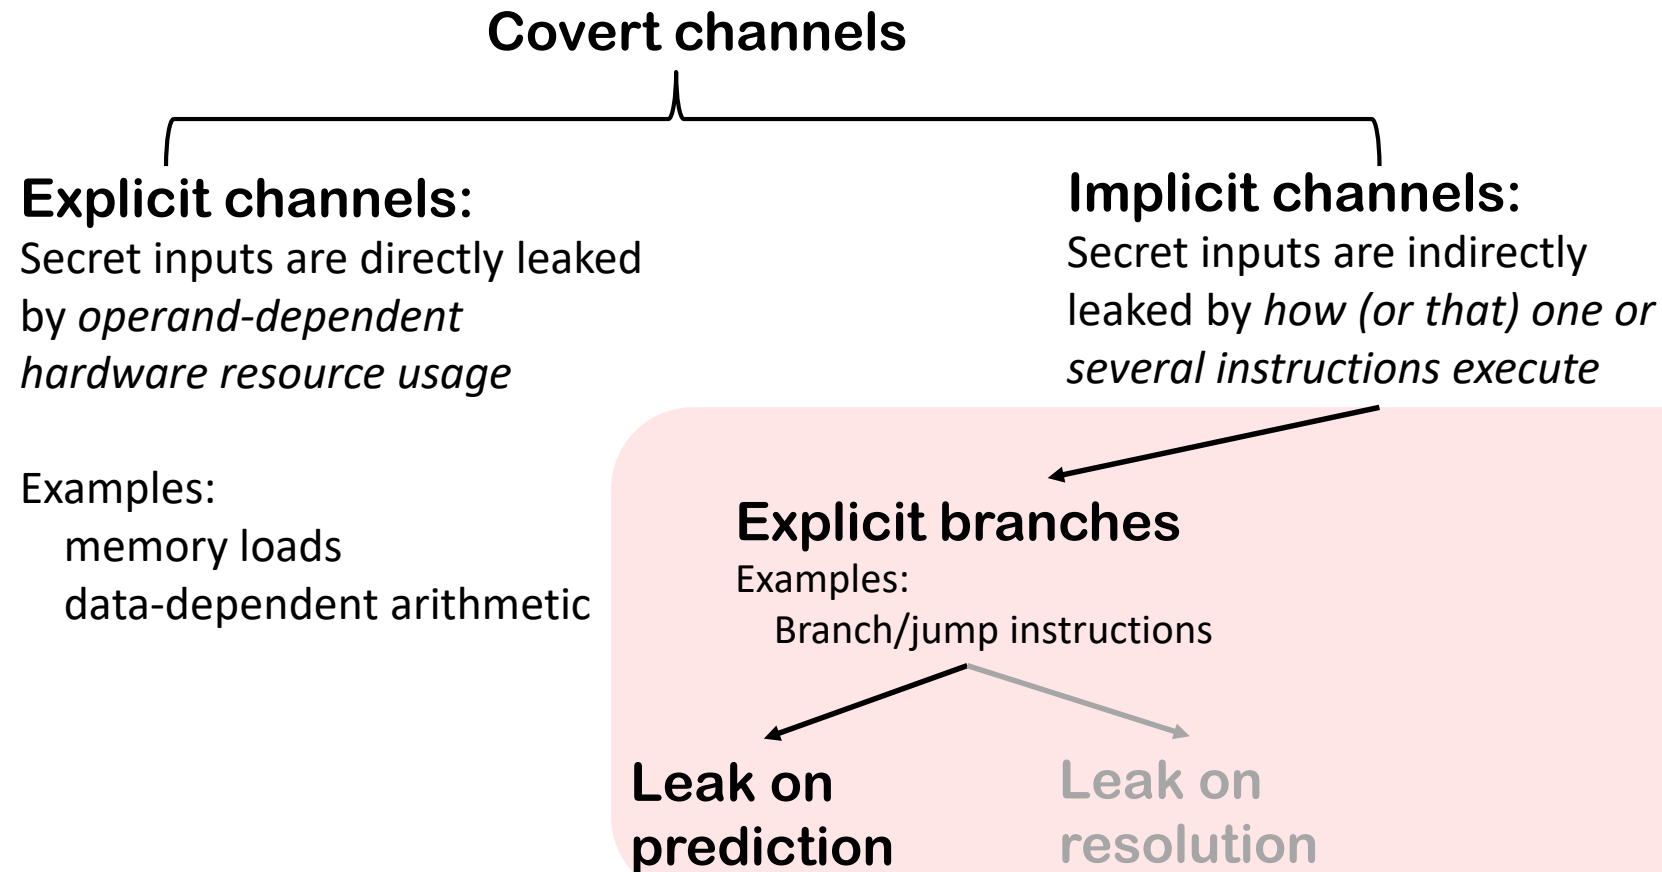
Examples:

- branch/jump instructions

Classification of Covert Channels



Classification of Covert Channels



Explicit Branches @ Prediction

Cause:

The predictor state becomes a function of secret

```
... ..  
... ..  
if ( secret )  
... ..  
... ..  
if ( public )  
    load [0x00];  
else  
    load [0x10];
```

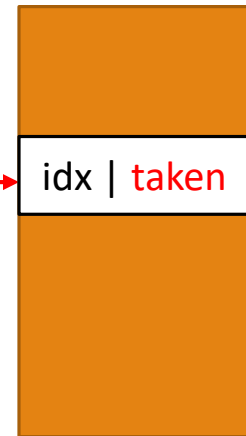
Explicit Branches @ Prediction

Cause:

The predictor state becomes a function of secret

```
... ..  
... ..  
if ( secret )  
... ..  
... ..  
if ( public )  
    load [0x00];  
else  
    load [0x10];
```

Resolve and update
branch predictor



Branch Predictor Unit (BPU)

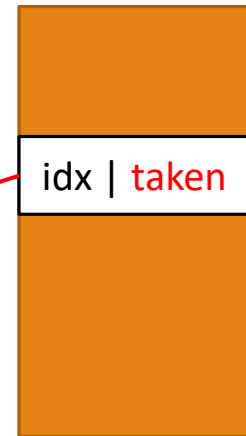
Explicit Branches @ Prediction

Cause:

The predictor state becomes a function of secret

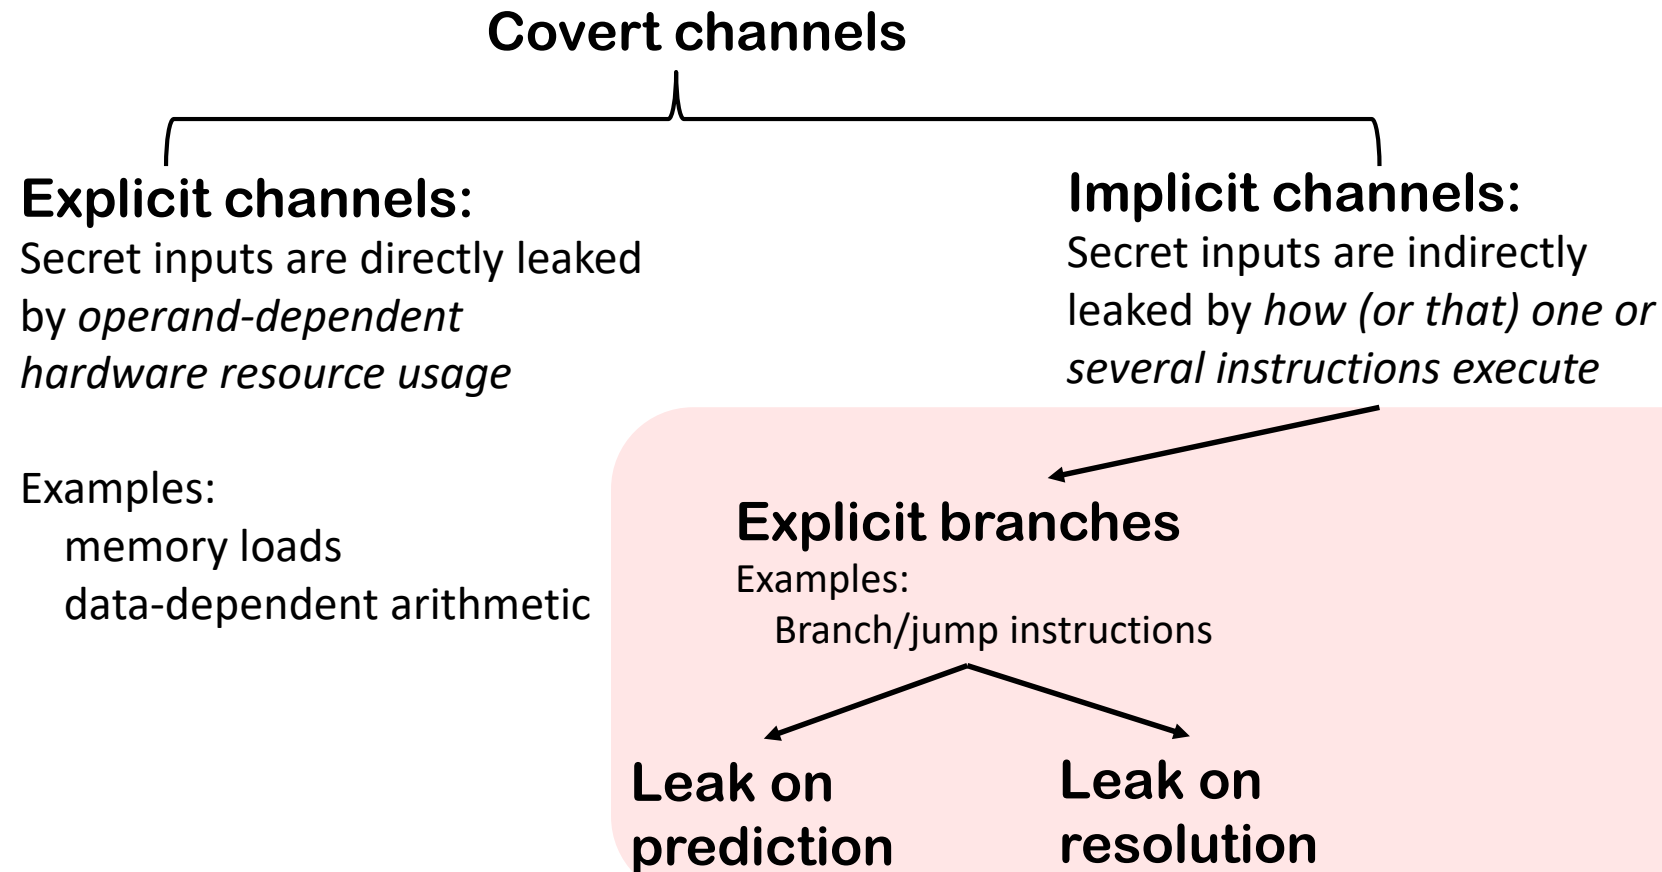
```
... ..  
... ..  
if ( secret )  
... ..  
... ..  
if ( public )  
    load [0x00];  
else  
    load [0x10];
```

Use BPU entry to predict



Branch Predictor Unit (BPU)

Classification of Covert Channels



Explicit Branches @ Resolution

Cause:

The resolution of a mis-speculation
triggers a pipeline squash and
alternation of control flow

```
if (secret) {  
    y++;  
}  
z = load [0x00]
```

Explicit Branches @ Resolution

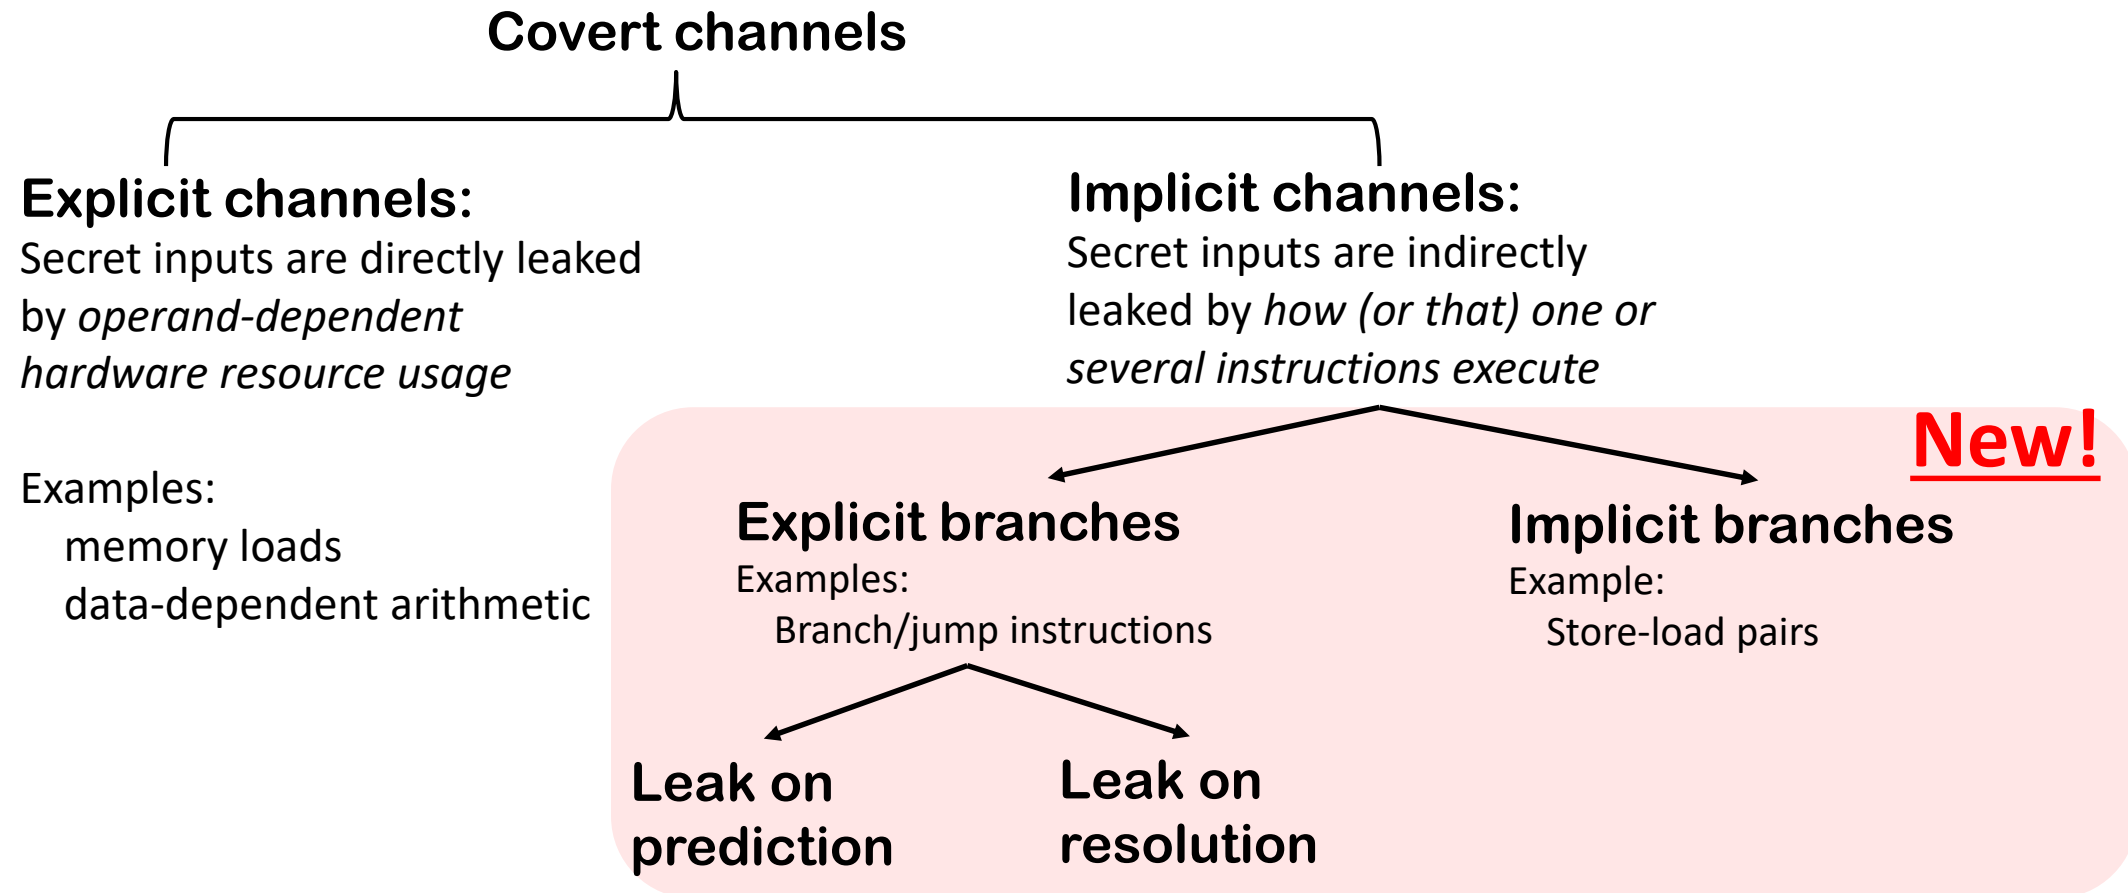
Cause:

The resolution of a mis-speculation
triggers a pipeline squash and
alternation of control flow

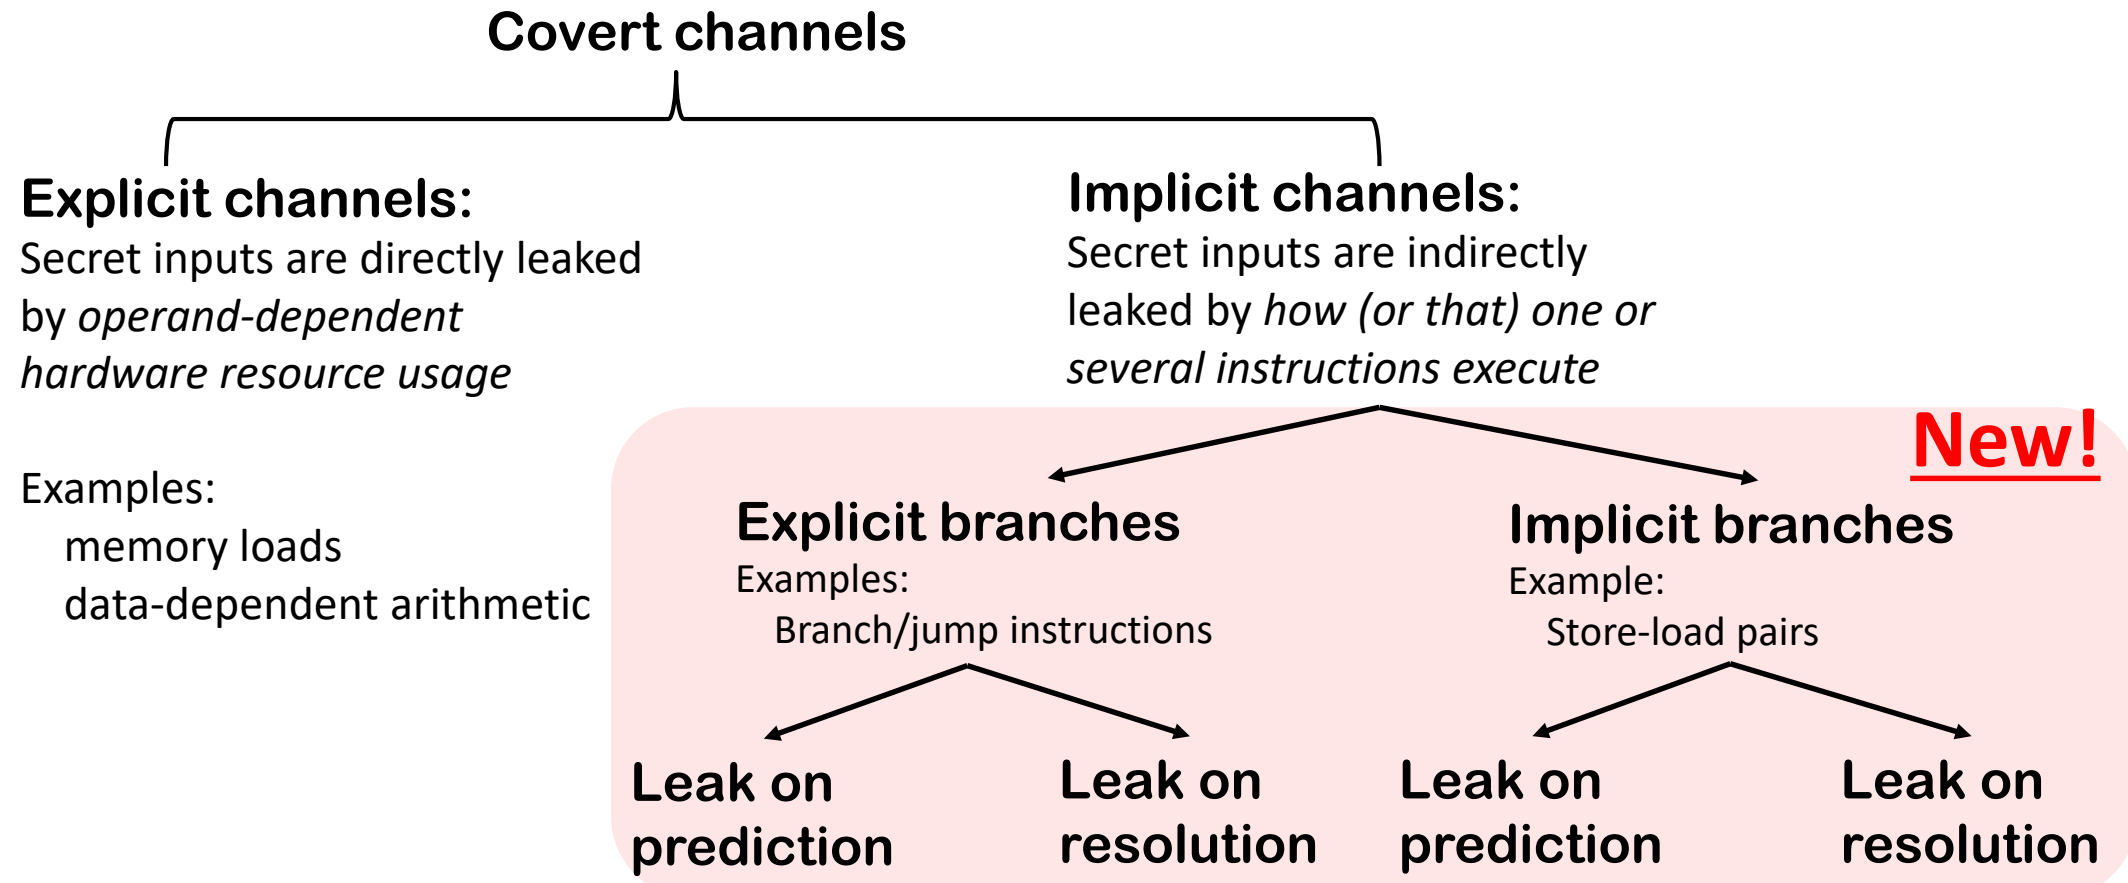
```
if (secret) {  
    y++;  
}  
z = load [0x00]
```

secret != prediction
→ squash
→ load executes twice!

Classification of Covert Channels



Classification of Covert Channels



Implicit Branches

Cause:

Non-control flow instructions create branch-like behaviors.

```
store [secret] = foo;
```

```
bar = load [0x00];
```

Implicit Branches

Cause:

Non-control flow instructions create branch-like behaviors.

```
store [secret] = foo;
```

```
bar = load [0x00];
```

Can be thought as:

```
if (secret == 0x00) {  
    forward from store queue  
}  
else {  
    cache_load [0x00]  
}
```

Identifying Secrets using Tainting/Untainting

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

```
if (addr < N) {
```

```
    // access instruction  
    a = load [addr];
```

```
    // simple arithmetic  
    b = a + 4;
```

```
    // cache/mem covert channel  
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

STT *taints*:

- 1) Output of speculative access instructions (a)

```
if (addr < N) {
```

```
    // access instruction  
    a = load [addr];
```

```
    // simple arithmetic  
    b = a + 4;
```

```
    // cache/mem covert channel  
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

STT *taints*:

- 1) Output of speculative access instructions (a)
- 2) Output of instructions with tainted inputs (b)

```
if (addr < N) {
```

```
    // access instruction  
    a = load [addr];
```

```
    // simple arithmetic  
    b = a + 4;
```

```
    // cache/mem covert channel  
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

STT *taints*:

- 1) Output of speculative access instructions (a)
- 2) Output of instructions with tainted inputs (b)

```
if (addr < N) { ← Resolved!
```

```
    // access instruction
    a = load [addr];
```

```
    // simple arithmetic
    b = a + 4;
```

```
    // cache/mem covert channel
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

STT *taints*:

- 1) Output of speculative access instructions (a)
- 2) Output of instructions with tainted inputs (b)

STT *untaints* when:

- 1) A speculative access instruction becomes non-speculative (a)

```
if (addr < N) { ← Resolved!
```

```
    // access instruction
    a = load [addr];
```

```
    // simple arithmetic
    b = a + 4;
```

```
    // cache/mem covert channel
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Identifying Secrets using Tainting/Untainting

Basic idea: taint all the secrets

- Speculatively accessed data (secrets by definition)
- And their dependents

STT *taints*:

- 1) Output of speculative access instructions (a)
- 2) Output of instructions with tainted inputs (b)

STT *untaints* when:

- 1) A speculative access instruction becomes non-speculative (a)
- 2) An instruction has all its input untainted (b)

```
if (addr < N) { ← Resolved!
```

```
    // access instruction
    a = load [addr];
```

```
    // simple arithmetic
    b = a + 4;
```

```
    // cache/mem covert channel
    load [b];
```

```
}
```

```
.....
```

```
.....
```

```
.....
```

speculative

Microarchitect Identifies ...

Instructions forming **explicit channels**

- E.g. load, data-dependent arithmetic

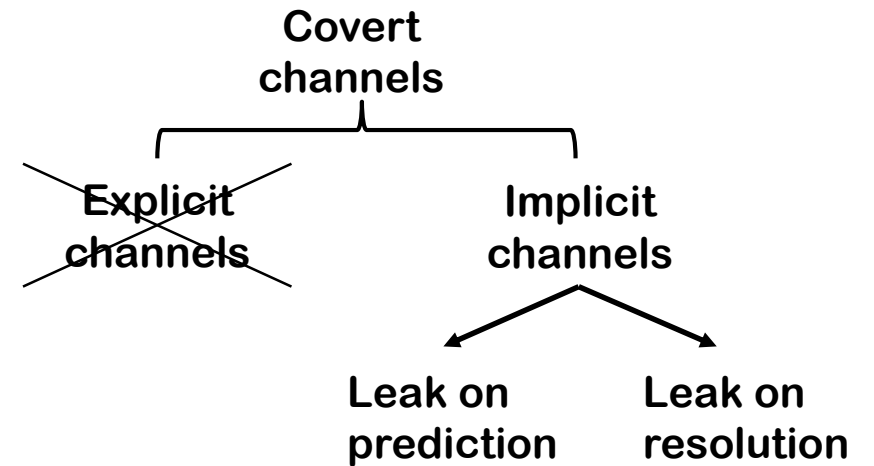
Instructions forming **implicit channels**

- E.g. control-flow instructions, store-load pairs

Blocking Covert Channels

Explicit channels:

- Delay execution until operands **untainted** (e.g., load address)



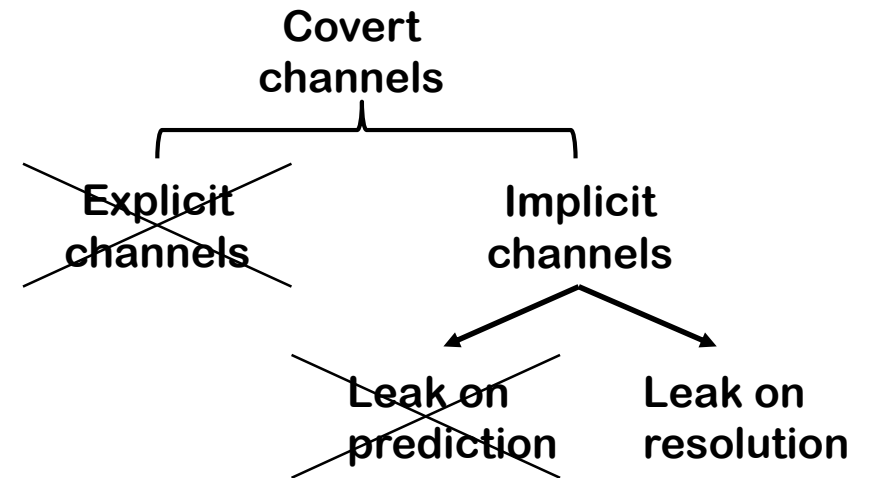
Blocking Covert Channels

Explicit channels:

- Delay execution until operands **untainted** (e.g., load address)

Implicit channels:

- Delay predictor update until branch predicate **untainted**



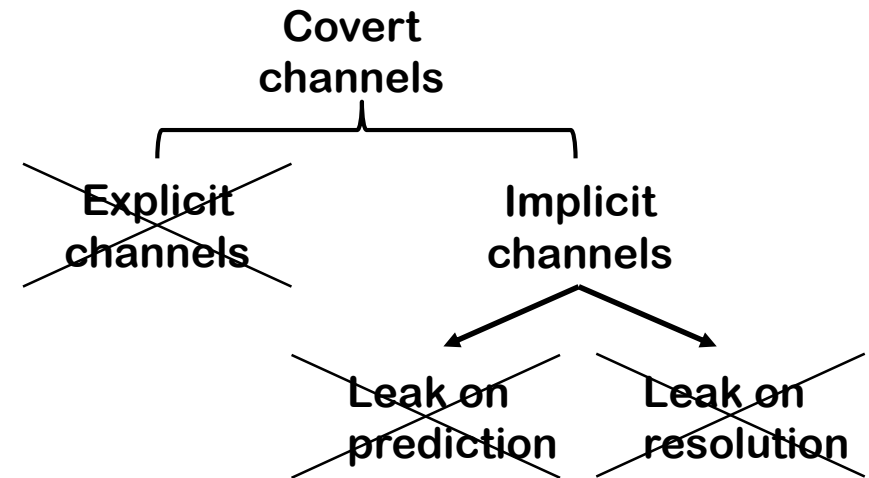
Blocking Covert Channels

Explicit channels:

- Delay execution until operands **untainted** (e.g., load address)

Implicit channels:

- Delay predictor update until branch predicate **untainted**
- Delay resolution until branch predicate **untainted**



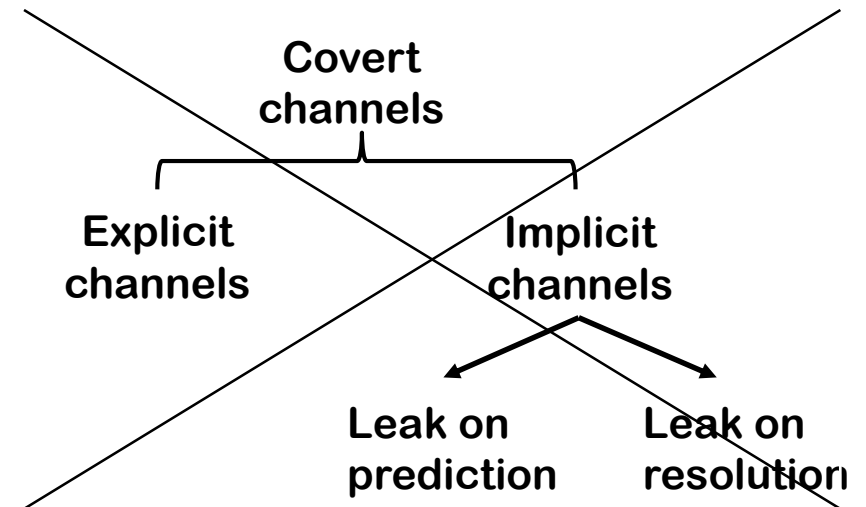
Blocking Covert Channels

Explicit channels:

- Delay execution until operands **untainted** (e.g., load address)

Implicit channels:

- Delay predictor update until branch predicate **untainted**
- Delay resolution until branch predicate **untainted**



Hardware Implementation of STT

Efficient Implementation of Tainting/Untainting Logic

program order

1) branch

2) **a** = load [0x00]

3) branch

4) **b** = load [0x04]

5) branch

6) **c** = **a** + **b**

7) load [**c**]

Delay execution!

speculative

Efficient Implementation of Tainting/Untainting Logic

program order

1) branch

2) **a** = load [0x00]

3) branch

4) **b** = load [0x04]

5) branch

6) **c** = **a** + **b**

7) load [**c**]

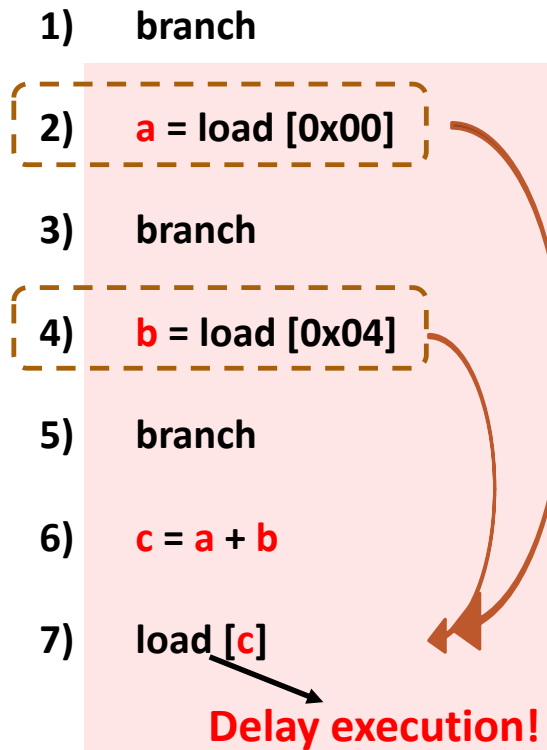
Delay execution!

speculative

Efficient Implementation of Tainting/Untainting Logic

Observation: All instructions turn non-speculative in-order

program order



speculative

Efficient Implementation of Tainting/Untainting Logic

Observation: All instructions turn non-speculative in-order

program order

1) branch → resolved!

2) **a** = load [0x00]

3) branch

4) **b** = load [0x04]

5) branch

6) **c** = **a** + **b**

7) load [**c**]

Delay execution!

speculative

Efficient Implementation of Tainting/Untainting Logic

Observation: All instructions turn non-speculative in-order

program order

1) branch

2) **a** = load [0x00]

3) branch

4) **b** = load [0x04]

5) branch

6) **c** = **a** + **b**

7) load [**c**]

→ resolved!

Execute!

speculative

Efficient Implementation of Tainting/Untainting Logic

Observation: All instructions turn non-speculative in-order

Each instruction tracks the “youngest access instruction” it depends on -- “***Youngest Root of Taint***” (***YRoT***)

program order

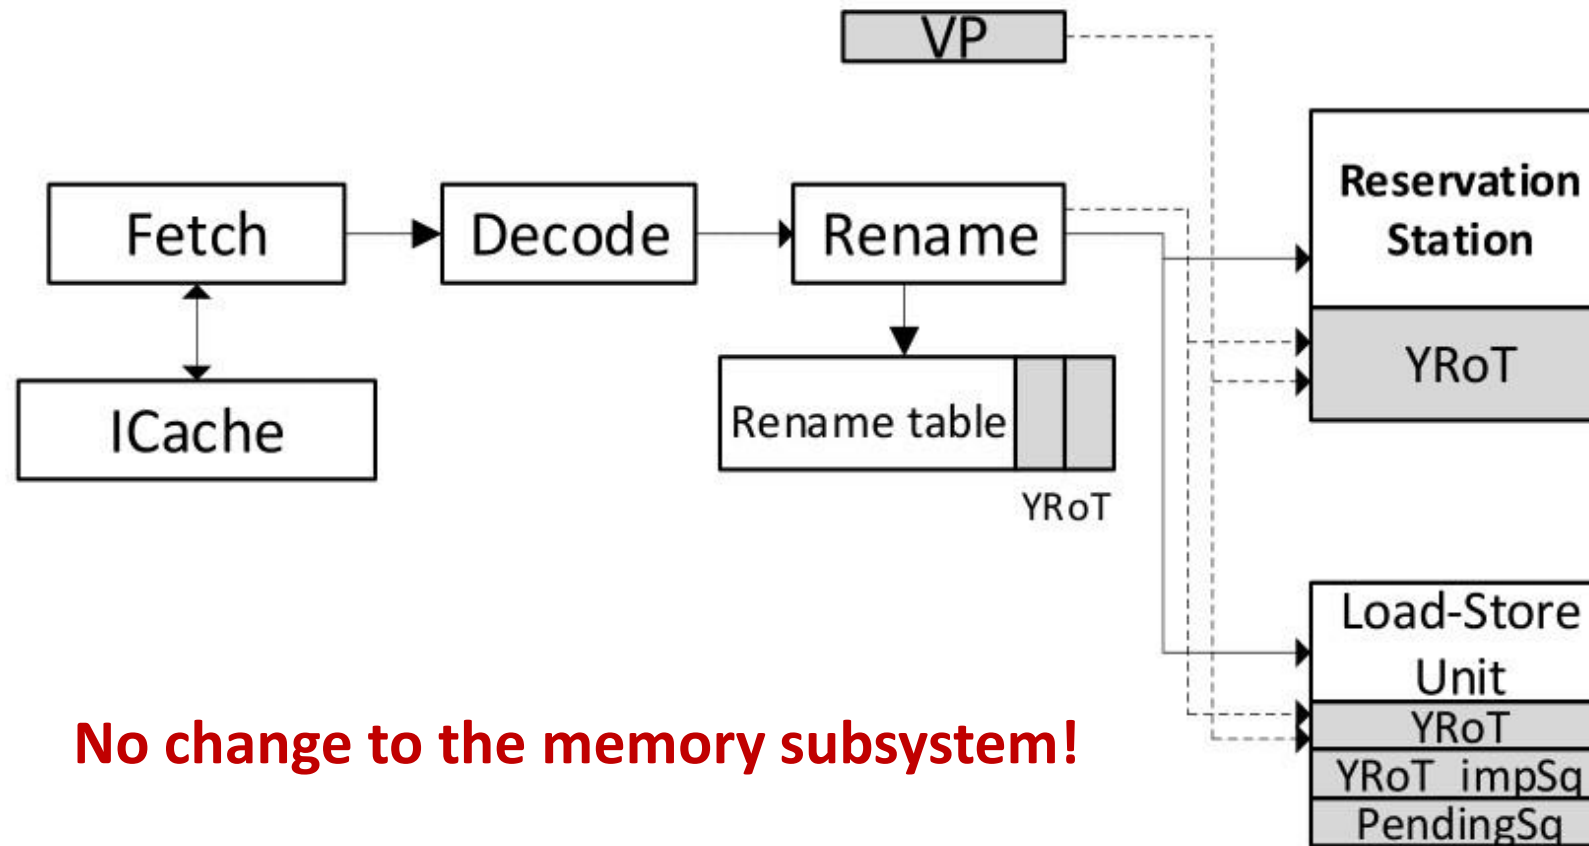
- 1) branch
- 2) a = load [0x00]
- 3) branch
- 4) **b = load [0x04]**
- 5) branch
- 6) **c = a + b**
- 7) load [**c**]

YRoT of 7 is 4

Execute!

speculative

Efficient Implementation of Tainting/Untainting Logic



Security Evaluation

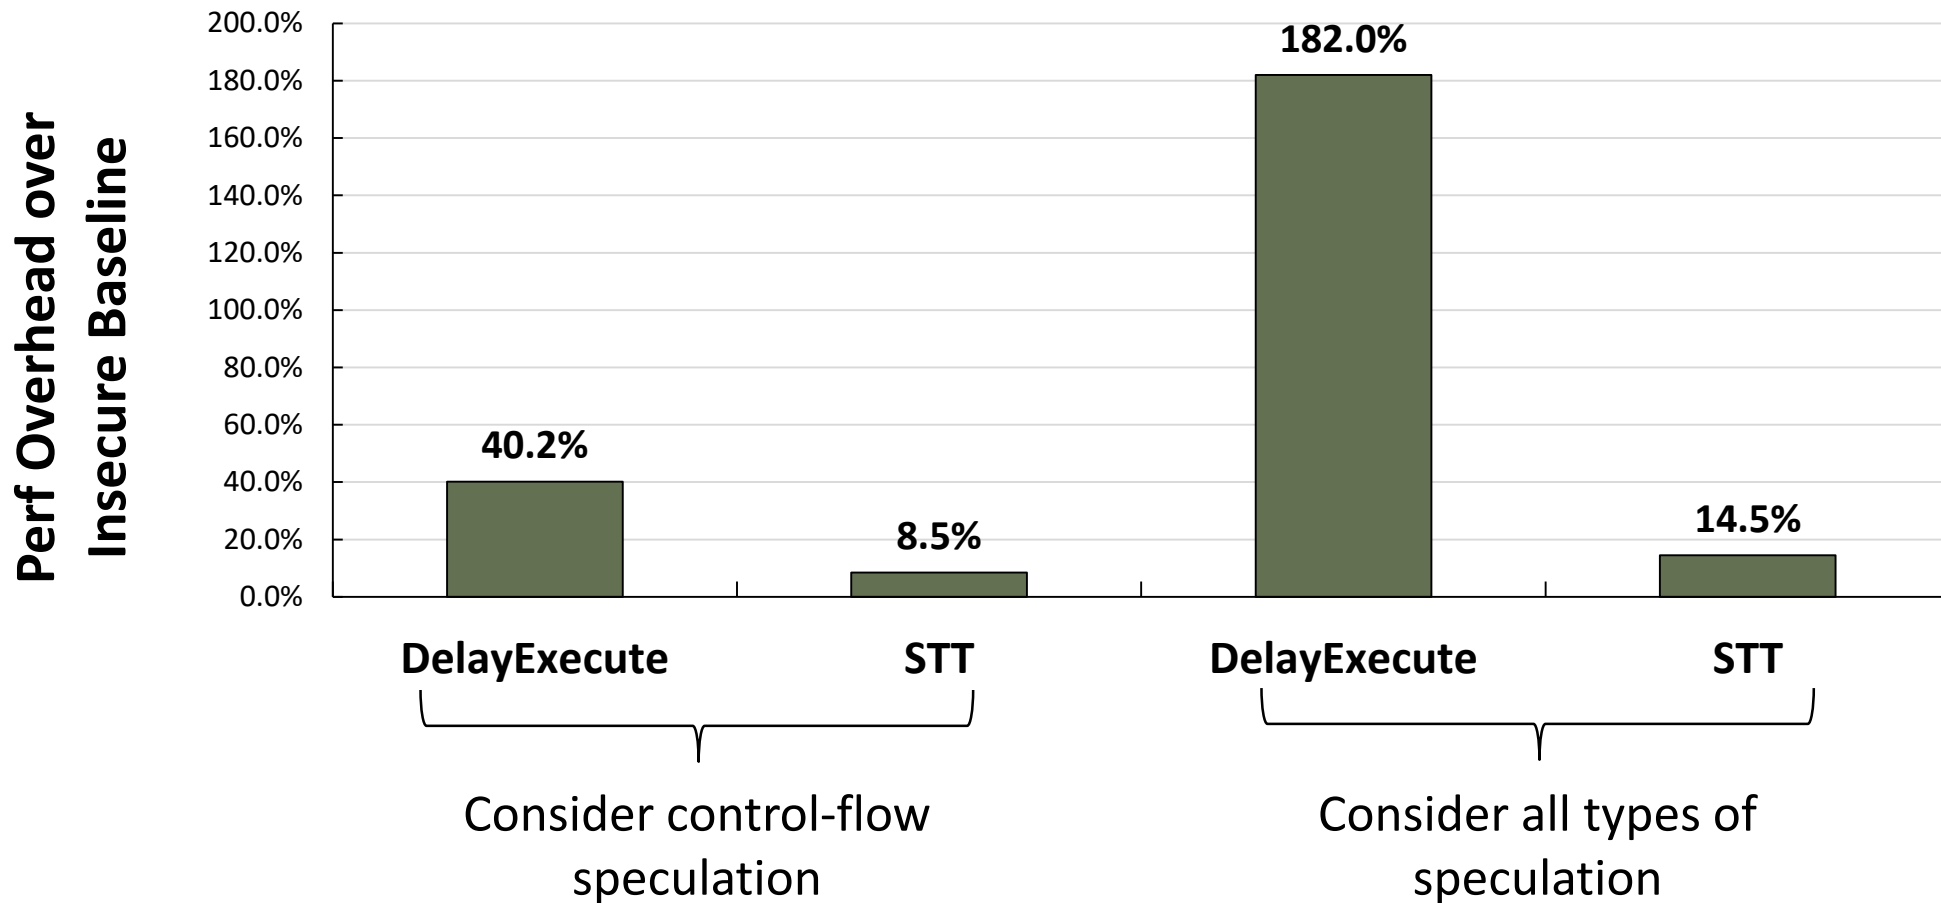
Security definition:

**Arbitrary speculative execution can only leak
retired register file state (not arbitrary program memory)**

To prove it: STT enforces a non-interference property w.r.t speculatively accessed data

The link to the **detailed formal analysis and security proof** is in the paper

Performance Evaluation on SPEC2006



Conclusion

STT Blocks leakage of speculatively accessed data over any uarch covert channels with:

- 1) High performance
- 2) Provable security protection
- 3) No software change; No memory subsystem change

Conclusion

STT Blocks leakage of speculatively accessed data over any uarch covert channels with:

- 1) High performance
- 2) Provable security protection
- 3) No software change; No memory subsystem change

Questions?

Backup slides

Threat Model

A powerful attacker who:

- Monitors covert channels
 - Cache, SIMD units, or any shared hardware resources
- From everywhere
 - Within same thread
 - Adjacent SMT context
 - Cross core
- At cycle granularity

TODO: compare with NDA and SpecShield

TODO: STT threat model vs. STT+ threat model

Outline

Introduction

Threat Model

Speculative Taint Tracking

Evaluation

Conclusion

Implementation of STT

A new classification of covert channels in HW

➔ Specify instructions with explicit or implicit channels

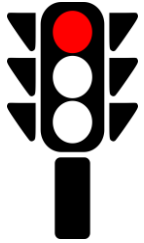
Putting it together



Executes all instructions w/ untainted inputs

Executes non-transmit instructions w/ tainted inputs

Predicts explicit/implicit branches w/ tainted predicates



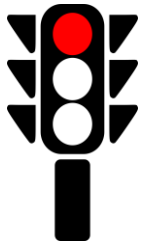
Putting it together



Executes all instructions w/ untainted inputs

Executes non-transmit instructions w/ tainted inputs

Predicts explicit/implicit branches w/ tainted predicates



Delay executing transmit instructions w/ tainted inputs

Delay resolution/predictor updates of explicit/implicit branches w/ tainted predicates



Block explicit channel

Putting it together



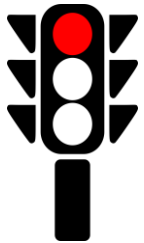
Executes all instructions w/ untainted inputs

TODO: spend a little bit more time on this slide?

Executes non-transmit instructions w/ tainted inputs

Predicts explicit/implicit branches w/ tainted predicates

TODO: what is transmit instruction?



Delay executing transmit instructions w/ tainted inputs



Block explicit channel

Delay resolution/predictor updates of explicit/implicit branches w/ tainted predicates



Block implicit channel

Microarchitecture Design of STT

- Each instruction has its “*Youngest Root of Taint*” (YRoT)


For each (transmit) instruction:

Input is secret

⇔ YRoT is still speculative

⇔ Visibility point is ahead of YRoT

program order

- 
- 1) branch
 - 2) `a = load [0x00]` // YRoT = -1
 - 3) branch_1
 - 4) `b = load [0x04]` // YRoT = -1
 - 5) branch
 - 6) `c = a + b` // YRoT = $\max(2, 4) = \underline{4} > \underline{1}$
 - 7) `d = load [c]` // YRoT = $\underline{4} > \underline{1}$

Microarchitecture Design of STT

- Each instruction has its “*Youngest Root of Taint*” (YRoT)

For each (transmit) instruction:

Input is secret

⇔ YRoT is still speculative

⇔ Visibility point is ahead of YRoT

program order

- 1) branch_0
- 2) a = load [0x00] // YRoT = -1
- 3) branch_1
- 4) b = load [0x04] // YRoT = -1
- 5) branch_2
- 6) **c** = a + b // YRoT = max(2, 4) = 4 > 3
- 7) d = load [**c**] // YRoT = 4 > 3

Microarchitecture Design of STT

- Each instruction has its “*Youngest Root of Taint*” (YRoT)

For each (transmit) instruction:

Input is secret

⇔ YRoT is still speculative

⇔ Visibility point is ahead of YRoT

program order

- 1) branch_0
- 2) a = load [0x00] // YRoT = -1
- 3) branch_1
- 4) b = load [0x04] // YRoT = -1
- 5) branch_2 ← VP
- 6) c = a + b // YRoT = max(2, 4) = 4 < 5
- 7) d = load [c] // YRoT = 4 < 5

Efficient Implementation of Tainting/Untainting Logic

Observation: All instructions turn non-speculative in-order

Each instruction tracks the “youngest access instruction” it depends on -- “**Youngest Root of Taint**” (YRoT)

For each instruction:

Input is tainted

TODO: remove this

⇔ Input depends on some speculative access instruction

⇔ YRoT is still speculative

⇔ Visibility point is ahead of YRoT

program order

1) branch

2) a = load [0x00]

3) branch

4) **b** = load [0x04]

5) branch

6) **c** = a + **b**

7) load [**c**]

Execute!

Visibility point (VP)

YRoT of 7 is 4

speculative

Security Evaluation

Arbitrary speculative execution can
only leak retired register file state
(not arbitrary program memory)

Security Evaluation

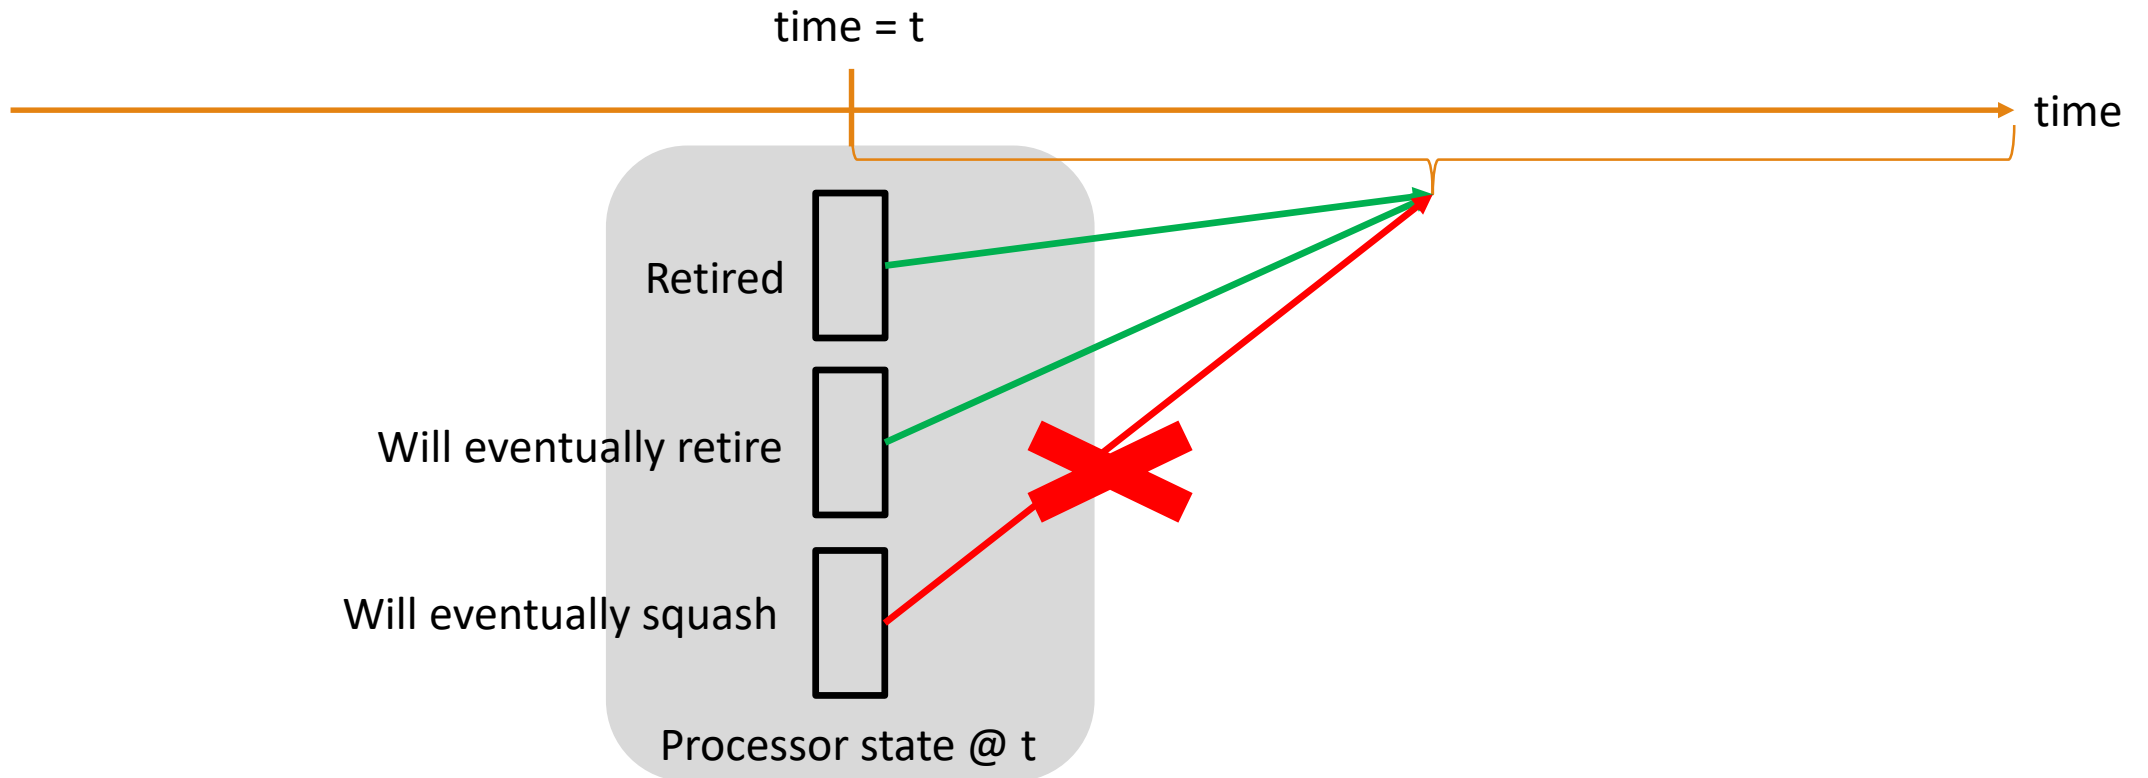
Arbitrary speculative execution can
only leak retired register file state
(not arbitrary program memory)



The Universal Read Gadget == many Spectre
variants (1, 3, 4, ..), MDS attacks, Meltdown, etc.

Security Evaluation

STT enforces a non-interference property w.r.t speculatively accessed data:



Security Evaluation

STT enforces a non-interference property w.r.t speculatively accessed data:

